Optimal sin taxation and market power

Martin O'Connell and Kate Smith*

November, 2022

Abstract

We study how market power impacts the efficiency and redistributive properties of sin taxation, with an empirical application to sugar-sweetened beverage taxation. We estimate a detailed equilibrium model of the UK drinks market, which we embed in a tax design framework to solve for optimal sugar-sweetened beverage tax policy. Positive price-cost margins on drinks create inefficiencies, which act to lower the optimal rate compared with a perfectly competitive setting. However, since profits mainly accrue to the rich, this is partially mitigated under social preferences for equity. Overall, ignoring market power when setting the optimal sugar-sweetened beverage tax rate leads to welfare gains that are 40% below those at the optimum.

Keywords: externality, corrective tax, market power, profits, redistribution

JEL classification: D12, D43, D61, D62, H21, H23, L13

Acknowledgements: The authors would like to thank participants at seminars at Bristol, Brussels, Chicago Booth, INRAE, Michigan, Oxford, Warwick, UCL, UIUC, WUSTL and Wisconsin-Madison for many very helpful discussions. The authors gratefully acknowledge financial support from the Economic and Social Research Council (ESRC) under the Centre for the Microeconomic Analysis of Public Policy (CPP; grant number ES/T014334/1), under the Open Research Area (ORA; grant number ES/VO13513/1) and the British Academy under pf160093. Data supplied by Kantar FMCG Purchase Panel. The use of Kantar FMCG Purchase Panel data in this work does not imply the endorsement of Kantar FMCG Purchase Panel in relation to the interpretation or analysis of the data. All errors and omissions remain the responsibility of the authors.

*O'Connell is at the University of Wisconsin-Madison and the Institute for Fiscal Studies and Smith is at the London School of Economics and the Institute for Fiscal Studies. **Correspondence**: moconnell9@wisc.educ and k.e.smith2@lse.ac.uk

1 Introduction

One-fifth of all consumer spending is undertaken in markets subject to taxes aimed, at least in part, at discouraging socially costly consumption.[1] Many of these markets are concentrated and likely to be characterized by firms that exercise market power. For instance, soft drink markets, recently the subject of new taxes in several jurisdictions (GFRP (2021)), are dominated by a small number of multi-product firms selling instantly recognizable and long-established brands viewed by consumers as imperfect substitutes.[2] The effective design of sin taxes requires balancing efficiency improvements and policymakers' redistributive goals. How best to do this depends on the competitive conditions in affected markets. Positive price-cost margins create allocative distortions, in addition to those associated with social costs of consumption; firms' strategic adjustment of margins affects the mapping between taxes, equilibrium prices and consumption; and concentration of profit ownership among the rich[3] impacts the distributional consequences of reforms.

In this paper we show how market power affects the optimal design of sin taxes, and undertake a substantive empirical application to the taxation of sugar-sweetened beverages. We analyze a tax design framework, which highlights how market power impacts both the efficiency and redistributive effects of policy, and motivates our empirical approach. Harnessing detailed consumer level data, we estimate an equilibrium model of the UK market for drinks, and embed this into the tax problem to solve for optimal sugar-sweetened beverage tax policy. We show that allocative distortions associated with positive margins on drinks reduce the optimal sin tax rate, compared with a perfectly competitive setting, but this reduction is partially offset by the fact that profits accrue disproportionately to those with lower social welfare weights (foreign residents and the rich). Overall, ignoring market power leads to welfare gains that are 40% below those achieved by optimal policy.

Our tax design framework entails an environment with heterogeneous consumers and strategic firms. We study a government's choice over a set of linear commodity taxes levied in a specific market of interest that comprises "sin" products, which generate social costs and alternatives, which do not create social costs. The products are supplied by multi-product firms that exercise market power, strategically

[1]Spending on alcohol, tobacco, soft drinks, fuel, and motoring (all of which are subject to some kind of excise duty in the UK – Levell et al. (2016)) accounts for 19% of consumer spending recorded in the UK's consumer expenditure survey (Living Costs and Food Survey (2017)).

[2]More generally, a large industrial organization literature demonstrates that market power is a typical feature of many markets (see Einav and Levin (2010)), and there is recent evidence of its growing importance – for instance, see Syverson (2019) and De Loecker et al. (2020).

[3]See, for example, evidence in Cooper et al. (2016).

reoptimize prices in response to tax policy, and earn positive profits, which are distributed to individuals in a potentially unequal manner. The government sets policy anticipating the strategic responses of firms,[4] accounts for spillovers to other tax bases, and may place more weight on the welfare of the poor than the rich. We show how optimal policy balances: (i) reducing the social costs of consumption, (ii) exacerbating inefficiencies arising from the exercise of market power, (iii) fiscal externalities, and (iv) redistributive effects that arises both from inequality in consumption patterns, and inequality in profit holdings.

We apply our framework to study the taxation of sugar-sweetened beverages, a policy that is increasingly popular due, in part, to the large externalities that arise from the healthcare costs associated with treating obesity. The combination of asymmetrically differentiated products, multi-product firms and heterogeneous consumers means that the optimal tax expressions that we derive do not straightforwardly translate into externally valid elasticities. They depend instead on product-specific price-cost margins (which are typically unobservable, see Bresnahan (1989)), firm portfolio effects, and the full set of product level own- and cross-price elasticities. We therefore estimate a detailed model of consumer choice and firm pricing competition in the UK drinks market and embed this into our tax design framework. We model the discrete choice that consumers make between the many products in the market (or allocating all of their spending outside the market), and estimate the model using micro longitudinal data. We incorporate rich preference heterogeneity, including by income and a measure of total dietary sugar, which are important for capturing both redistributive and externality correcting aspects of tax policy. We exploit price changes that are agreed in advance by drinks firms and retailers, but that create differential variation across consumers, as a source of identifying price variation. We identify product level marginal costs by coupling our demand estimates with a game-theoretic pricing model.

Our demand and supply estimates show that drinks products are highly differentiated and, due to this and the multi-product nature of their portfolios, drinks firms exercise a substantial degree of market power. We find that the average Lerner index at observed prices is around 0.5 for both sugar-sweetened drinks and alternatives. Our estimates show that, in response to a change in the price of one sugar-sweetened drink, consumers are most willing to switch to a similar sugar-sweetened drink. However, in response to a rise in the price of all sugary drinks, substitution to alternative (non sugar-sweetened) drinks is substantial. Market power among alternative drinks is therefore relevant for sugar-sweetened beverage

[4] This is also the case in Miravete et al. (2018), who show quantitatively how firms' pricing responses change the peak and shape of the Laffer curve in the liquor market.

tax policy. Our model allows us to determine how firms adjust prices (and hence price-cost margins) in response to tax policy. We find that pass-through of a tax on sugar-sweetened drinks is slightly above 100%, and show that this is in line with (out-of-sample) observed price changes following the introduction of the UK's Soft Drinks Industry Levy, as well as with findings in other jurisdictions.

We use our estimated model to first solve for the optimal sin tax rate set by a planner concerned only with inefficiencies arising from externalities and market power, and that is indifferent to the distribution of incidence across individuals. Consumption of sugar-sweetened beverages is strongly linked to diet-related disease, which creates externalities through increased societal costs of funding both public and insurance-based health care (see Allcott et al. (2019b)). We use the best available evidence from the epidemiological literature to quantify these costs. The efficiency-maximizing tax rate balances the reduction in social costs of consumption – the more effective tax is at combating externality distortions, all else equal, the higher will be the optimal rate – against exacerbating distortions from the exercise of market power. If equilibrium price-cost margins for the sin products are high relative to untaxed alternative products, this will act to lower the optimal rate.[5] We find that the efficiency-maximizing tax rate on sugar-sweetened beverages is 4 pence per 10 oz, which leads to a 19% rise in their average price. If the government ignores all distortions from market power it would set a suboptimally high rate of 12 pence per 10 oz, which leads to efficiency *losses*.

We then consider optimal sin tax policy, assuming that the government does have redistributive concerns. In this case, both the share of profits collected by the government in corporate and dividend taxes, and the distribution of net-of-tax profits across individuals, affect optimal policy. We measure the allocation of corporate profits using information from national accounts and the distribution of dividend income across households. We find that under our baseline social preferences, the optimal sin tax rate on sugar-sweetened beverages is 6 pence per 10 oz,[6] which is 50% higher than the efficiency-maximizing rate. This increase reflects the net impact of two offsetting forces. On the one hand, sugar-sweetened beverage consumption is highest among relatively low income consumers, which acts to lower the optimal rate. On the other hand, post-tax profits are disproportionately owned

[5]That the market power correction lowers the optimal rate if, on average, sin products have high margins *relative to alternatives* is consistent with the results in Kaplow (2021), who shows that in a Mirrleesian economy in which goods have fixed positive price-cost margins it is Pareto-improving to marginally lower/raise the tax rate on any single good that has an above/below average margin whilst off-setting the distributional effect through a distribution neutral adjustment to the labor tax schedule.

[6]In practice, US and UK sugar-sweetened beverage taxes range from 7 to 15 pence per 10 oz.

by high income consumers (or flow overseas), which increases the progressivity of the tax, thus raising the optimal rate. Overall, the second effect dominates. We show that these results apply under a wide range of social preferences, including when the government places zero welfare weight on the post-tax profits flowing to individuals (in which case ignoring market power when setting policy would lead it to overshoot the optimal rate by 70% – with the resulting welfare gains 26% below optimal – due to spillovers to the corporate and dividend tax bases).

We also explore how the nature of social costs, including their overall size, degree of concentration among a specific group of individuals, and composition between externalities and internalities impact optimal policy. In addition, we simulate optimal policy under alternative tax instruments, including a multi-rate system and a sugar tax. In all cases market power has a first order impact on the efficiency and redistributive impacts of taxation, and ignoring leads to substantial unrealized welfare gains.

Our paper contributes to a growing literature on quantifying the impacts of sugar-sweetened beverage taxation. This includes a set of papers that use the implementation of these taxes to estimate their effect on prices and quantities,[7] and a set that use estimates of consumer demand based on periods and locations with no tax in place to simulate the introduction of taxes similar to those used in practice.[8] Included in the second set of papers is Dubois et al. (2020), who use novel UK panel data to estimate individual-specific preference parameters for "on-the-go" soft drinks. We extend this work by estimating consumer demand across the whole non-alcoholic drinks sector, including drinks brought into the home, modeling the supply-side of the market, and studying optimal policy design. Like us, Allcott et al. (2019a) ask: what is the optimal sugar-sweetened beverage tax? They develop a novel characterization of the optimal corrective commodity tax rate under general preference heterogeneity and an optimally set labor tax schedule, and incorporate consumer misoptimization into the corrective tax motive. Our work complements theirs by showing how market power impacts optimal sin tax policy.

Our work also adds to the emerging literature that uses empirically rich models of specific markets to evaluate how imperfect competition impacts public policy design. This literature includes work showing how the existence, and firms' strategic adjustment, of mark-ups impacts: revenue- and efficiency-maximizing taxation

[7]See, for instance, Bollinger and Sexton (2018) and Rojas and Wang (2017) who study the Berkeley tax, Seiler et al. (2021) and Roberto et al. (2019) who study the Philadelphian tax, and Grogger (2017) who study the Mexican tax. For a full survey of the recent literature see Griffith et al. (2019).

[8]These papers include Bonnet and Réquillart (2013), Wang (2015), Harding and Lovenheim (2017), Chernozhukov et al. (2019).

in alcohol markets (Conlon and Rao (2015), Miravete et al. (2018, 2020)); the effectiveness and optimal design of carbon abatement policy (Fowlie et al. (2016), Preonas (2019)); and the design of demographic-targeted subsidies in health insurance markets (Tebaldi (2017), Polyakova and Ryan (2019), Einav et al. (2019)). We contribute to this literature by quantifying how market power alters both the efficiency and redistributive effects, and hence optimal design, of sugar-sweetened beverage taxation. More generally, our results illustrate the importance of accounting for market power on commodity tax design.

The rest of the paper is structured as follows. In Section 2 we outline a tax design framework that highlights the key ingredients determining optimal policy. In Section 3 we summarize key characteristics of the UK drinks market. In Sections 4 and 5 we present our empirical model of market demand and supply, and our estimates. In Section 6 we present our optimal tax results.

2 Sin tax design

We begin by outlining a tax design framework, which serves to highlight the key determinants of optimal sin tax policy with market power, and which we use when solving for optimal sugar-sweetened beverage taxation. We consider a setting in which a government chooses tax policy in a market (for instance, the drinks market) with multiple products, including a set associated with social costs (e.g., sugar-sweetened beverages). We allow for the possibility that these products are sold by firms that exercise market power. The government sets the tax rates while accounting for interactions with other parts of the tax system, and balances distortions from social costs and market power with equity considerations.

2.1 Set-up

Individuals

There is a continuum of individuals (or consumers) indexed i. Individuals supply labor in a competitive labor market to generate pre-tax earnings, z^i, which are subject to a non-linear earnings tax, $\mathcal{T}(z^i)$. Each individual also potentially receives income arising from their holding of profits, which are generated by the sale of consumption goods. We denote total profits by Π and i's share of profits by $\delta^i \geq 0$, where $\int_i \delta^i di = 1$. Individual profit holdings are subject to a (potentially non-linear) tax, $T_\Pi(\delta_i\Pi)$.[9]

Net income is spent on consumption goods: $\mathbf{x}_S^i = \{x_j^i\}_{j\in\mathcal{S}}$ is a quantity vector of "sin" products, belonging to the set $\mathcal{S}$, consumption of which potentially creates social costs; $\mathbf{x}_N^i = \{x_j^i\}_{j\in\mathcal{N}}$ is a quantity vector of products belonging to the set $\mathcal{N}$, which are in the same market as those in $\mathcal{S}$, but do not give rise to social costs; let $\mathcal{M} = \mathcal{S}\bigcup\mathcal{N}$ and $n(\mathcal{M}) = J$. x_O^i denotes the quantity of a composite consumption good that represents all goods outside of those in market $\mathcal{M}$. Consumers face the tax-inclusive price vector $\mathbf{p} = (\mathbf{p}_\mathcal{S}, \mathbf{p}_\mathcal{N}, 1)$, which embeds the normalization that the price of the composite good, x_O^i, is 1. The individual's budget constraint is $\sum_{j\in\mathcal{M}} p_j x_j^i + x_0^i = z^i - \mathcal{T}(z^i) + \delta^i\Pi - T_\Pi(\delta^i\Pi)$. We assume that the earnings tax is piece-wise linear, denote $\frac{d\mathcal{T}}{dz^i} \equiv \tau_z^i$, and define virtual labor income by $G \equiv \tau_z^i z^i - \mathcal{T}(z^i)$. We denote the sum of virtual and unearned (profit) income by $Y^i \equiv G^i + \delta^i\Pi - T_\Pi(\delta^i\Pi)$.

Each individual chooses a bundle, $(\mathbf{x}_S^i, \mathbf{x}_N^i, x_O^i, z^i)$, to maximize utility, $U^i(\mathbf{x}_\mathcal{S}^i, \mathbf{x}_\mathcal{N}^i, x_O^i, z^i)$, subject to their budget constraint. Consumer i's product demands are denoted by $x_j^i = x_j^i(\mathbf{p}, (1-\tau_z^i), Y^i)$, earnings supply by $z^i = z^i(\mathbf{p}, (1-\tau_z^i), Y^i)$, and indirect utility by $V^i = V^i(\mathbf{p}, (1-\tau_z^i), Y^i)$. We denote the marginal utility of income by $\alpha^i \equiv \partial V^i/\partial Y^i$.

Firms

Each product $j \in \mathcal{M}$ is produced by a single firm (firms can produce multiple products). We denote the market demand for product j by $X_j(\mathbf{p}_\mathcal{M}) = \int_i x_j^i di$, and we denote the product's marginal cost by c_j. We consider a system of linear excise taxes that apply to the products in market $\mathcal{M}$, $\boldsymbol{\tau}_\mathcal{M}$. Note that constant marginal costs means that if market $\mathcal{M}$ were perfectly competitive, prices would shift mechanically with commodity taxes.

We allow for the possibility that firms exercise market power, meaning they can set price above marginal cost and face positive demand. In equilibrium (where all firms choose their strategies to maximize their profit function), the tax-exclusive price for any product j ($\tilde{p}_j \equiv p_j - \tau_j$), can be written $\tilde{p}_j(\mathbf{c}_\mathcal{M}, \boldsymbol{\Upsilon}_\mathcal{M}; \boldsymbol{\tau}_\mathcal{M}) = c_j + \mu_j(\mathbf{c}_\mathcal{M}, \boldsymbol{\Upsilon}_\mathcal{M}; \boldsymbol{\tau}_\mathcal{M})$, where μ_j denotes the equilibrium price-cost margin of product j.[10]

[9] $T_\Pi(\delta_i\Pi)$ captures both corporate taxes and individual (e.g., dividend) taxation. For simplicity we write the earning and profit taxes as additively separable. This is unimportant for our analysis, which would not be materially affected by non-separabilities in $\mathcal{T}(.)$ and $T_\Pi(.)$.

[10] In this section we remain agnostic about the precise nature of the imperfect competition (e.g., whether products are offered by monopolist, competing oligopolists or colluding oligopolists). In Section 4 we assume the firms in the UK drinks market compete in a Nash-Bertrand pricing game.

For the composite good, which has aggregate demand $X_O = \int_i x_O^i di$, we assume that its price remains fixed in response to the introduction of an excise tax system in market $\mathcal{M}$, but we allow for the possibility of a non-zero price-cost margin, μ_O. Total variable profits in the economy are given by $\Pi \equiv \sum_{j \in \mathcal{M} \bigcup 0} \mu_j X_j$.[11]

Government

We consider a government that chooses a system of linear commodity taxes in market $\mathcal{M}$. The most general system entails a set of product specific taxes, $\boldsymbol{\tau}_{\mathcal{M}} = \{\tau_j\}_{j \in \mathcal{M}}$. In practice, tax rates tend not to vary across disaggregate products, due to prohibitive implementation issues. We therefore focus on more constrained systems closer to those used in practice. We assume that when introducing the excise tax system the government does not change other parts of the tax system (in particular, the earnings and profit taxes).

Such tax reforms are often motivated by the existence of social costs associated with consumption. We allow for a budgetary externality associated with consumer i's consumption of products in set $\mathcal{S} \subset \mathcal{M}$ that takes the form $\Phi^i \equiv \Phi^i(\boldsymbol{x}_{\mathcal{S}}^i)$, where $\Phi^i(.)$ is weakly increasing in each of its arguments. We denote the marginal externality of consumer i's consumption of good $j \in \mathcal{S}$ by $\frac{d\Phi^i}{dx_j^i} \equiv \phi_j^i$. Below we discuss the implications for tax policy of other forms of social costs, such as internalities.

The optimal choice of excise taxes requires the government to balance reducing inefficiencies associated with consumption externalities with the inefficiencies arising from the exercise of market power. It must also account for any spillovers to existing tax bases, and, depending on its preferences for equity, it may factor in distributional consequences of tax reform. We consider a government with the social welfare function:

$$W = \int_i \omega^i V^i + \lambda(T_D^i + T_\Pi(\delta^i \Pi) - \Phi^i) di, \tag{2.1}$$

where ω^i is the Pareto weight on consumer i, λ is the marginal value of government revenue, and where tax revenue raised from individual i is given by revenue from distortionary taxes:

$$T_D^i = \sum_{j \in \mathcal{M}} \tau_j x_j^i + \mathcal{T}(z^i)$$

[11]We hold fixed firms' entry and product design decisions. In the case of a sugar-sweetened beverage tax we believe this assumption is mild. In what follows we model demand and supply over the major brands that comprise 90% of the UK beverage market. These are long established and all available in both sugar-sweetened and diet varieties, meaning removing all sugar from a sugar-sweetened product to avoid the tax is equivalent to removing the product completely from the market – which our model suggests would lead to significant reductions in variable profits. In Section 6.3 we consider a tax levied directly on sugar content and allow for product redesign.

and from the tax on their profit holdings, $T_\Pi(\delta^i\Pi)$.

Market power has important implications for tax design. First, the existence of positive margins distorts resource allocations. Second, the existence of positive profits, depending on how they are distributed across individuals, may impact the distributional consequences of taxation. Third, as firms re-optimize their strategies in response to a tax change, the tax-exclusive prices of all products in the market may change in response to a change in the tax rate levied on any one product (with one implication of this being tax changes are not necessary shifted one-for-one to the products subject to the tax change).[12]

2.2 Optimal policy

We focus on the case where the government sets a single tax rate, $\tau_{\mathcal{S}}$, applied to the set of sin goods. This serves to clarify the key forces that determine optimal tax policy, facilitates comparison with existing sin tax results derived under perfect competition and is an interesting case as, in practice, governments often implement market specific excise tax systems that set a single rate.

The optimal sin tax rate, $\tau^*_{\mathcal{S}}$, satisfies the *implicit* function:

$$\begin{aligned}\tau^*_{\mathcal{S}} =& \underbrace{\bar{\phi} + \frac{\text{cov}(\phi^i_j, dx^i_j/d\tau_{\mathcal{S}})}{(1/n(\mathcal{S})) \times d\mathbb{X}_{\mathcal{S}}/d\tau_{\mathcal{S}}}}_{\text{externality correction}} - \underbrace{(\bar{\mu}_{\mathcal{S}} - \bar{\mu}_{\mathcal{N}}\Theta_{\mathcal{N}} - \mu_O\Theta_O)}_{\text{market power correction}} \\ &+ \underbrace{\frac{1}{d\mathbb{X}_{\mathcal{S}}/d\tau_{\mathcal{S}}}\left[\text{cov}\left(g^i, \sum_{j\in\mathcal{M}} x^i_j\rho_j - \delta^i(1-\tau^i_\Pi)\frac{d\Pi}{d\tau_{\mathcal{S}}}\right)\right]}_{\text{distributional concerns}} - \underbrace{\frac{d(\int_i \mathcal{T}(z^i)di)/d\tau_{\mathcal{S}}}{d\mathbb{X}_{\mathcal{S}}/d\tau_{\mathcal{S}}}}_{\text{tax base erosion}}.\end{aligned} \tag{2.2}$$

(see Appendix A). $\bar{\phi} \equiv \int_i \frac{1}{n(\mathcal{S})}\sum_{j\in\mathcal{S}}\phi^i_j di$ denotes the average marginal consumption externality in the population and $d\mathbb{X}_{\mathcal{S}}/d\tau_{\mathcal{S}} = \sum_{j\in\mathcal{S}} dX_j/d\tau_{\mathcal{S}}$ is the impact of a marginal tax change on total consumption of the set of sin products. $\bar{\mu}_{\mathcal{X}} \equiv \sum_{j\in\mathcal{X}}\mu_k \frac{dX_j/d\tau_{\mathcal{S}}}{\sum_{j'\in\mathcal{X}} dX_{j'}/d\tau_{\mathcal{S}}}$ is the weighted average margin for products in set $\mathcal{X} = \{\mathcal{S}, \mathcal{N}, O\}$, where the weights are each product's contribution to the marginal impact of the sin tax on equilibrium consumption of all products in that set. $\Theta_{\mathcal{X}} \equiv \frac{d\mathbb{X}_{\mathcal{X}}/d\tau_{\mathcal{S}}}{d\mathbb{X}_{\mathcal{S}}/d\tau_{\mathcal{S}}}$ is the fraction of reduced consumption of products in set $\mathcal{S}$ diverted to those in $\mathcal{X} = \{\mathcal{N}, O\}$ due to a marginal tax rise, and $\rho_j \equiv \frac{dp_j}{d\tau_{\mathcal{S}}}$ is the impact of a marginal tax change on the equilibrium consumer price of product j. $g^i \equiv \frac{\omega^i\alpha^i}{\lambda}$ are social marginal welfare weights. Variation in g^i across individuals reflects the

[12]Related is previous work on optimal taxation and imperfect competition by Auerbach and Hines (2002), who study a setting with a representative consumer, homogeneous products and Cournot competition. Our set-up differs as we allow for heterogeneous consumers, differentiated products and social preferences for equity.

government's preferences for equity: a government with a preference for reducing inequality will assign low weights to the rich and high weights to the poor.

Equation (2.2) expresses the optimal tax rate as a function of wedges from non-government (externality and market power) distortions, tax derivatives (for quantities, prices and profits) and government distributional preferences. It comprises four components, three reflecting efficiency and one equity considerations. Although these components are written additively, it should be noted that the terms on the right hand side of the expression (including product-level margins) depend on the tax rate, and that market power influences each component of the formula, because the strategic responses of firms are key in driving the quantity, price and profit tax derivatives. We describe each component of the formula in turn.

Externality correcting component. This equals the average marginal externality across consumers and sin products, plus an adjustment that reflects the covariance between the consumer-product specific marginal externality and the sensitivity of the individual's consumption of the product to a change in the sin tax rate. This covariance captures how effective the tax is at reducing the most socially costly consumption. The more a tax rise reduces consumption by consumers and/or of products associated with high marginal externalities, the better targeted it will be at the most externality generating consumption and the higher will be the optimal tax rate.[13] This mirrors the logic in the optimal externality correcting tax rate with heterogeneous externalities derived in Diamond (1973). However, an important difference is that the sensitivity of consumption to changes in the sin tax rate also depends on the equilibrium pricing response of firms; i.e., $\frac{dx_j^i}{d\tau_{\mathcal{S}}} = \sum_{j' \in \mathcal{M}} \frac{\partial x_j^i}{\partial p_{j'}} \frac{dp_{j'}}{d\tau_{\mathcal{S}}}$.

Market power correcting component. All else equal, higher equilibrium margins on the products in set $\mathcal{S}$ act to lower the optimal tax rate. This reflects the argument made by Buchanan (1969), who pointed out the appropriate tax rate on an externality-generating monopolist lies below the full Pigouvian rate. However, if taxing products in set $\mathcal{S}$ induces an increase in consumption of other products also supplied non-competitively, then distortions arising from the exercise of market power on these alternatives, all else equal, act to raise the optimal tax rate. The strength of this effect depends on the weighted average equilibrium margin on these

[13]Note, $d\mathbb{X}_{\mathcal{S}}/d\tau_{\mathcal{S}}$ will generally be negative; an increase in the sin tax rate will lower equilibrium consumption of those products. If $\text{cov}(\phi_j^i, dx_j^i/d\tau_{\mathcal{S}})$ is negative, so a tax rise tends to achieve relatively large consumption reductions among products/consumers with large marginal externalities, the externality correction will exceed $\bar{\phi}$. Hence, all else equal, the more the total reduction in sin good consumption is concentrated among the most socially costly consumption, the higher will the externality correction component of the optimal tax rate.

non-taxed alternatives, the extent to which a marginal tax rise diverts equilibrium consumption towards them. For products in the set $\mathcal{N}$, these two forces are captured by the terms $\bar{\mu}_{\mathcal{N}}$ and $\Theta_{\mathcal{N}}$ (analogous expressions capture the influence of any market power distortions from the numeraire good).

Distributional concerns. These enter equation (2.2) through the covariance of consumers' social marginal welfare weights with the reduction in their utility resulting from a marginal increase in the tax rate. The more a marginal tax rise results in relatively large utility falls among those with high social marginal welfare weights (low income consumers when the planner has a preference for equity), the lower will be the optimal tax rate. This covariance term is scaled by the marginal effect of the tax on total consumption of the sin products. Hence, the strength of the distributional channel depends on how responsive equilibrium consumption of the set of sin products is to a marginal sin tax rise: the more sensitive is equilibrium consumption, the smaller the impact of equity considerations on the optimal tax rate.

To highlight how market power distortions interact with distributional concerns, it is informative to consider this term under perfect competition: $\frac{1}{d\mathbb{X}_{\mathcal{S}}/d\tau_{\mathcal{S}}}$ $\text{Cov}\left(g^i, \sum_{j\in\mathcal{S}} x^i_j\right)$. The reduction in a consumer's utility due to a marginal tax rise equals their total consumption on the taxed sin goods. All else equal, the more that those with high social marginal welfare weights consume a large quantity of the set of sin goods, the lower is the optimal tax rate on them.[14] Imperfect competition has two consequences for the utility impact of a marginal tax rate increase, and hence on the distributional component of equation (2.2). First, consumption of the sin products is replaced by consumption of all products in market $\mathcal{M}$, weighted by the marginal impact of tax on each product's equilibrium price. If firms hold their tax-exclusive prices fixed in response to tax changes, then this term collapses back to total consumption of the sin goods. Second, the impact of the tax on the size of individuals' net-of-tax profit holdings also matters. If the tax leads to a reduction in profits, and profit holdings are disproportionately held by those with low social marginal welfare weights (the wealthy), this will act to make the tax more progressive and will increase the optimal rate.

Tax base erosion. This term arises because we assume that the government holds fixed the earnings tax schedule. All else equal, the more that a marginal increase in

[14] Allcott et al. (2019a) show that if the government also optimizes the earnings tax schedule, what matters is the cross-sectional correlation in social marginal welfare weights and consumption of the taxed goods *net of income effects*.

the sin tax leads to a reduction in labor tax revenue, the lower will be the optimal rate. Whether this term (and hence the loss from not re-optimizing the earnings tax alongside introducing the sin tax) is large or small is context dependent. To highlight what drives this term we assume that income effects on labor supply are negligible (see Saez et al. (2012) for empirical support of this), which allows us to re-write $d(\int_i \mathcal{T}(z^i)di)/d\tau_{\mathcal{S}}$ as

$$\frac{d(\int_i \mathcal{T}(z^i)di)}{d\tau_{\mathcal{S}}} = \int_i \frac{\tau_z^i}{1-\tau_z^i}\zeta_z^i \sum_{j\in\mathcal{M}} \xi_j^i x_j^i \rho_j di,$$

where ζ_z^i is the individual (compensated) elasticity of taxable earnings and ξ_j^i is the individual elasticity of demand for product j with respect to earnings (see Appendix A). This expression highlights that a key determinant of the tax base erosion component of equation (2.2) is the strength of income effects for products in market $\mathcal{M}$. In Appendix B we provide empirical evidence that these are very small in the context of the UK drinks market, nonetheless, they may be significant in other settings.

Other forms of social costs

Sin taxes are often partially motivated by internalities – costs that individuals impose on themselves that they ignore when making choices. Internalities affect social welfare through individual utilities, in contrast to a budgetary externality, which affects welfare through its impact on the government's budget. Suppose individual i's consumption of sin products gives rise to an internality, $I^i = I(\mathbf{x}_{\mathcal{S}}^i)$, which impacts their normative but not decision utility, and let $\psi_j^i \equiv \frac{\partial V^i/\partial I^i}{\alpha^i}\frac{\partial I^i}{\partial x_j^i}$ denote the monetary cost the individual imposes on themselves per additional unit of consumption of sin good j due to this bias. In this case, there are two modifications to the optimal policy condition (2.2) – see Appendix A. First, the externality correcting term will now reflect externality and internality correction; ϕ_j^i is replaced with $\phi_j^i + \psi_j^i$. Second, the distributional concerns component contains an additional term capturing the covariance of internality changes and social marginal welfare weights; all else equal, for a government with preferences for equity, the more internality reductions resulting from a tax rise are concentrated among the relatively poor, the more valuable is a given fall in the average internality and hence the higher is the optimal tax rate. The impact of these considerations are studied in detail in Allcott et al. (2019a).

2.3 Discussion of empirical implementation

A common approach to empirical tax analysis is to write the expression of interest in terms of a set of externally valid elasticities, or sufficient statistics. An advantage of this approach is that the elasticities can be estimated using quasi-experimental variation, with transparent identification arguments (Chetty (2009)). However, an important restriction of the application of sufficient statistics to optimal tax formulae (as opposed to marginal tax reforms) is that they require implicitly assuming an iso-elastic preference model – see Kleven (2021).

There are two challenges with implementing a sufficient statistics approach in our context. First, the tax derivatives do not straightforwardly map into price elasticities. For a number of the derivatives this could potentially be overcome by making simplifying assumptions about tax pass-through.[15] Alternatively, one could use data covering the introduction of a new (or change in an existing) tax to directly estimate the tax derivatives (policy elasticities in the language of Hendren (2016)). The second challenge is that the optimal tax formula depends on equilibrium product-level price-cost margins. Marginal costs, and hence margins, are typically not straightforwardly observable in economic data. However, they can be inferred based on a profit-maximizing model of the firms operating in the market, coupled with estimates of the own- and cross-price elasticities of all the products in the market.

Our approach is therefore to specify and estimate an equilibrium model of the market of interest (the UK drinks market). This enables us to simulate the impact of an arbitrary tax policy on equilibrium margins, consumption and profits. To validate the model, we compare, where possible, its predictions to existing estimates of relevant elasticities, as well as to evidence from the introduction of the UK's sugary drinks tax. We embed the equilibrium model into the tax problem and consider optimal policy under different government preferences for redistribution.

3 The drinks market

Sugar-sweetened beverage taxation is a natural setting in which to study how taxes designed to reduce social costs interact with market power. In many jurisdictions, taxes on drinks are explicitly motivated as a tool for improving public health, in

[15] For instance, in general, the impact of a marginal tax rise on consumption of the set of sin products takes the form $\frac{d\mathbb{X}_\mathcal{S}}{d\tau_\mathcal{S}} = \sum_{j\in\mathcal{S}}\sum_{j'\in\mathcal{M}}\frac{\partial X_j}{\partial p_{j'}}\frac{dp_{j'}}{d\tau_\mathcal{S}}$. However, under the assumption of fixed tax pass-through (denoted ρ) across products in $\mathcal{S}$ and fixed tax-exclusive prices for other goods, this collapses to $\rho \times \bigtriangledown\mathbb{X}_\mathcal{S}$ where we use $\bigtriangledown\mathbb{X}_\mathcal{S}$ to denote the marginal impact of consumption goods in $\mathcal{S}$ with respect to a marginal price rise for all these products.

part due to substantial social costs associated with their consumption. We discuss the nature and measurement of these costs in Section 6. It is also the case that the market is concentrated and comprises a set of highly recognizable branded products. It is likely therefore that firms exercise considerable market power.

3.1 Data

We model behavior in the UK market for non-alcoholic drinks. This market includes carbonated drinks (often referred to as sodas), fruit concentrates, and sports and energy drinks. We refer collectively to these as soft drinks. Soft drinks brands usually come in sugar-sweetened and artificially sweetened (i.e., diet) varieties. Sugar-sweetened beverage taxes typically apply to sugar-sweetened varieties of soft drinks. The market also includes pure fruit juices and flavored milk. We use micro data on the drinks purchases of a sample of consumers living in Great Britain collected by the market research firm Kantar. The data contain information on household level purchases for "at-home" consumption (Kantar Worldpanel), as well as purchases made by individuals for "on-the-go" consumption (Kantar On-The-Go Survey). Together "at-home" and "on-the-go" consumption account for over 90% of drinks consumption by volume.[16]

There are 30,405 households in the at-home sample who record, by means of a barcode scanner, all grocery purchases made and brought into the home. The data are broadly representative of British households (in Appendix B we compare them with the nationally representative Living Cost and Food Survey) and cover 2008 to 2012. Individuals in the on-the-go data record all purchases they make from shops and vending machines for out-of-home consumption using a cell phone app. The data cover 2010 to 2012 and comprise 2,862 individuals (aged 13 and upwards) randomly drawn from the Worldpanel households. In both datasets, we observe households/individuals over many months. The data contain detailed information – including brand, flavor, size and nutrient composition – on the UPCs (barcodes) purchased, the store in which the purchase took place, and transaction level prices.

3.2 Consumers and purchasing patterns

We use the term consumer to refer to households in the at-home segment, and individuals in the on-the-go segment. Figure 3.1 highlights variation in soft drinks purchases across two key dimensions. Panel (a) shows that consumers that get a

[16]The remainder occurs in restaurants and bars, which are not covered by our data. Numbers are computed using the Living Cost and Food Survey.

high fraction of their *total* dietary calories from added sugar purchase significantly more sugar-sweetened beverages than other consumers. Policymakers have typically focused on changing the behavior of consumers with dietary sugar above a particular threshold, due to elevated health risks (e.g., the World Health Organization (2015) advice focuses primarily on those with dietary sugar above 10%). The more that a sugar-sweetened beverage tax is able to achieve large consumption reductions among those consumers that create relatively high marginal externalities through their sugar-sweetened beverage intake, the more effective it will be at reducing externality distortions.

Panel (b) shows that there is a negative cross-sectional correlation between sugar-sweetened beverage consumption and equivalized household income – richer households consume less sugar-sweetened beverages, and therefore have consumption baskets less exposed to a sugar-sweetened beverage tax, than lower income consumers. The extent to which this is driven by preference heterogeneity (correlated with income) versus causal income effects determines how much sugar-sweetened beverage taxation induces labor supply distortions (captured by the tax base erosion term in equation (2.2)), and hence impacts optimal policy. In Appendix B we show that after removing consumer fixed effects (and hence relying on within household income transitions to estimate any income effects), the consumption gradient in equivalized income flattens completely. In our demand model we control flexibly for equivalized income, and in our optimal tax analysis, we treat variation in drinks demand across the income distribution as preference heterogeneity.

Figure 3.1: *Variation in volume of sugar-sweetened beverages consumed at-home*

(a) Dietary sugar intake　　　　(b) Household income

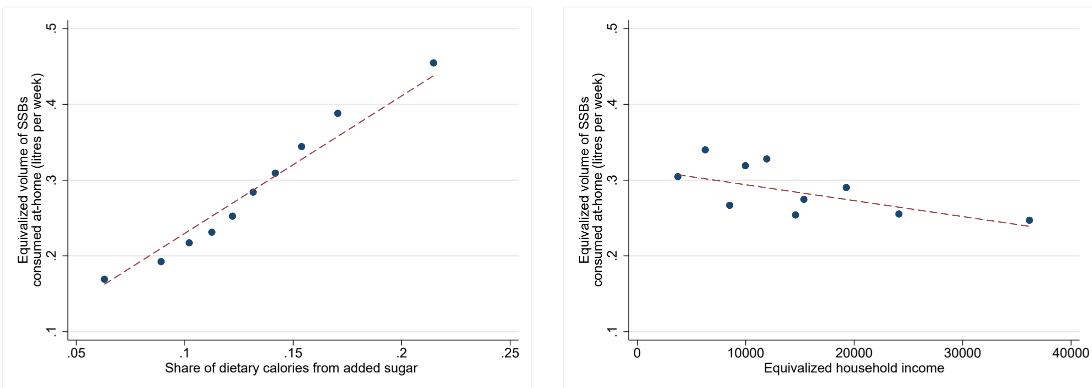

Notes: The left hand panel shows mean volume of sugar-sweetened beverages purchased per person per week and consumed at home by deciles of the share of dietary calories from added sugar (from food consumed at home). The right hand panel shows mean volume of sugar-sweetened beverages purchased per person per week and consumed at home by deciles of equivalized (using the OECD-equivalence scale) household income. Analogous figures for the sugar from soft drinks consumed on-the-go are shown in Appendix B.

3.3 Firm and brands

The drinks market is highly concentrated. In Figure 3.2 we show the cumulative market share of the ten largest UK producers. The two largest firms (Coca Cola Enterprises and Britvic have a combined market share of almost 60%).[17]

Firms each typically own several separate brands. For instance, Coca Cola Enterprises most popular brand is Coca Cola, but it also owns 9 other brands with market share of at least 1%. Each soft drinks brand is typically available in sugar-sweetened ("regular") and artificially sweetened ("diet" and/or "zero") variants. Each brand-variant is available in multiple pack-sizes. In our equilibrium model of the market we focus on the set of brands with more than a 1% market share in either the at-home or on-the-go segment, as well as the main fruit juice and flavored milk brands, which together comprise over 75% of total spending on non-alcoholic drinks. Table B.3 in Appendix B lists brand-variants and the number of sizes they are available in.

Figure 3.2: *Market share of the largest firms in the drinks market*

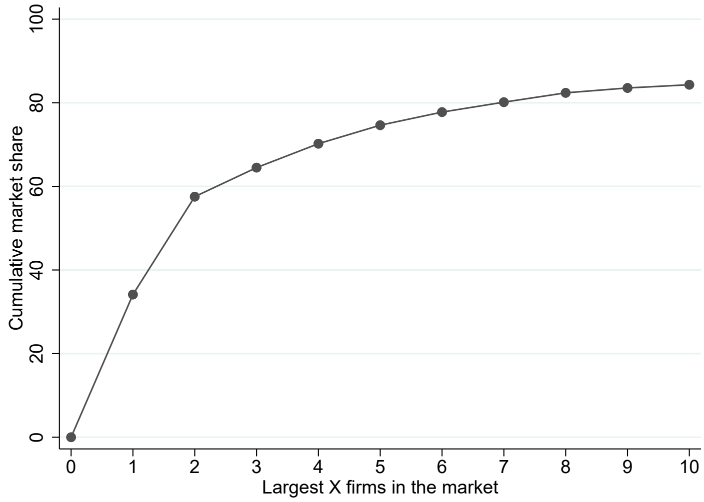

Notes: The line shows the cumulative market share for the X largest firms in the market, where X is shown on the horizontal axes. Market shares are shown for 2012.

[17] Drinks producers are known in the industry as bottlers. They buy concentrate from upstream firms (e.g., Coca Cola Enterprises obtains concentrate from The Coca Cola Company) and use this as an input to produce soft drinks products (see Luco and Marshall (2020)). For modeling firm behaviour in the UK drinks market, it is the bottlers – all of which are national – who are the relevant agents: as well as producing the products, they are responsible for negotiating product shelf prices and placement with retailers, and for promotional activity (see Competition Commission (2013)).

3.4 Drinks firm-retailer relations

Drinks firms do not sell directly to consumers, rather retailers act as intermediaries between drinks firms and consumers. The majority of expenditure is undertaken in national grocery chains (see Table B.4 in Appendix B). In the UK the main grocery chains set prices nationally (see Competition Commission (2000)), which means that if we observe a product's price in one store, we know the price faced by consumers in other stores belonging to the same retailer chain. Retailers typically offer all brand-varieties, though the available pack sizes can vary.

We do not directly observe the contracting relationship between the drinks firms and retailers. However, a 2013 report into the soft drinks market by the UK competition authority provides evidence on the nature of these relations (see Competition Commission (2013)). They cite evidence that annual bilateral "Joint Business Plans" are agreed between a drinks firm and retailer setting out wholesale prices, payments related to product visibility, recommended retail prices, and agreements on the number, type and timings of promotions. This evidence of non-linear contracting suggests drinks firms and retailers avoid double marginalization. We therefore treat drinks firms as (effectively) setting final consumer prices, an outcome consistent with optimal non-linear contracting – see Villas-Boas (2007) and Bonnet and Dubois (2010). We also exploit the fact that the promotions are agreed on in advance (and are not coordinated across retailers) as a useful identifying source of price variation (see Section 4.3).

4 Equilibrium model of the drinks market

We estimate a model of consumer demand in the drinks market using a discrete choice framework in which consumer preferences are defined over product characteristics (Gorman (1980), Lancaster (1971), Berry et al. (1995)). This enables us to model demand for the many differentiated products in the market, while incorporating rich preference heterogeneity, including by total dietary sugar and income. We identify product level marginal costs by combining the demand estimates with the equilibrium conditions from an oligopoly pricing game (Berry (1994), Nevo (2001)). The estimates of the primitives of demand and supply enable us to simulate the impact of tax policy on equilibrium quantities and prices, and hence consumer utilities and profits. Appendix C provides additional details on the demand specification.

4.1 Consumer demand

Choice problem

We model which, if any, drink product a consumer (indexed i) chooses on a "choice occasion", where choice occasion refers to a week in which a household purchases groceries in the at-home segment, or a day on which an individual buys a cold beverage (including bottled water) in the on-the-go segment. We treat the decisions that households make in the at-home segment and individuals make in the on-the-go segment separately, allowing for all preferences to vary freely with each type of choice situation. In Appendix C we provide evidence that recent purchases of drinks by a household in the at-home segment do not influence either the propensity to buy or quantity purchased by household members in the on-the-go segment. Choice in the on-the-go segment is between single portion size of products (e.g., 330ml cans and 500ml bottles); choice in the at-home segment is between multi-portion sizes. For notational parsimony we suppress a market segment index.

We index the drinks products by $j = \{1, \ldots, J\}$. Products vary by brand, indexed by $b = \{1, \ldots, B\}$, whether or not they contain sugar (for instance, the brand Coke is available in Regular, Diet and Zero variants), and their size, indexed by $s = \{1, \ldots, S\}$. Brand-variants can be purchased in different sizes for two reasons: (i) the availability of different pack sizes (or UPCs), and (ii) the purchase of multiple packs. For each brand-variant we define sizes as the set of available pack sizes and the most common multiple pack purchases of UPCs. The consumer chooses between the available drinks products and choosing not to buy a drink, which we denote by $j = 0$. On around 42% of at-home and 60% of on-the-go choice occasions, a household purchases a drink (i.e., $j > 0$).[18] We take as given the consumer's decision over which retailer (indexed r) to shop with. Choice sets are retailer specific (and denoted Ω_r) due to some variation in available pack-sizes by retailer.

Consumer i in period t, with total period budget y_{it}, solves the utility maximization problem:

$$V(y_{it}, \mathbf{p}_{rt}, \mathbf{x}_t, \epsilon_{it}; \boldsymbol{\upsilon}_i) = \max_{j \in \{\Omega_r \cup 0\}} \nu(y_{it} - p_{jrt}, \mathbf{x}_{jt}; \boldsymbol{\upsilon}_i) + \epsilon_{ijt}. \tag{4.1}$$

[18] Consumers are sometimes observed purchasing multiple (typically) two brand-variants on a single choice occasion. On 40% (10%) of occasions in which a consumer chooses a drink in the at-home (on-the-go) segment, multiple are chosen. In this case, we randomly sample one, assuming that, conditional on consumer specific preferences, these purchases are independent, e.g., because they are bought for different household members.

where $\mathbf{p}_{rt} = (\mathbf{p}_{1rt}, \ldots, \mathbf{p}_{Jrt})$ is the price vector faced by the consumer, $\mathbf{x}_{jt}$ are other characteristics of product j, and $\mathbf{x}_t = (\mathbf{x}_{1t}, \ldots, \mathbf{x}_{Jt})$ (note $p_0 = 0$ and $\mathbf{x}_{0t} = 0$); $\boldsymbol{\upsilon}_i$ is a vector of consumer level preference parameters; and $\epsilon_{it} = (\epsilon_{i0t}, \epsilon_{i1t}, \ldots, \epsilon_{iJt})$ is a vector of idiosyncratic shocks.

The function $\nu(.)$ captures the payoff the consumer gets from selecting option j. Its first argument, $y_{it} - p_{jrt}$, is spending on the numeraire good, i.e., spending outside the drinks market. We assume that preferences are quasi-linear, so $y_{it} - p_{jrt}$ enters $\nu(.)$ linearly and y_{it} differences out when the consumer compares different options; we therefore suppress the dependency of $\nu(.)$ on y_{it}. An implication of quasi-linearity is that a change in the price of any drinks product does not induce an income effect. Given the small share of total consumer expenditure allocated to drinks products, this is a mild assumption. However, we do allow equivalized household income to shift consumer preferences, υ_i. This enables our model to capture how demands vary across consumers with different incomes.

We assume that ϵ_{ijt} is distributed i.i.d. type I extreme value. Under this assumption the probability that consumer i selects product j in period t, conditional on prices, product characteristics and preferences, is given by:

$$\sigma_j(\mathbf{p}_{rt}, \mathbf{x}_t; \boldsymbol{\upsilon}_i) = \frac{\exp(\nu(p_{jrt}, \mathbf{x}_{jt}; \boldsymbol{\upsilon}_i))}{1 + \sum_{j' \in \Omega_r} \exp(\nu(p_{j'rt}, \mathbf{x}_{j't}; \boldsymbol{\upsilon}_i))}, \tag{4.2}$$

and the consumer's expected utility is given by:

$$v(\mathbf{p}_{rt}, \mathbf{x}_t; \boldsymbol{\upsilon}_i) = \ln \sum_{j \in \Omega_r} \exp\{\nu(p_{jrt}, \mathbf{x}_{jt}; \boldsymbol{\upsilon}_i)\} + C, \tag{4.3}$$

where C is a constant of integration.

Utility specification

We allow for rich preference heterogeneity, with both observed and unobserved consumer characteristics. This is important for two reasons. First, it is well established that the inclusion of rich, and, in particular, unobserved preference heterogeneity is crucial in enabling models of this type to recover realistic patterns of consumer substitution across products (see Berry et al. (1995), Nevo (2001)). Second, in our application it is important that we capture variation in preferences across different consumers that is relevant for tax policy. This includes variation across consumers whose consumption is likely to create different externalities at the margin, and across consumers to which a government may assign different social marginal welfare weights.

We therefore partition consumers into groups (indexed d) based on two variables. First, whether their total share of dietary calories (measured in the preceding year) is below 10%, between 10 and 15% and in excess of 15%. Second, we allow for at-home preferences to vary by whether the household contains children and on-the-go preferences to vary by whether the individual is younger or older than 30. We allow all preference parameters to vary by these groups; see Table C.1 in Appendix C for details. We also allow preferences over key product attributes, including price, to vary with equivalized household income (which we denote $\tilde{y}_i$).

The payoff function $\nu(.)$ for consumer i belonging to consumer group $d(i)$ and for product j belonging to brand $b(j)$ and of size $s(j)$ takes the form:

$$\nu(.) = -\alpha_{i0} p_{jrt} + \sum_{k>0}^{K} \alpha_{ik} x_{jk} + \zeta_{d(i)b(j)s(j)rt},$$

where, for product attribute, $k = 0, \ldots, K$:

$$\alpha_{ik} = \bar{\alpha}_{d(i)k} + \alpha^{\tilde{y}}_{d(i)k} \tilde{y}_i + \sigma^{\alpha}_{d(i)k} \eta_{ik},$$

$\zeta_{d(i)b(j)s(j)rt}$ denotes an unobserved brand-size attribute (which may vary by consumer group, retailer and time) and η_{ik} is a standard normal random variable.[19]

$\bar{\alpha}_{d(i)k}$ denotes the baseline preference for product attribute k among consumer group d, $\alpha^{\tilde{y}}_{d(i)k}$ captures how preferences for the attribute vary across consumers with different equivalized household incomes, and $\sigma^{\alpha}_{d(i)k}$ captures the dispersion in unobserved preferences for the attribute. In addition to price, product attributes include sugar content, drink types (cola, lemonade, pure fruit juice etc.) size and advertising.[20] We allow preferences over price, sugar, branded soft drinks, non-branded drinks and pure fruit juice to vary with equivalized income, and for unobserved heterogeneity in preferences for drinks (relative to the outside option), price, sugar, size, cola, lemonade, non-branded drinks and fruit juice. We allow for correlation between preferences for sugar and for soft drinks (relative to the outside option). We allow both the baseline and dispersion parameters to vary across consumer groups d, so the overall random coefficient distribution is a flexible mixture of normal distributions.

[19]For price we assume the coefficient is log-normally distributed.

[20]We measure monthly advertising expenditure across TV, radio and online in the AC Nielsen Advertising Digest. We compute product specific stocks based on a monthly depreciation rate of 0.8. This is similar to the rate used in Dubois et al. (2018) on similar data in the potato chips market.

We decompose the unobserved product attribute, $\zeta_{d(i)b(j)s(j)rt}$, into a set of detailed fixed effects:

$$\zeta_{d(i)b(j)s(j)rt} = \xi^{(1)}_{d(i)b(j)s(j)} + \xi^{(2)}_{d(i)b(j)t} + \xi^{(3)}_{d(i)s(j)t} + \xi^{(4)}_{d(i)s(j)r} + \xi^{(5)}_{d(i)b(j)r}$$

All the fixed effects are consumer group specific. They include brand-size, brand-time (year-quarter), size-time (year-quarter), brand-retailer and size-retailer effects. These control for shocks to demand that may be correlated with price setting.

4.2 Supply model

We model drinks firms as setting prices in a simultaneous move Nash-Bertrand game.[21] We do not explicitly model drinks firms-retailer relationships, but, based on the evidence of non-linear contracting in these relations, we assume manufacturers set consumer prices, which is consistent with efficient contracting. Let $\boldsymbol{p}_m = (p_{1m}, \dots, p_{Jm})$ denote the prices that drinks firms set in market (year) m.[22]

Market demand for product j is given by:

$$q_{jm}(\mathbf{p}_m) = \int_i \sigma_j(\mathbf{p}_m, \mathbf{x}_m; \boldsymbol{v}_i) dF(\boldsymbol{v}) M_m,$$

where M_m denotes the potential size of the market. We denote the marginal cost of product j in market m as c_{jm}.

We index the drinks firms by $f = (1, \dots, F)$ and denote the set of products owned by firm f by $\mathcal{J}_f$. Firm f's total variable profits in market m are

$$\Pi_{fm}(\boldsymbol{p}_m) = \sum_{j \in \mathcal{J}_f} (p_{jm} - c_{jm}) q_{jm}(\boldsymbol{p}_m). \tag{4.4}$$

Under Nash-Bertrand competition, the equilibrium prices satisfy the set of first order conditions: $\forall f$ and $\forall j \in \mathcal{J}_f$,

$$q_{jm}(\boldsymbol{p}_m) + \sum_{j' \in \mathcal{J}_f} (p_{j'm} - c_{j'm}) \frac{\partial q_{j'm}(\boldsymbol{p}_m)}{\partial p_{jm}} = 0. \tag{4.5}$$

[21] Around one-fifth of the market consists of "store brands". These are no-frills, low priced alternatives to the branded products. We treat these as a competitive fringe, with (tax-exclusive) prices that remain fixed.

[22] In practice, for a given product-year a drinks firm and retailer agree on a base price $\bar{p}$ and a sale price p_S, with the former applying ρ proportion of weeks (see Section 3.4). Instead of modeling choice over $(\bar{p}, p_S, \rho)$, we model choice over $p = (1-\rho)\bar{p} + \rho p_S$. This average price exhibits little variation across retailers. Cross-retailer variation in the price of a given product at a point in time is driven by non-synchronization of sales (see next section). Hence, we specify the relationship between prices in the supply game, p_{jm}, and those faced by consumers in retailer r week $t \in m$ as $p_{jrt} = p_{jm} + e_{jrt}$, where $\mathbb{E}[e_{jrt} | (j, m)] = 0$.

From this system of equations we can solve for the implied marginal cost, c_{jm}, and hence the equilibrium price-cost margin, $\mu_{jm} = p_{jm} - c_{jm}$, for each product in each market. For any set of product taxes we can use the system of equations (4.5), replacing c_{jm} with $c_{jm} + \tau_j$, to solve for counterfactual equilibrium prices p'_{jm} and margins $\mu'_{jm} = p'_{jm} - \tau_j - c_{jm}$.

4.3 Identification

We begin by discussing identification of the baseline price parameters ($\bar{\alpha}_{d(i)0}$), before discussing the other preference parameters.

We assume that the detailed fixed effects that we include in the model (in addition to the advertising controls) absorb taste variation that is relevant for price-setting. The brand-size effects, $\xi^{(1)}_{d(i)b(j)s(j)}$, absorb the influence of unobserved product attributes that are not captured by the included observable product attributes. Variation in taste for brands or particular sizes over time, due, for instance, to seasonal patterns, are captured by the brand-time, $\xi^{(2)}_{d(i)b(j)t}$, and size-time, $\xi^{(3)}_{d(i)s(j)t}$, effects. In addition, we control separately for product level advertising, which will capture the effect on demand of the (overwhelmingly national) advertising in the UK drinks market. Finally, tastes for brands or sizes may vary across retailers, which is captured by the brand-retailer, $\xi^{(4)}_{d(i)b(j)r}$, and size-retailer, $\xi^{(5)}_{d(i)s(j)r}$, effects.

The price variation that we exploit is *product-level time-series variation that is differential across retailers*. In particular, while differences in the average price (over time) that different retailers set for a given product are small, the degree of co-movement in prices for the same product in different retail chains over time is low. The differential price movements are generally driven by time-limited price reductions, which vary in timing and depth across retailers; see Appendix C. These price reduction strategies, which are agreed in advance at bilateral meetings between drinks firms and retailers (see Section 3.4), creates, from consumers' perspective, randomness in the prices they face.

A threat to our identification strategy is that consumers respond to promotions by intertemporally switching their purchases, i.e., stocking up when the price is low and consuming from this stock when the price is high. This would lead to over-estimates of own price elasticities and likely to under-estimates of cross price elasticities (Hendel and Nevo (2006a)). In Appendix C we show that, when purchasing on sale, consumers are more likely to choose a different brand, container type and size relative to their previous purchase, but they do not systematically change the timing of their purchases. In other words, sales in the UK drinks market primarily lead to *intra*-temporal substitution rather than intertemporally switching.

We do not model the decision that consumers make over which retailer to shop with, though we do allow brand and pack-size preferences to vary flexibly by retailer. We therefore assume that the choice of retailer is not driven by time-varying, drink product specific factors. Drinks expenditure comprises only a small share of total grocery expenditure (4%), so we think this is a reasonable assumption. In addition, we find that when purchasing on sale, there is no economically meaningful increase in the likelihood that the consumers shopped with a different retailer compared with their previous purchase, which supports our assumption (see Appendix C).

For the non-price product attributes the baseline preference is identified as long as the attribute exhibits within brand-size variation (otherwise it is absorbed by the brand-size fixed effects). For instance, a given consumer group will have a stronger preference for sugar if, conditional on brand-size, they more regularly choose sugar rather than diet products.

We observe in our micro data many consumers of different total dietary calories, household composition and income making repeated choice while facing different price vectors. We use this variation to incorporate rich preference heterogeneity into the model across these observable consumer attributes. We also include rich unobserved preference heterogeneity. The panel structure of our data is helpful in pinning down the spread parameters governing the unobserved preference distributions. For instance, all else equal, a higher degree of within-consumer covariation in the price of chosen options across choice occasions, the higher the dispersion in price preferences, which is captured by the spread parameter, $\sigma^{\alpha}_{d(i)0}$.

5 Demand and supply estimates

5.1 Consumer substitution patterns

We estimate the demand model outlined in Section 4.1 using simulated maximum likelihood, and report the coefficient estimates in Appendix D. The estimated coefficients exhibit some intuitive patterns: those with relatively high overall added sugar in their diets have stronger preferences for sugary drinks products, and those with lower incomes are more sensitive to price, have stronger preferences for soft drinks and weaker preferences for pure fruit juice. The variance parameters of the random coefficients are significant both statistically and in size, indicating an important role for unobserved preference heterogeneity. We allow for correlation in preferences (within consumer group) between tastes for drinks (relative to the outside option) and preferences for sugar, and find a negative relationship.

The estimated preference parameters determine our demand model predictions of how consumers switch across products as prices change. The mean own-price elasticity is around -2.1 (in both the at-home and on-the-go segments), though with significant variation around this: 25% of products have own-price elasticities with magnitude greater than 2.5, a further 25% of products have own-price elasticities with magnitude less than 1.6. The distribution of the cross-price elasticities exhibits a high degree of skewness, with the mean close to the 75^{th} percentile. This reflects consumers' willingness to switch between products close together in product characteristic space.

Table 5.1: *Selected elasticities for cola products*

	Coke						Pepsi				Non-colas		
	Regular		Diet		Zero		Regular		Max		SSBs	Diet	Fruit juice
	2l b.	8 pk.	2l b.	8 pk.	2l b.	8 pk.	2l b.	8 pk.	2l b.	8 pk.			
Regular													
2l bottle	-2.204	0.018	0.011	0.009	0.011	0.009	0.023	0.017	0.012	0.009	0.007	0.003	0.005
8x330ml can	0.036	-2.832	0.017	0.022	0.017	0.021	0.035	0.042	0.017	0.021	0.013	0.006	0.008
Diet													
2l bottle	0.010	0.007	-2.185	0.014	0.019	0.014	0.010	0.007	0.020	0.014	0.003	0.006	0.003
8x330ml can	0.014	0.018	0.026	-2.777	0.027	0.032	0.014	0.017	0.026	0.031	0.006	0.011	0.005

Notes: Numbers show price elasticities of market demand (for products listed in top row) in the most recent year covered by our data (2012) with respect to price changes for two specific pack sizes of Coke Regular and Diet Coke (shown in first column). "Non-colas" exclude Coke and Pepsi and are means over products belonging to each of the sets, sugar-sweetened beverages (SSBs), diet drinks and fruit juices.

To illustrate this, Table 5.1 shows product level elasticities associated with a price change for two popular sizes – a 2l bottle and a 8 pack of 330ml cans – of Coke Regular and Diet Coke. It shows the impact on demand for each of the 2l bottle and 8×330ml packs of Coke and Pepsi, and the mean elasticities for other (non-cola) sugar-sweetened and diet beverages, and for pure fruit juice. The table highlights a number of intuitive patterns: (i) consumers are more willing to switch across cola products of the same variety (e.g., within Regular) than they are to other varieties (e.g., Diet or Zero) or to non-cola drinks; (ii) consumers are more willing to switch between products of the same size/pack type than they are to different sizes; (iii) consumer substitution from sugary varieties of Coke to sugary non-cola drinks (both sugar-sweetened beverages and fruit juice) is stronger than it is from Diet Coke. In Appendix D we report further details of product level elasticities.

In Table 5.2 we summarize the effects of increasing the price of all sugar-sweetened beverages by 1%. This leads to a 1.41% fall in liters demanded of sugar-sweetened beverages. Around 33% of the reduction in demand for sugar-sweetened beverages is diverted to alternative drinks. As we discuss in Section 2, if alternative products are supplied non-competitively the degree of switching to them in response to a marginal tax rise is an important determinant of the optimal policy. The diversion ratio in Table 5.2 does not directly tell us this since (i) it reflects only demand responses to a price change, but not supply side pricing responses to a tax change and (ii) it is evaluated at observed prices and not at the optimal tax rate. Nonetheless, as we show in Section 6, the relatively high degree of substitution between the two product sets indicated by the diversion ratio plays an important role in determining optimal sugar-sweetened beverage taxation. The 1% increase in the price of sugar-sweetened beverages (at observed prices) leads to essentially no change in overall drinks expenditure.

Table 5.2: *Switching due to an increase in the price of sugar-sweetened beverages*

Own price elasticity for sugar sweetened beverages	-1.41
	[-1.46, -1.37]
% lost demand diverted to substitute drinks	32.8
	[32.3, 33.6]
% change in overall drinks expenditure	0.047
	[0.033, 0.060]

Notes: We simulate the effect of a 1% price increase for all sugar-sweetened beverage products. The first row shows the % reduction in volume demanded of sugar-sweetened beverages, the second row shows how much of the volume reduction is diverted to substitute drinks products, and the third row shows the percent change in total drinks expenditure. Numbers are for the most recent year covered by our data (2012). 95% confidence intervals are given in square brackets.

5.2 Estimated costs and margins

We use the first order conditions of the firms' profit-maximization problem (equation (4.5)) to solve for product marginal costs, and hence the price-cost margins and Lerner indexes (margin over price) at observed prices. In Table 5.3 we show the averages of these for sugar-sweetened beverages and alternative products.

The average Lerner index is 0.57 for sugar-sweetened beverages and 0.55 for alternative products. This indicates that firms exercise a significant degree of market power when setting the prices both of sugar-sweetened beverages and alternative drinks. As we illustrate in Section 6, failing to account for distortions from the exercise of this market power leads to substantial unrealized welfare gains when setting tax policy. In Appendix D we show that there is substantial variation in equilibrium margins across brand and that average margins are declining in product size (since, on average, price per liter is strongly declining in size, while marginal cost per liter is flatter across the size distribution).

It is important to emphasize that distortions from the exercise of market power are endogenous to tax policy. For instance, a tax levied on sugar-sweetened beverages may exacerbate distortions from market power on these products if firms respond by raising their tax-exclusive prices (and hence product level margins) and/or consumers respond by downsizing to smaller high margin products.

Table 5.3: *Summary of costs and margins*

	Sugar-sweetened beverages	Alternative products
Price (£/l)	1.09	1.07
Marginal cost (£/l)	0.42	0.44
Price-cost margin (£/l)	0.67	0.63
Lerner index (margin/price)	0.57	0.55

Notes: We recover marginal costs for each product in each market. The table shows the average price, marginal cost, price-cost margin (all expressed in per liter terms) and Lerner index among sugar-sweetened beverages and substitute products, constructed using quantity weights. We report the values for the most recent year covered by our data (2012).

5.3 Discussion of demand and supply estimates

Our demand and supply estimates enable us to capture rich consumer level substitution patterns across products, as well as product level price-cost margins. These are key inputs into our empirical study of optimal taxation with market power. However, this necessarily entails making identifying and functional form assumptions about the nature of demand and firm competition. Here we compare our estimates with those in the existing literature, and to alternative information on price-cost margins. We also provide validation of our model using data on price changes following the introduction of the UK's soft drinks tax.

We estimate an own-price elasticity for sugar-sweetened beverages of 1.41. We calculate this by simulating an increase in the prices of all sugar-sweetened beverages by 1% and recovering the change in demand for those products, allowing for substitution between sugary beverages, to alternative drinks and to not buying drinks. Allcott et al. (2019a) employ an alternative approach, using US scanner data and an instrumental variable methodology applied to quarterly purchases of sugar-sweetened beverages. They estimate an own-price elasticity for sugar-sweetened drinks in line with ours (between -1.37 and -1.48, depending on the specification). The margins we recover imply the average Lerner index in the market is 0.56, which is broadly consistent with the gross margins reported in accounting data, which are between 35-70% (see Competition Commission (2013)).

Pass-through

We use our demand and supply estimates to simulate how firms choose to adjust their margins when a tax is introduced. A potential concern with simulated pass-through of a hypothetical tax is that it can be influenced by functional form assumptions. In particular, the curvature of market level demand is an important determinant of tax pass-through (see Weyl and Fabinger (2013)). A feature of logit demand models with no heterogeneity in preference parameters is that they heavily restrict demand curvature. However, the addition of preference heterogeneity breaks the link between the curvature of individual- and market-level demand curves, allowing for more flexibility in the latter, as curvature now also depends on how the composition of individuals along the market demand curve changes (Griffith et al. (2018)).

We provide direct evidence that our model succeeds in generating realistic tax pass-through predictions by simulating the introduction of the UK's Soft Drinks Industry Levy (SDIL) in 2018 and comparing this with what happened in practice – see Appendix D.2 for full details. This tax was introduced more recently than the period covered by our data, so we use a weekly database of UPC level prices, collected from the websites of six major UK supermarkets, that cover the period 12 weeks before and 18 weeks after the introduction of the tax. We estimate the average within-product price changes before and after the introduction of the tax. Figure 5.1 shows the results for those products with sugar content above 8g per 100ml, which were subject to the higher rate (results for the low tax band and tax-exempt products are shown in the Appendix). We find evidence that the tax was slightly overshifted, with average pass-through rates of 105-108% and no change in the price of untaxed products. The price changes predicted by our model are very close to the observed price changes.

To illustrate the extent to which our pass-through predictions are driven by the logit shocks in our demand model we also simulate pass-through shutting down all preference heterogeneity in the model. Doing this results in simulated tax pass-through that is around 50% on average. This illustrates that modeling preference heterogeneity is key in enabling us to capture realistic pass-through patterns.

Modest over-shifting of a drinks tax is broadly consistent with evidence from ex-post evaluations of these policies. For example, the Philadelphian tax was found to be fully passed through to prices (Seiler et al. (2021), Cawley et al. (2020)), and in Mexico the tax was fully to slightly more than fully passed through to prices (Grogger (2017), Colchero et al. (2015)). An exception is Berkeley, where pass-through of the tax is estimated to be statistically insignificant or low (e.g. Rojas and Wang (2017), Bollinger and Sexton (2018)).

Figure 5.1: *Out-of-sample model validation: UK Soft Drinks Industry Levy*

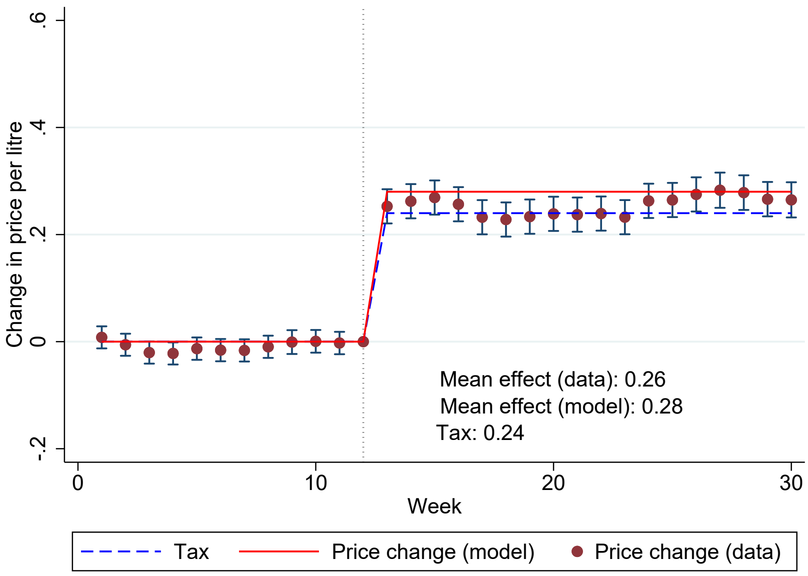

Notes: The red circles show the estimated price changes (relative to the week preceding the introduction of the tax). Appendix D.2 contains full details of the specification. 95% confidence intervals shown. The blue line shows the value of the tax, and the red solid line shows the predicted price changes from our estimated demand and supply model.

6 Optimal tax results

In this section we combine our empirical model of the UK drinks market with the tax design framework that we outline in Section 2 to quantitatively solve for optimal tax policy; see Appendix E for the solution algorithm. We begin by considering the efficiency-maximizing sin tax rate, where the government seeks to minimize allocative distortions, but is indifferent to the incidence of the tax across individuals. Next we consider the optimal sin tax rate when the government has distributional concerns, highlighting the importance of inequality in consumption and profit holdings across individuals in determining optimal policy. Finally, we consider the potential welfare gains from multi-rate taxation and levying tax directly on product sugar content rather than volumetrically.

External costs of sugar-sweetened beverage consumption. The welfare effect of tax policy depends on the magnitude of externalities from sugar-sweetened beverages. Consumption of these products can increase health care costs, the bulk of which are not borne by the individual (see Scientific Advisory Committee on Nutrition (2015)). As in Allcott et al. (2019a), we measure the monetary value of the average externality using epidemiological evidence from Wang et al. (2012). They estimate the impact of a reduction in sugar-sweetened beverage consumption on US healthcare costs;[23] we translate this into a UK-equivalent figure. Provision of healthcare in the UK is almost entirely public, meaning that increased healthcare costs directly impact the government's budget. We also use evidence from the UK based on an epidemiological study (Briggs et al. (2013)) and information on the NHS costs of treating obesity. Both approaches provide a similar picture and lead to an estimate of the average externality of approximately 4 pence per 10g of sugar – see Appendix E for details of these calculations. Based on a review of the medical consequences of added sugar intake, the World Health Organization (2003) recommends that added sugar should make up less than 10% of dietary calories. We therefore assume that externalities from sugar-sweetened beverage consumption are due to those people with more than 10% of their dietary calories from sugar; this implies the externality per 10g of sugar consumption is 5 pence for the 80% of consumers with dietary sugar above 10% of their calorie intake. The marginal externality from sugar-sweetened beverage consumption therefore varies across products based on their sugar content, and varies across consumers based on the total amount of added sugar in their diet. The average marginal externality due to 10 oz of sugar-sweetened beverage consumption, for those with dietary sugar exceeding 10% of total calories, is 14 pence. Below we show how our results vary with alternative assumptions about the nature of externalities.

If a marginal change in drinks tax policy causes consumption changes outside the drinks market, distortions from the exercise of market power in the supply of non-drinks products will impact optimal policy. In our baseline results we assume that the numeraire good is competitively supplied. However we show below that, since the impact of a marginal tax change on total drinks spending (and hence numeraire good consumption) is small, our results are numerically insensitive to this assumption.

6.1 Efficiency-maximizing policy

Efficiency-maximizing tax policy minimizes the allocative distortions resulting from governmental (distortionary tax) and non-governmental (externality and market

[23]The majority of these savings come from the reduced costs of treating diabetes and cardiovascular disease. They do not account for the potential impact on, for example, social security from people living longer.

power) distortions. It is indifferent to the distribution of welfare gains across individuals. The efficiency-maximizing policy corresponds to the maximum of the social welfare function (equation (2.1)) when all individuals have social marginal welfare weights equal to one.

To illustrate the impact of market power on the efficiency properties of sin taxation, we consider a single tax rate applied to the set of sugar-sweetened beverages (i.e., the sin products, which comprise product set $\mathcal{S}$). We summarize the impact of the tax in the first row of Table 6.1. The efficiency-maximizing tax rate is 4.19 pence per 10 oz; at the time of writing, US and UK sugar-sweetened beverage taxes range from 7 to 15 pence per 10 oz. The tax leads to an average increase in the price of sugar-sweetened drinks of 19.0%, and little change, on average, in the price of alternative drinks. This results in a 28.7% reduction in consumption of sugar-sweetened beverages, and a 7.0% increase in the consumption of alternative drinks.[24] The tax leads to substantial losses in consumer surplus (£510m, compared to total expenditure in the overall drinks market of approximately £9b) and moderate profit losses (£190m, or around 3.5% fall in market variable profits). However, this is more than made up for by a £386m (25%) fall in externality costs and £409m in excise tax revenue. Overall, economic efficiency increases by £94m.

The efficiency-maximizing tax rate consists of two components – one reflecting distortions from externalities, the other distortions from the exercise of market power for sin and alternative (substitute) products. To illustrate the role that each plays in determining efficiency-maximizing tax policy, it is instructive to consider two naive policies, one that ignores distortions from market power for all goods, and one that ignores them for alternative products.[25]

A government that completely ignores distortions from market power would set a tax rate equal to 11.97 pence per 10 oz. This policy ignores the fact that equilibrium prices are set in excess of marginal costs and results in a sub-optimally high rate. The second row in Table 6.1 shows that this leads to a fall in consumer surplus and profits that is over twice as large as under efficiency-maximizing policy, with the combined loss outweighing the fall in externality costs plus tax revenue, meaning that overall economic efficiency falls by £71m. A government that takes account of market power distortions but only among the taxed sin products, would choose to set a tax rate that is approximately zero; the positive equilibrium margins

[24]This fall in sugar-sweetened beverage consumption is due to an 14.3% reduction in the probability, on average, a consumer purchases from this product set and, conditional on buying, a reduction in volume of 11.9%.

[25]In the first case the government sets τ that satisfies $\tau = \bar{\phi} + \frac{n(\mathcal{S})}{d\mathbb{X}_{\mathcal{S}}/d\tau_{\mathcal{S}}}\text{cov}(\phi^i_j, dx^i_j/d\tau_{\mathcal{S}})$, in the second case it sets the rate that satisfies $\tau = \bar{\phi} + \frac{n(\mathcal{S})}{d\mathbb{X}_{\mathcal{S}}/d\tau_{\mathcal{S}}}\text{cov}(\phi^i_j, dx^i_j/d\tau_{\mathcal{S}}) - \bar{\mu}_{\mathcal{S}}$.

for the sin goods offset the externality associated with their consumption. However, this policy results in a sub-optimally low tax rate. It ignores the fact that other products (substitutes to sin products) are supplied non-competitively. As we show in Tables 5.2 and 5.3, substitution to alternative drinks products is substantial, and these products (like sin products) have substantial mark-ups. Efficiency-maximizing policy depends on the average margins for sin goods relative to those on alternatives.

Table 6.1: *Efficiency-maximizing single rate sin tax policy*

	Tax rate (p/10oz)	% change in Price	% change in Cons.	Change (relative to zero tax) in: (£m) Consumer surplus	Total profits	Excise tax revenue	External cost savings	Total efficiency
Optimal	4.19	19.0%	-28.7%	-510	-190	409	386	94
	[3.93, 4.54]	[19.3%, 18.4%]	[-27.8%, -29.3%]	[-526, -494]	[-198, -181]	[398, 420]	[361, 402]	[77, 110]
Pigouvian	11.97	55.8%	-59.2%	-1199	-429	749	808	-71
	[11.81, 12.17]	[58.2%, 54.4%]	[-58.2%, -60.2%]	[-1244, -1167]	[-446, -407]	[733, 772]	[765, 842]	[-104, -39]

Notes: Optimal refers to efficiency-maximizing policy. Pigouvian refers to policy set by a government that ignores distortions from the exercise of market power. Price and consumption changes are for sugar-sweetened beverages. Consumer surplus, Total profits, External costs, Excise tax revenue and Total efficiency numbers are per annum and report the change relative to no drinks taxation. Total profits are inclusive tax revenues from taxation of corporate profits. Total efficiency = Consumer surplus+Total profits+Excise tax revenue+External cost savings. 95% confidence intervals are given in square brackets.

6.2 Optimal policy with distributional concerns

When the government has distributional concerns, it must balance efficiency with equity considerations. In this case, the distribution of the effects of tax policy across individuals, which depends both on the distribution of consumption and profit holdings, matters for optimal policy.

Profits flow to the government (via corporate and dividend taxes), and to domestic and overseas residents. Measuring stock ownership across the income distribution is challenging and remains a topic of considerable debate. Recent papers, including Saez and Zucman (2016) and Smith et al. (2020), use a combination of dividend income and realized capital gains to estimate wealth in publicly traded stocks. In this spirit, we use information from the UK national accounts and the distribution of dividend income to allocate profits to different groups. The effective average corporate (see Bilicka and Devereux (2012)) and dividend tax rates leads to the government collecting 29% of profits. Using data from the national accounts, we set the fraction flowing overseas to 30%. We assume that the remaining 41% is distributed to UK residents in proportion to the share of (net-of-tax) dividend income received by households in equivalized income bands. The net-of-tax profit holdings of domestic residents is concentrated among the relatively wealthy; households with equivalized income below £10k make up around 25% of the population, but receive less than 3% of post-tax domestic dividend income; households with equivalized income above £45k comprise 5% of the population, but receive more than 21%. Sugar-sweetened beverage taxation will mainly affect the profits of drinks firms. We assume that profit holdings in these firms are approximated by profit holding in the economy more generally (which is reasonable based on diversified investment portfolios). However, as a robustness exercise we also show results in the case when post-tax profits flow to individuals with zero marginal social welfare weights. In Appendix E, we provide full details of these calculations.

We parameterize the social marginal welfare weights as $g^i = (\tilde{y}_i)^{-\vartheta}$, where $\tilde{y}_i$ is equivalized household income and ϑ captures the degree of inequality aversion in government preferences. As our baseline we set $\vartheta = 1$. In all calculations, we assume the government places a social marginal welfare weight of zero on the portion of profits that flows to overseas individuals.

We first illustrate how the distribution of consumption and profit holdings affect policy for a fixed set of government preferences for equity, before showing how the strength of these preferences affect the optimal rate. In the final part of this section, we show how the market structure and nature of externalities affects optimal policy.

How the distribution of consumption and profits affects policy

In Table 6.2 we summarize the impact of distributional concerns on the optimal sin tax rate. The first row of the table shows the optimal rate, and its impact on welfare, under our measure of the true distribution of profit holdings (as described above). The optimal tax rate in this case is 5.97 pence per 10 oz – over 40% higher than the efficiency-maximizing rate (4.19p/10oz) – and it achieves an increase in social welfare of £167m. There are three channels through which distributional concerns affect the optimal rate. We conduct two thought experiments, based on counterfactual distributions of profits, to highlight the relative importance of these different channels.

First, we consider what the optimal rate would be if the government were to collect all profits as tax revenue (row two of the table). This isolates the impact of distributional concerns arising purely from consumption patterns. Since sugar-sweetened beverages are more popular with low income households, the optimal tax in this case lies *below* the efficiency-maximizing rate (it is 3.23 pence per 10 oz). A similar effect is highlighted by Allcott et al. (2019a). Second, we consider optimal policy if the share of profits flowing to overseas individuals equals the true share, but the government collects all domestically owned profits as tax revenue (row three). This introduces a second channel through which distributional concerns impact optimal policy; since a share of profits are owned by foreigners, who, under our social preference specification, are assigned social marginal welfare weights of zero, the government places less weight on the fall in profits resulting from tax policy. This leads the optimal tax rate to increase from 3.23 (when government is assumed to tax fully foreign as well as domestic profits) to 4.71 pence per 10 oz.

Comparing the final row with the first illustrates the importance of the third channel – the impact of the unequal distribution of domestic profits in combination with the government's preference for equity. This leads the optimal tax rate to increase further (from 4.71 to 5.97 pence per 10 oz). Domestic profits are disproportionately in the hands of high income households, meaning the incidence of profit losses associated with the sin tax is mainly on those with relatively low social marginal welfare weights. All else equal, this increases the progressivity of the tax, raising the optimal rate and the size of associated welfare gains.

Table 6.2: *Impact of distributional concerns on optimal single rate sin tax policy*

		Change (relative to zero tax) in:						
		Welfare components (£m)						
	Tax rate	Private welfare, from:		Tax revenue:		Ext. cost	Total
	(p/10oz)	Cons.	Profits	Sin tax	Profit tax	savings	welfare
True profit distribution	5.97	-747	-43	522	-74	509	167
All profits taxed at 100%	3.23	-439	0	336	-152	311	56
Domestic profits taxed at 100%	4.71	-611	0	445	-147	424	111

Notes: Numbers summarize the effect of policy when the social marginal welfare weight on foreign individuals is 0 and on domestic individuals is $1/\tilde{y}_i$, showing effects under the true distribution of profit holdings (row (1)) and counterfactual distributions (rows (2) and (3)). Welfare numbers are per annum and report the change relative to no drinks taxation. Total welfare = Private welfare+Tax revenue+External cost savings. 95% confidence intervals are given in square brackets.

In summary, relative to an efficiency-maximizing government, distributional concerns lead to an increase in the optimal tax rate. This is because post-tax profits are mainly in the hands of foreign and relatively wealthy domestic residents, who have relatively low or zero social marginal welfare weights. This effect more than offsets the influence of distributional concerns over consumption patterns, which, all else equal, act to lower the optimal rate since sugar-sweetened beverages are disproportionately consumed by relatively low income (high social marginal welfare weight) consumers.

How the strength of preferences for equity affects policy

In Figure 6.1 we summarize how differences in the strength of social preferences for equity impact the optimal tax rate (panel (a)), the associated welfare gain (panel (b)), and the fraction of possible welfare gains forgone if the government ignores distortions from the exercise of market power when setting policy[26] (panel (c)). On the horizontal axis we plot efficiency-maximizing policy, and policy when the social marginal welfare weight on foreign profits is 0 and the weights on domestic consumers are $(\tilde{y}_i)^{-\vartheta}$ for $\vartheta = \{0, 1, 2, 3, 4\}$.

Relative to the efficiency-maximizing policy, when $\theta = 0$, the government places a social marginal welfare weight of zero on foreign profits, which acts to raise the optimal tax rate. When $\vartheta > 0$, the government is no longer indifferent to the incidence of the tax across domestic individuals. Thus, the fact that consumption is concentrated among those with low incomes, all else equal, acts to lower the optimal rate. However, this is offset by the concentration of post-tax domestic profit holdings among those with high incomes (meaning they bear a disproportionately large share of profit reductions from sin taxation), which acts to raise the optimal rate. ϑ controls the strength these effects exert on optimal policy. The optimal rate (and associated welfare gain) peaks at $\vartheta = 1$. When $\vartheta > 1$ the impact of consumption inequality on the optimal rate dominates the effect of inequality in domestic profits holding. Overall, the optimal tax rate and associated welfare gain are relatively insensitive to the inequality aversion parameter, and for all values of ϑ, the optimal tax rate and associated welfare gain are larger than under efficiency-maximizing policy.

Panel (c) shows that the welfare losses from ignoring distortions due to the exercise of market power are mitigated under social preferences for equity, but they remain substantial. Policy set by an efficiency-maximizing planner that ignores market power distortions leads to forgone welfare gains of more than 150%. The forgone welfare gains from this form of naive policy when there are distributional concerns are between 64% (when $\vartheta = 0$) and 35% (when $\vartheta \geq 2$). With stronger preferences for equity, since post-tax profits are disproportionately in the hands of those with relatively low or zero social marginal welfare weights, the government's welfare function places less weight on profits (than under efficiency-maximization), and therefore ignoring market power is less costly. However, even with strongly inequality averse preferences there remains a substantial loss from ignoring market power – this is because some profits are collected by the government as tax revenue,

[26] Solving for the τ that satisfies $\tau = \bar{\phi} + \frac{n(\mathcal{S})}{d\mathbb{X}_{\mathcal{S}}/d\tau_{\mathcal{S}}}\text{cov}(\phi^i_j, dx^i_j/d\tau_{\mathcal{S}}) + \frac{1}{d\mathbb{X}_{\mathcal{S}}/d\tau}\text{Cov}\left(g^i, \sum_{j\in\mathcal{S}} x^i_j\right)$.

and because low and moderate income consumers have small, but non-zero, profit holdings.

Figure 6.1: *Impact of different preferences for equity on optimal sin tax policy*

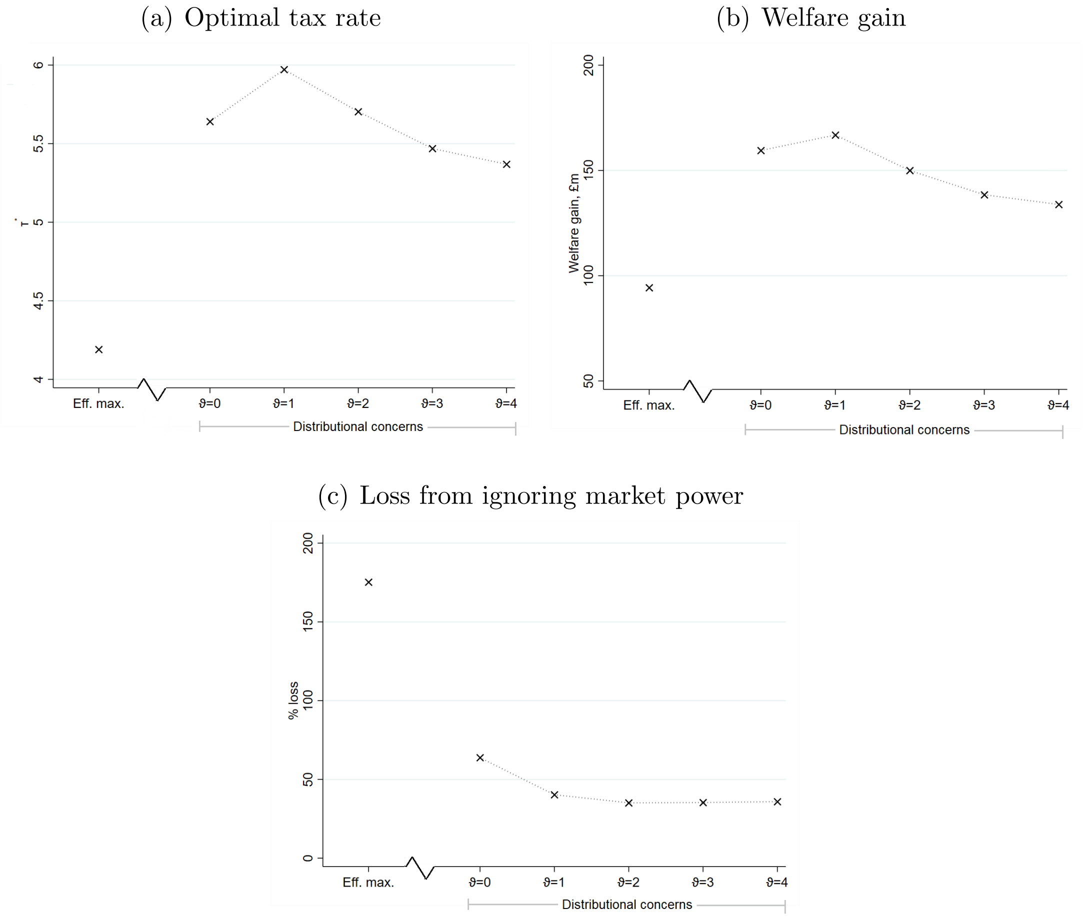

Notes: The horizontal axis plots social preferences. Eff. max. corresponds to an efficiency-maximizing government. Under distributional concerns the social marginal welfare weights on overseas individuals are 0 and the weights on domestic individuals are $(\tilde{y}_i)^{-\vartheta}$*. Panel (a) shows how the optimal tax rate (p/10oz) varies, panel (b) shows how the gain in welfare (per annum) relative to no drinks tax varies and panel (c) shows how the % welfare loss from policymaking that ignores market power varies.*

Another possibility is that, while the government has inequality averse preferences with respect to consumption, it places a social marginal welfare weight of zero on all *post*-tax profits, perhaps due to uncertainty in measuring the distribution of post-tax profits across individuals. Even in this case, market power and the realization of profits has an important bearing on optimal sin tax policy, because sin taxation leads to spillovers to the profit (i.e., corporate and dividend) tax bases. The optimal sin tax rate in this case is 6.90 pence per 10 oz. If the government ignores market power when setting the tax rate (and hence the spillovers to the corporate and dividend tax bases) it would set a rate of 11.26, which would result in unrealized welfare gains of 26% (see row (2) in Table 6.3). This illustrates that, even if post-tax profits holdings are more concentrated in the hands of the rich than our measure of the profit distribution suggests (which is based on post-tax dividend income), market power continues to have an important bearing on optimal tax policy.

How market structure affects policy

Our baseline results assume that the numeraire good is competitively supplied. However, it is possible that when consumers switch away from drinks they switch to other non-competitively supplied goods. This acts to raise the optimal tax rate through an efficiency channel, as it dampens the market power correction element of the optimal tax (which equals the average margin of sin goods relative to alternatives). Yet it acts to lower the optimal rate through an equity channel, as switching to a non-competitively supplied numeraire good leads to off-setting profit gains mainly for the rich. Row (2) in Table 6.3 shows results when the numeraire good is supplied non-competitively, with a margin equal to that implied by the estimate of the UK economy-wide mark-up in De Loecker and Eeckhout (2018). It shows that the optimal tax rate is 6% higher and the associated welfare gain marginally higher, than in our baseline with a competitively supplied numeraire. The modest impact of adding market power for the numeraire on optimal policy is due to the relatively small fall in overall drinks expenditure (and hence rise in numeraire good consumption) induced by a marginal tax rise.

In row (3) of Table 6.3 we consider optimal tax policy under the counterfactual market structure of a perfectly competitive drinks market. Without any tax in place, moving from the true market structure to perfect competition raises welfare by £622m; the gains from eliminating market power distortions dwarf the resulting increase in externality costs. The optimal tax rate under perfect competition is 10.18 pence per 10 oz and the resulting welfare gain is £515 million, which is around 3 times as large as under the true market structure (row (1)). Competition and optimal tax policy therefore exhibit a form of complementarity.

How the nature of externalities affects policy

Rows (5)-(10) of Table 6.3 show how changes in the shape of the function mapping sugar-sweetened beverage consumption to externalities impact tax policy. Rows (5) and (6) vary the overall magnitude of external costs, setting them at 25% above and 25% below the baseline value, implied by Wang et al. (2012). Higher consumption externalities lead to an increase in the optimal tax rate and the associated welfare gains. However, even with larger externality costs, the losses from ignoring distortions from the exercise of market power remain substantial (at 23%).

Row (7) and (8) show results when we vary the convexity of the externality function (holding fixed the average marginal externality across consumers). Row (7) assumes marginal externalities are equal for all consumers; row (8) assumes that only those with total dietary sugar above 15% (compared with 10% in the baseline) generate externalities through their consumption. The consumption of people with higher overall dietary sugar is moderately more responsive to tax changes. Therefore, the more concentrated externality generation is among overall high sugar consumers, the more effective the sin tax is at reducing externality distortions. Hence, more convexity in the externality function leads to a higher optimal rate and larger resulting welfare gains.

We also consider the consequence of externality spillovers. In particular, consumers respond to the sugar-sweetened beverage tax by, in part, switching to pure fruit juices and flavored milk. These are typically exempt from drinks taxes partly because these products contain other (positive) nutrients that may offset negative consequences of sugar intake. However, this is subject to debate. In row (9) we show the consequences for optimal policy when externalities are associated with the untaxed alternative sources of sugar. This type of externality leakage diminishes the effectiveness of tax in reducing externality distortions, leading to a lower optimal rate and smaller welfare gains (relative to the baseline, where consumption of these alternative products is assumed not to generate externalities). Row (10) quantifies the gains of including pure fruit juices and flavored milk in the tax base, in the case where consumption of these products is associated with externalities. This results in a higher optimal rate and welfare gains that are £43m larger – untaxed externalities lead to sizeable forgone welfare gains.

In the final row of Table 6.3 we illustrate how optimal policy changes if social costs also include internalities. We assume externalities are the same as in the baseline case, but that there are also internalities, such that the average social cost of 10g of addition sugar consumption is 25% larger than in the baseline (and the same as in row (5)). We assume that the internality associated with an additional unit of sugar consumption (in monetary terms) for someone at the bottom of the income distribution is 30% larger than for someone at the top (based on the evidence from Allcott et al. (2019a)). In this case, the optimal tax rate is 8.85 pence per 10 oz. This is marginally higher than the optimal rate when the additional 25% of social costs are due to externalities (row (5)). This is because the presence of internalities (negatively related in to income in severity) act to make the tax more progressive, raising the optimal rate. It remains the case that ignoring market power results in substantial forgone welfare gains.

Table 6.3: *Optimal tax policy and the costs of ignoring market power*

		Optimal policy		Ignoring market power		
		Tax rate (p/10oz)	Change in welfare (£m)	Tax rate (p/10oz)	Change in welfare (£m)	% loss
(1)	Baseline	5.97	167	11.26	100	40%
		[5.51, 6.37]	[142, 188]		[58, 129]	[31, 59]
(2)	Zero weight on post-tax profits	6.90	213	11.26	157	26%
		[4.31, 7.34]	[185, 238]		[114, 188]	[21, 38]
Market structure						
(3)	Numeraire good market power	6.34	169	11.26	116	32%
		[5.76, 6.87]	[140, 194]		[70, 150]	[23, 50]
(4)	Perfect competition	10.18	515	10.18	515	0%
		[9.73, 10.67]	[457, 567]		[457, 567]	-
Externalities						
(5)	25% larger	8.44	312	14.63	239	23%
		[7.87, 8.84]	[274, 343]		[187, 281]	[18, 32]
(6)	25% smaller	3.49	61	8.01	-2	103%
		[3.05, 3.79]	[47, 72]		[-31, 18]	[75, 166]
(7)	Linear	5.35	138	10.33	73	47%
		[4.92, 5.73]	[116, 159]		[37, 105]	[34, 68]
(8)	More convex	7.88	267	13.98	194	27%
		[4.33, 5.48]	[226, 314]		[139, 262]	[17, 39]
(9)	Leakage	5.24	153	10.12	105	31%
		[4.78, 5.63]	[131, 172]		[69, 132]	[23, 48]
(10)	Broader base (+leakage)	6.25	196	12.97	109	44%
		[5.77, 6.65]	[172, 223]		[66, 152]	[32, 61]
(11)	+Internalities	8.85	338	15.23	263	22%
		[8.05, 9.22]	[303, 367]		[217, 303]	[17, 28]

Notes: The first two columns summarize optimal policy, the final three summarize policy set by government that ignores distortions from market power (with the final column showing the % of welfare gains from optimal policy forgone). All numbers are based on social marginal welfare weights on foreign individuals of 0 and on domestic individuals of $1/\tilde{y}_i$. Row (1) repeats numbers under our central calibration of the numeraire good margin and externalities. Row (2) sets social marginal welfare weights on post-tax profits to zero. The remaining rows present results under alternative market structures and externality functions, with details described in the text. Welfare numbers are per annum and report the change relative to no drinks taxation. 95% confidence intervals are given in square brackets.

6.3 Optimal policy with alternative tax instruments

Our focus to this point is on a volumetric single tax rate levied on sugar-sweetened beverages. This form of taxation is very common among the taxes that have actually been implemented.[27] In this section we consider alternative tax instruments. In

[27] As of mid 2021, of the 44 countries and 9 US cities that have implemented sugar-sweetened beverage taxes, only Mauritius, South Africa and Sri Lanka have taxes levied directly on sugar

particular, we consider a system of multiple volumetric tax rates applied in the drinks market, and a tax that is levied directly on product sugar content.

Multi-rate system

Alcohol taxation typically involves separate rates for different alcohol types. We consider a similar system applied in the drinks market, in which the government can set rates that vary across broad drink types (see Appendix E.4 for details). Under the constraint that policy cannot lead to a deterioration in the government's budget, inclusive of the budgetary externality, optimal policy involves subsidizing non-sugar sweetened drinks paid for by a combination of taxes on sugar-sweetened drinks and reductions in externalities. Overall it raises welfare by 80% more than than under the optimal single rate system. When subsidies are prohibited, optimal policy entails varying the tax rate across different sugar sweetened beverage types. Relative to under the optimal single sugar sweetened beverage tax rate, the average price rise for sin products and fall in externalities is similar. However, rate differentiation enables the multi-rate system to better target inefficiencies, resulting in a smaller loss in private welfare than under the single rate system and a welfare rise that is 17% higher.

A sugar tax

An alternative to the volumetric sugar-sweetened beverage taxes typically used is to levy a tax directly on the sugar in drinks. This has the advantage that the tax is proportional to the total quantity of the externality generating attribute that a product contains. We compute the optimal tax on the sugar in sweetened beverages, assuming firms do not change their products' sugar contents. Compared with the optimal single rate volumetric tax, it leads to slightly higher prices, larger falls in consumer welfare and lower tax revenue, however this is more than offset by larger reductions in externalities. Overall, the sugar tax raises welfare by 23% more than the single rate volumetric tax (see Table E.4 in the Appendix). The foregone welfare gains from ignoring market power are similar under the sugar and volumetric taxes.

One possible advantage of a sugar tax is that it would incentivize input substitution, as a firm can lower its exposure to the tax by reducing the sugar content of its products. The impact of input substitution on welfare will depend on how the change impacts production costs and consumers' valuation of the products. In the Appendix we follow Barahona et al. (2021) by assuming firms can lower sugar

content. In the large majority of cases, sugar-sweetened beverage taxes entail a single rate. Exceptions are the UK and Portugal, which have banded systems with two rates.

content in their products at the cost of raising production costs, but without altering product taste. Our simulations suggest that the smaller is the increase in production costs from removing sugar, the larger are the welfare gains associated with a sugar tax over a volumetric tax rise (see Appendix E.5). While firms choose to reformulate products to maximize profits, the associated externality reductions are sufficiently high that social welfare increases. An important avenue for future research is to estimate the consumer welfare impact of firms altering their products to avoid taxes.

7 Conclusion

In this paper we show how market power affects the efficiency and redistributive properties of sin taxation. Allocative distortions from the exercise of market power lead optimal policy to depend on the magnitude of equilibrium price-cost margins on sin products, relative to alternatives. The relative concentration of profit holdings in the hands of the wealthy, leads policy to be more progressive than if no profits were realized, counteracting the regressive incidence of the tax based on consumption patterns alone. We show, in an application to sugar sweetened beverage taxation, that market power exerts a quantitatively significant impact on optimal policy and ignoring it leads to substantial unrealized welfare gains. We believe our results are both of direct relevance for the design and implementation of sugar-sweetened beverage taxation and yield important lessons for the design of sin taxes levied in other specific markets. We conclude by suggesting two promising avenues for future research.

Allcott et al. (2019a) highlight how the corrective motive arising from consumer misoptimization interacts with government redistributive objectives in sin tax design: if consumers with high social marginal welfare weights are more likely to make mistakes, this acts to make tax policy more progressive. An interaction between consumer misoptimization and market power arises if firms exploit consumer mistakes, for instance, through persuasive advertising, exploiting consumer self-control problems or obfuscating the unhealthy nature of products (e.g., see Spiegler (2010)). An interesting avenue for future research is to explore the implications of such strategies for sin tax design.

We focus in this paper on illustrating how the non-competitive structure of a market can have an important bearing on the welfare effects of taxes levied in that market. We show how optimal policy varies with competitive conditions, as well as for different levels of profit taxes, under the assumption that the government optimizes sin taxation while treating competition policy fixed. When considering tax policy in a single, relatively small, market this assumption is both reasonable and realistic. However, in other contexts, when tax policy directly affects a large swath of the economy, such as that targeted at climate change, the case for (and gains from) also adjusting competition policy are likely to be larger. Several papers highlight the importance of market power in markets associated with greenhouse gas emissions (e.g, see Bushnell et al. (2008) for electricity and Hastings (2004) for gasoline). A second promising avenue for future work is to consider the joint determination of commodity taxation with other aspects of policy, including competition policy and corporate taxation, in tackling systemic externalities, such as climate change.

References

Allcott, H., B. B. Lockwood, and D. Taubinsky (2019a). Regressive Sin Taxes, with an Application to the Optimal Soda Tax. *Quarterly Journal of Economics 134*(3), 1557–1626.

Allcott, H., B. B. Lockwood, and D. Taubinsky (2019b). Should We Tax Soda? An Overview of Theory and Evidence. *Journal of Economic Perspectives 33*(3), 23.

Auerbach, A. J. and J. R. Hines (2002). Taxation and Economic Efficiency. In *Handbook of Public Economics*, Volume 3, pp. 1347–1421.

Barahona, N., C. Otero, S. Otero, and J. Kim (2021). Equilibrium Effects of Food Labeling Policies.

Berry, S., J. Levinsohn, and A. Pakes (1995). Automobile prices in market equilibrium. *Econometrica 63*(4), 841–890.

Berry, S. T. (1994). Estimating discrete-choice models of product differentiation. *RAND Journal of Economics 25*(2), 242–262.

Bilicka, K. and M. Devereux (2012). CBT Corporate tax ranking 2012. *Oxford University Centre for Business Taxation*, 37.

Bollinger, B. K. and S. Sexton (2018). Local Excise Taxes, Sticky Prices, and Spillovers: Evidence from Berkeley's Soda Tax. *SSRN Electronic Journal*.

Bonnet, C. and P. Dubois (2010). Inference on vertical contracts between manufacturers and retailers allowing for nonlinear pricing and resale price maintenance. *RAND Journal of Economics 41*(1), 139–164.

Bonnet, C. and V. Réquillart (2013). Tax incidence with strategic firms in the soft drink market. *Journal of Public Economics 106*, 77–88.

Bresnahan, T. F. (1989). Empirical studies of industries with market power. In *Handbook of Industrial Organization*, Volume 2, pp. 1011–1057.

Briggs, A. D. M., O. T. Mytton, A. Kehlbacher, R. Tiffin, M. Rayner, and P. Scarborough (2013, October). Overall and income specific effect on prevalence of overweight and obesity of 20% sugar sweetened drink tax in UK: Econometric and comparative risk assessment modelling study. *BMJ 347*, f6189.

Browning, M. J. (1983). Necessary and Sufficient Conditions for Conditional Cost Functions. *Econometrica 51*(3), 851–856.

Buchanan, J. M. (1969). External Diseconomies, Corrective Taxes, and Market Structure. *American Economic Review 59*(1), 174–177.

Bushnell, J. B., E. T. Mansur, and C. Saravia (2008). Vertical Arrangements, Market Structure, and Competition: An Analysis of Restructured US Electricity Markets. *American Economic Review 98*(1), 237–266.

Cawley, J., D. Frisvold, A. Hill, and D. Jones (2020). The Impact of the Philadelphia Beverage Tax on Prices and Product Availability. *Journal of Policy Analysis and Management 39*(3), 605–628.

Chernozhukov, V., J. Hausman, and W. Newey (2019). Demand analysis with many prices. *NBER Working Paper 02138*.

Chetty, R. (2009). Sufficient Statistics for Welfare Analysis: A Bridge Between Structural and Reduced-Form Methods. *Annual Review of Economics 1*, 451–488.

Colchero, M. A., J. C. Salgado, M. Unar-Munguía, M. Molina, S. Ng, and J. A. Rivera-Dommarco (2015). Changes in Prices After an Excise Tax to Sweetened Sugar Beverages Was Implemented in Mexico: Evidence from Urban Areas. *PLOS ONE 10*(12), e0144408.

Competition Commission (2000). *Supermarkets: A Report on the Supply of Groceries from Multiple Stores in the United Kingdom.* London: The Stationary Office.

Competition Commission (2013). *AG Barr / Britvic Merger Inquiry.* London: The Stationary Office.

Conlon, C. T. and N. S. Rao (2015). The Price of Liquor is Too Damn High: Alcohol Taxation and Market Structure. *NYU Wagner Research Paper No. 2610118*.

Cooper, M., J. McClelland, J. Pearce, R. Prisinzano, J. Sullivan, D. Yagan, O. Zidar, and E. Zwick (2016). Business in the United States: Who Owns it and How Much Tax Do They Pay? *Tax Policy and the Economy*, 41.

De Loecker, J. and J. Eeckhout (2018). Global Market Power. *NBER Working Paper 24768*.

De Loecker, J., J. Eeckhout, and G. Unger (2020). The Rise of Market Power and the Macroeconomic Implications. *Quarterly Journal of Economics 135*(2), 561–644.

Diamond, P. A. (1973). Consumption externalities and corrective imperfect pricing. *The Bell Journal of Economics 4*(2), 526–538.

Dubois, P., R. Griffith, and M. O'Connell (2018). The Effects of Banning Advertising in Junk Food Markets. *Review of Economic Studies 85*(1), 396–436.

Dubois, P., R. Griffith, and M. O'Connell (2020). How well targeted are soda taxes? *American Economic Review 110*(11).

Einav, L., A. Finkelstein, and P. Tebaldi (2019). Market Design in Regulated Health Insurance Markets: Risk Adjustment vs. Subsidies. *Mimeo*.

Einav, L. and J. Levin (2010, May). Empirical Industrial Organization: A Progress Report. *Journal of Economic Perspectives 24*(2), 145–162.

Fowlie, M., M. Reguant, and S. P. Ryan (2016). Market-Based Emissions Regulation and Industry Dynamics. *Journal of Political Economy 124*(1), 55.

Global Food Research Program (GFRP) (2021). Sugary drinks taxes around the world. http://globalfoodresearchprogram.web.unc.edu/resources/.

Gorman, W. M. (1980). A Possible Procedure for Analysing Quality Differentials in the Egg Market. *Review of Economic Studies 47*(5), 843–856.

Griffith, R., L. Nesheim, and M. O'Connell (2018). Income effects and the welfare consequences of tax in differentiated product oligopoly. *Quantitative Economics 9*(1), 305–341.

Griffith, R., M. O'Connell, K. Smith, and R. Stroud (2019). The evidence on the effects of soft drinks takes. *IFS Briefing Note 255*.

Grogger, J. (2017). Soda Taxes And The Prices of Sodas And Other Drinks: Evidence From Mexico. *American Journal of Agricultural Economics 99*(2), 481–498.

Harding, M. and M. Lovenheim (2017). The effect of prices on nutrition: Comparing the impact of product- and nutrient-specific taxes. *Journal of Health Economics 53*, 53–71.

Hastings, J. S. (2004). Vertical Relationships and Competition in Retail Gasoline Markets: Empirical Evidence from Contract Changes in Southern California. *American Economic Review 94*(1), 317–328.

Hendel, I. and A. Nevo (2006a). Measuring the implications of sales and consumer inventory behavior. *Econometrica 74*(6), 1637–1673.

Hendel, I. and A. Nevo (2006b). Sales and Consumer Inventory. *RAND Journal of Economics 37*(3), 543–561.

Hendren, N. (2016). The Policy Elasticity. *Tax Policy and the Economy 30*(1), 51–89.

Joint WHO/FAO Expert Consultation (Ed.) (2003). *Diet, Nutrition, and the Prevention of Chronic Diseases.* Number 916 in World Health Organisation Technical Report Series. Geneva: World Health Organization.

Kaplow, L. (2021). Market Power and Income Taxation. *American Economic Journal: Economic Policy 13*(4), 329–354.

Kleven, H. J. (2021). Sufficient Statistics Revisited. *Annual Review of Economics 13*, 33.

Lancaster, K. (1971). *Consumer Demand: A New Approach.* Columbia University Press New York.

Levell, P., M. O'Connell, and K. Smith (2016). Chapter 9: Excise Duties. In *Green Budget 2016*, Volume 2016, pp. 201–231. London: IFS.

Luco, F. and G. Marshall (2020). The Competitive Impact of Vertical Integration by Multiproduct Firms. *American Economic Review 110*(7), 2041–2064.

Miravete, E. J., K. Seim, and J. Thurk (2018). Market Power and the Laffer Curve. *Econometrica 86*(5), 1651–1687.

Miravete, E. J., K. Seim, and J. Thurk (2020). One Markup to Rule Them All: Taxation by Liquor Pricing Regulation. *American Economic Journal: Microeconomics 12*(1), 1–41.

Nevo, A. (2001). Measuring Market Power in the Ready-to-Eat Cereal Industry. *Econometrica 69*(2), 307–342.

OECD (2019). Prices in the health sector. OECD Indicators.

Piketty, T., E. Saez, and G. Zucman (2018). Distributional National Accounts: Methods and Estimates for the United States. *Quarterly Journal of Economics 133*(2), 553–609.

Polyakova, M. and S. Ryan (2019). Subsidy Targeting with Market Power. *NBER Working Paper 26367*.

Preonas, L. (2019). Market Power in Coal Shipping and Implications for U.S. Climate Policy.

Public Health England (2017). Health matters: Obesity and the food environment. https://www.gov.uk/government/publications/health-matters-obesity-and-the-food-environment/health-matters-obesity-and-the-food-environment–2.

Roberto, C. A., H. G. Lawman, M. T. LeVasseur, N. Mitra, A. Peterhans, B. Herring, and S. N. Bleich (2019). Association of a Beverage Tax on Sugar-Sweetened and Artificially Sweetened Beverages With Changes in Beverage Prices and Sales at Chain Retailers in a Large Urban Setting. *JAMA 321*(18), 1799.

Rojas, C. and E. Y. Wang (2017). Do Taxes for Soda and Sugary Drinks Work? Scanner Data Evidence from Berkeley and Washington. *SSRN Electronic Journal*.

Saez, E., J. Slemrod, and S. H. Giertz (2012). The Elasticity of Taxable Income with Respect to Marginal Tax Rates: A Critical Review. *Journal of Economic Literature 50*(1), 3–50.

Saez, E. and G. Zucman (2016). Wealth Inequality in the United States since 1913: Evidence from Capitalized Income Tax Data. *Quarterly Journal of Economics 131*(2), 519–578.

Scientific Advisory Committee on Nutrition (2015). SACN Carbohydrates and Health Report. https://www.gov.uk/government/publications/sacn-carbohydrates-and-health-report.

Seiler, S., A. Tuchman, and S. Yao (2021). The Impact of Soda Taxes: Pass-through, Tax Avoidance, and Nutritional Effects. *Journal of Marketing Research 58*(1), 24–49.

Smith, M., O. Zidar, and E. Zwick (2020). Top Wealth in America: New Estimates and Implications for Taxing the Rich. pp. 117.

Spiegler, R. (2010). *Bounded Rationality and Industrial Organization*. Oxford University Press.

Syverson, C. (2019). Macroeconomics and Market Power: Context, Implications, and Open Questions. *Journal of Economic Perspectives 33*(3), 23–43.

Tebaldi, P. (2017). Estimating Equilibrium in Health Insurance Exchanges: Price Competition and Subsidy Design under the ACA. *University of Chicago Working Paper*.

Villas-Boas, S. B. (2007). Vertical Relationships between Manufacturers and Retailers: Inference with Limited Data. *Review of Economic Studies 74*(2), 625–652.

Wang, E. Y. (2015). The impact of soda taxes on consumer welfare: Implications of storability and taste heterogeneity. *RAND Journal of Economics 46*(2), 409–441.

Wang, Y. C., P. Coxson, Y.-M. Shen, L. Goldman, and K. Bibbins-Domingo (2012). A Penny-Per-Ounce Tax On Sugar-Sweetened Beverages Would Cut Health And Cost Burdens Of Diabetes. *Health Affairs 31*(1), 199–207.

Weyl, E. G. and M. Fabinger (2013). Pass-Through as an Economic Tool: Principles of Incidence under Imperfect Competition. *Journal of Political Economy 121*(3), 528–583.

World Health Organization (2015). *Sugar Intake for Adults and Children.* Geneva.

Appendix:
For online publication

Optimal sin taxation and market power

Martin O'Connell and Kate Smith

November, 2022

A Optimal tax formulae

A.1 Optimal policy

The social welfare function (equation (2.1)) is:

$$W = \int_i \omega^i V^i + \lambda(T_D^i + T_\Pi(\delta^i \Pi) - \Phi^i) di,$$

where ω^i are Pareto weights, λ is the marginal value of government revenue, and where individual indirect utilities, V^i, tax revenue from commodity and labor taxes, T_D^i, tax revenue from profits taxation, $T_\Pi(\delta^i \Pi)$ and the budgetary externality, Φ^i, are as defined in Section 2.1 of the paper.

Suppose the government has available a set of linear taxes, $\tau_1, \ldots, \tau_K$, where $K \leq J$ and $\mathcal{J}_k \subseteq \mathcal{M}$ is the set of products subject to rate τ_k. Let the product tax rates be a function of a policy parameter, θ, where changes in θ can capture any arbitrary changes in the tax rates. The optimal excise tax system satisfies $dW/d\theta = 0$,[28] which implies:

$$\int_i \frac{dV^i/d\theta}{\alpha^i} + \frac{dT_D^i}{d\theta} + \tau_\Pi^i \delta^i \frac{d\Pi}{d\theta} - \frac{d\Phi^i}{d\theta} + \left(g^i - 1\right) \frac{dV^i/d\theta}{\alpha^i} di = 0,$$

where α^i is the marginal utility of income of individual i and g^i is the social marginal welfare weight for the individual. The utility, tax revenue, profit and externality

[28]We focus here on the case where the government does not face any constraints on its policy choice (aside from the excise taxes being linear). The optimal single rate tax entails an improvement in the government's budget. However, in the multi-rate systems that we consider the unconstrained optimal entails a deterioration in the government's budget. When simulating the optimal multi-rate system we therefore add a constraint to the problem – either the government's budget net of externality costs must not deteriorate, or, the stronger condition, that all tax rates must be non-negative.

derivatives are:

$$
\begin{aligned}
\frac{dV^i/d\theta}{\alpha^i} &= -\sum_{k=1}^{K}\sum_{j\in\mathcal{J}_k}\left(\frac{d\tilde{p}_j}{d\theta}+\frac{d\tau_k}{d\theta}\right)x_j^i+\delta^i(1-\tau_\Pi^i)\frac{d\Pi}{d\theta}\\
\frac{d\Pi}{d\theta} &= \sum_{j\in\mathcal{M}}\left(\mu_j\frac{dX_j}{d\theta}+\frac{d\tilde{p}_j}{d\theta}X_j\right)+\mu_O\frac{dX_0}{d\theta}\\
\frac{dT_D^i}{d\theta} &= \sum_{k=1}^{K}\sum_{j\in\mathcal{J}_k}\left(\frac{d\tau_k}{d\theta}x_j^i+\tau_k\frac{dx_j^i}{d\theta}\right)+\tau_z^i\frac{dz^i}{d\theta}\\
\frac{\partial T_D^i}{\partial\theta} &= \sum_{k=1}^{K}\sum_{j\in\mathcal{J}_k}\frac{d\tau_k}{d\theta}x_j^i\\
\frac{d\Phi^i}{d\theta} &= \sum_{j\in\mathcal{S}}\phi_j^i\frac{dx_j^i}{d\theta}.
\end{aligned}
$$

Substituting these into the first order condition gives:

$$
\underbrace{\int_i\left(\frac{dT_D^i}{d\theta}-\frac{\partial T_D^i}{\partial\theta}\right)di}_{\text{fiscal externality}}+\underbrace{\sum_{j\in\mathcal{M}\bigcup O}\mu_j\frac{dX_j}{d\theta}}_{\text{market power distortions}}-\underbrace{\int_i\sum_{j\in\mathcal{S}}\phi_j^i\frac{dx_j^i}{d\theta}di}_{\text{externality distortions}}+\underbrace{\int_i(g^i-1)\frac{dV^i/d\theta}{\alpha^i}di}_{\text{distributional concerns}}=0.
$$

A tax reform impacts welfare through four channels; it entails a fiscal externality as behavioral adjustments alter tax revenue raised from commodities and labor, it alters the allocative distortions from both market power and externalities due to changes in commodity demands, and it affects the distribution of utilities across individuals (who may have different social marginal welfare weights). When policy is set optimally these factors must exactly offset in response to a small change in the tax system.

Equivalently, we can write the first order condition in terms of the tax rates, $\tau_1,\ldots,\tau_K$: for all $k'=1,\ldots,K$

$$
\begin{aligned}
\sum_{k=1}^{K}\tau_k\sum_{j\in\mathcal{J}_k}\frac{dX_j}{d\tau_{k'}}+\int_i\sum_{j\in\mathcal{M}}\left(\mu_j-\phi_j^i\right)\frac{dx_j^i}{d\tau_{k'}}di+\mu_O\frac{dX_0}{d\tau_{k'}}+\int_i\tau_z^i\frac{dz^i}{d\tau_{k'}}di+&\\
\int_i(g^i-1)\left(-\sum_{j\in\mathcal{M}}x_j^i\frac{dp_j}{d\tau_{k'}}+\delta^i(1-\tau_\Pi)\frac{d\Pi}{d\tau_{k'}}\right)di=0.& \qquad \text{(A.1)}
\end{aligned}
$$

A.2 Single sin tax rate

When a single tax rate is applied to the set of sin products, $\mathcal{S}$, condition (A.1) reduces to:

$$\int_i \sum_{j\in\mathcal{S}} \left(\mu_j + \mathbb{1}\{j\in\mathcal{S}\}\tau_{\mathcal{S}} - \phi_j^i\right) \frac{dx_j^i}{d\tau_{\mathcal{S}}} di + \mu_O \frac{dX_0}{d\tau_{\mathcal{S}}} + \int_i \tau_z^i \frac{dz^i}{d\tau_{\mathcal{S}}} di +$$
$$\int_i (g^i - 1)\left(-\sum_{j\in\mathcal{M}} x_j^i \frac{dp_j}{d\tau_{\mathcal{S}}} + \delta^i(1-\tau_\Pi)\frac{d\Pi}{d\tau_{\mathcal{S}}}\right) di = 0.$$

Letting $d\mathbb{X}_{\mathcal{S}}/d\tau_{\mathcal{S}} \equiv \sum_{j\in\mathcal{S}} dX_j/d\tau_{\mathcal{S}}$ and $\bar{\phi} \equiv \int_i \frac{1}{n(\mathcal{S})}\sum_{j\in\mathcal{S}} \phi_j^i di$, defining

$$\bar{\mu}_{\mathcal{X}} \equiv \sum_{j\in\mathcal{X}} \mu_j \frac{dX_j/d\tau_{\mathcal{S}}}{\sum_{j'\in\mathcal{X}} dX_{j'}/d\tau_{\mathcal{S}}} \quad \text{for } \mathcal{X} = \{\mathcal{S},\mathcal{N},O\}$$
$$\Theta_{\mathcal{X}} \equiv \frac{d\mathbb{X}_{\mathcal{X}}/d\tau_{\mathcal{S}}}{d\mathbb{X}_{\mathcal{S}}/d\tau_{\mathcal{S}}} \quad \text{for } \mathcal{X} = \{\mathcal{S},\mathcal{N}\}$$
$$\rho_j \equiv \frac{dp_j}{d\tau_{\mathcal{S}}}$$

and rearranging the first order condition, we obtain equation (2.2):

$$\tau_{\mathcal{S}}^* = \underbrace{\bar{\phi} + \frac{\text{cov}(\phi_j^i, dx_j^i/d\tau_{\mathcal{S}})}{(1/n(\mathcal{S}))\times d\mathbb{X}_{\mathcal{S}}/d\tau_{\mathcal{S}}}}_{\text{externality correction}} - \underbrace{(\bar{\mu}_{\mathcal{S}} - \bar{\mu}_{\mathcal{N}}\Theta_{\mathcal{N}} - \mu_O\Theta_O)}_{\text{marjet power correction}}$$
$$+ \underbrace{\frac{1}{d\mathbb{X}_{\mathcal{S}}/d\tau_{\mathcal{S}}}\left[\text{cov}\left(g^i, \sum_{j\in\mathcal{M}} x_j^i\rho_j - \delta^i(1-\tau_\Pi^i)\frac{d\Pi}{d\tau_{\mathcal{S}}}\right)\right]}_{\text{distributional concerns}}$$
$$- \underbrace{\frac{d(\int_i \mathcal{T} di)/d\tau_{\mathcal{S}}}{d\mathbb{X}_{\mathcal{S}}/d\tau_{\mathcal{S}}}}_{\text{tax base erosion}}.$$

A.3 Characterization of the tax base erosion component

The base erosion term in equation (2.2) can be expressed in terms of income and price elasticities. To see this first note that the base erosion term can be written:

$$\frac{d(\int_i \mathcal{T}(z^i)di)}{d\tau_{\mathcal{S}}} = \int_i \tau_z^i \frac{dz^i}{d\tau_{\mathcal{S}}} di = \int_i \tau_z^i \sum_{j\in\mathcal{M}} \frac{\partial z^i}{\partial p_j}\frac{dp_j}{d\tau_{\mathcal{S}}} di$$

Assume income effects on labor supply are negligible (see Saez et al. (2012) for support of this). Using Slutsky symmetry and the Slutsky decomposition we can re-write $\frac{\partial z^i}{\partial p_j}$:

$$\frac{\partial z^i}{\partial p_j} = -\frac{\partial \tilde{x}^i_j}{\partial(1-\tau^i_z)} = -\frac{\partial x^i_j}{\partial(1-\tau^i_z)} + z^i \frac{\partial x^i_j}{\partial Y^i}, \tag{A.2}$$

where $\tilde{x}^i_j$ denotes compensated demand for good j, and, as in the paper, Y^i is the sum of the consumer's virtual labor income and their profit income.

We make use of the conditional cost function (see Browning (1983)). The consumer's conditional cost function is defined:

$$e^i(\mathbf{p}, u, \bar{z}^i) = \min_{\mathbf{x}^i}\{\mathbf{p}\mathbf{x}^i : \text{ s.t. } U^i(\mathbf{x}^i, \bar{z}^i) = u\},$$

and gives the minimum expenditure necessary to achieve a given level of utility, *holding labor supply fixed* at $\bar{z}^i$. The associated conditional compensated demand for product j is given by $\tilde{x}^i_j = \frac{\partial e^i(\mathbf{p},u,\bar{z}^i)}{\partial p_j}$. Inverting the conditional expenditure function yields the conditional indirect utility function: $V^i(\mathbf{p}, Y^i_{\bar{z}}, \bar{z}^i)$, where $Y^i_{\bar{z}} \equiv Y^i + (1-\tau^i_z)\bar{z}^i = \mathbf{p}\mathbf{x}^i$. Substituting this into the conditional compensated demand for product j, yields the conditional uncompensated demand $x^i_j = \tilde{f}^i_j(\mathbf{p}, \tilde{Y}^i_{\bar{z}}, \bar{z}^i)$. Let z^i denote the optimal labor supply choice and $\tilde{Y}^i(z^i) \equiv Y^i + (1-\tau^i_z)z^i$ denote total income at this level of labor supply. Then $x^i_j = \tilde{f}^i_j(\mathbf{p}, \tilde{Y}^i(z^i), z^i) = f^i_j(\mathbf{p}, 1-\tau^i_z, Y^i)$ where f^i_j is the unconditional compensated demand.

Consider the derivative of $x^i_j = \tilde{f}^i_j(\mathbf{p}, \tilde{Y}^i(z^i), z^i)$ with respect to $(1-\tau^i_z)$:

$$\begin{aligned}
\frac{\partial x^i_j}{\partial(1-\tau^i_z)} =& \frac{\partial \tilde{f}^i_j}{\partial \tilde{Y}^i} \frac{d\tilde{Y}^i}{d(1-\tau^i_z)} + \frac{\partial \tilde{f}^i_j}{\partial z^i} \frac{\partial z^i}{\partial(1-\tau^i_z)} \\
=& \frac{\partial \tilde{f}^i_j}{\partial \tilde{Y}^i}\left(z^i + (1-\tau^i_z)\frac{\partial z^i}{\partial(1-\tau^i_z)}\right) + \frac{\partial \tilde{f}^i_j}{\partial z^i} \frac{\partial z^i}{\partial(1-\tau^i_z)} \\
=& \frac{\partial \tilde{f}^i_j}{\partial \tilde{Y}^i} z^i + \left(\frac{\partial \tilde{f}^i_j}{\partial \tilde{Y}^i}(1-\tau^i_z) + \frac{\partial \tilde{f}^i_j}{\partial z^i}\right)\frac{\partial z^i}{\partial(1-\tau^i_z)} \\
=& \frac{\partial \tilde{f}^i_j}{\partial \tilde{Y}^i} z^i + \frac{\partial f^i_j}{\partial z^i} \frac{\partial z^i}{\partial(1-\tau^i_z)}
\end{aligned} \tag{A.3}$$

Where: equality (2) follows from the definition of $\tilde{Y}^i(z^i)$ and our assumption that $\mathcal{T}$ is piecewise linear; equality (3) follows from rearranging; and equality (4) follows from the definition of $\tilde{Y}^i(z^i)$ and $\tilde{f}^i_j(\mathbf{p}, \tilde{Y}^i(z^i), z^i) = f^i_j(\mathbf{p}, 1-\tau^i_z, Y^i)$.

As we have assumed that there are no income effects on labor supply:

$$\frac{\partial \tilde{f}^i_j}{\partial \tilde{Y}^i} = \frac{\partial f^i_j}{\partial Y^i}. \tag{A.4}$$

Combining conditions (A.2)-(A.4) yields:

$$\begin{aligned}\frac{\partial z^i}{\partial p_j} &= -\frac{\partial f_j^i}{\partial z^i}\frac{\partial z^i}{\partial(1-\tau_z^i)} \\ &= -\xi_j^i \zeta_z^i \frac{x_j^i}{(1-\tau_z^i)}\end{aligned}$$

where $\xi_j^i \equiv \frac{\partial f_j^i}{\partial z^i}\frac{z^i}{x_j^i}$ is the elasticity of good j with respect to labor earnings and $\zeta_z^i \equiv \frac{\partial z^i}{\partial(1-\tau_z^i)}\frac{(1-\tau_z^i)}{z^i}$ is the elasticity of taxable earnings.

Hence the tax base erosion terms can be written:

$$\frac{d(\int_i \mathcal{T}(z^i)di)}{d\tau_{\mathcal{S}}} = \int_i \frac{\tau_z^i}{1-\tau_z^i}\zeta_z^i \sum_{j\in\mathcal{M}} \xi_j^i x_j^i \rho_j di.$$

A.4 Incorporating internalities

Suppose individual i's consumption of sugar-sweetened drinks gives rise to an internality, $I^i = I(\mathbf{x}_{\mathcal{S}}^i)$, which impacts their utility but that they ignore when making decisions. In this case we write the consumer's indirect utility function: $V^i = V^i(\nu^i(\mathbf{p},(1-\tau_z^i),Y^i),I^i)$.

Now there is an additional channel through which policy impacts social welfare: through its impact on internalities. The monetary impact of a marginal policy change on individual i is given by:

$$\frac{dV^i/d\theta}{\alpha^i} = -\sum_{j\in\mathcal{M}} x_j^i \frac{dp_j}{d\theta} + \delta^i(1-\tau_\Pi^i)\frac{d\Pi}{d\theta} + \sum_{j\in\mathcal{S}} \psi_j^i \frac{dx_j^i}{d\theta},$$

where $\psi_j^i \equiv \frac{\partial V^i/\partial I^i}{\alpha^i}\frac{dI^i}{dx_j^i}$ is the monetary cost the individual imposes on themself per additional unit of consumption of sin product j. The first order condition for optimal is then becomes:

$$\underbrace{\int_i \left(\frac{dT_D^i}{d\theta} - \frac{\partial T_D^i}{\partial\theta}\right) di}_{\text{fiscal externality}} + \underbrace{\sum_{j\in\mathcal{M}\bigcup O} \mu_j \frac{dX_j}{d\theta}}_{\text{market power distortions}} - \underbrace{\int_i \sum_{j\in\mathcal{S}} (\phi_j^i + \psi_j^i)\frac{dx_j^i}{d\theta} di}_{\substack{\text{externality and}\\ \text{internality distortions}}} + \underbrace{\int_i (g^i-1)\frac{dV^i/d\theta}{\alpha^i} di}_{\text{distributional concerns}} = 0.$$

There are two differences with the condition under no internalities (equation A.1). First the externality distortion term is adjusted to also capture the internality distortion. This term reflects how policy reform affects economic efficiency through its impact on externality and internality distortions. The second difference is that the distributional concerns term is now influenced by the covariance of internality changes and social marginal welfare weights; all else equal, for a government with preferences for equity, the more internality reductions are concentrated among the relatively poor, the more valuable is a given fall in the average internality.

B Data

B.1 Sample

Our at-home sample (drawn from the Kantar Worldpanel) comprises 30,405 households and our on-the-go sample (drawn from the Kantar On-The-Go Survey) comprises 2,862 individuals. We omit a small number of consumers that record irregularly. Specifically, in the at-home segment we focus on households that record purchases of any groceries in at least 10 weeks per year and who make at least one drink purchase. In the on-the-go segment we focus on individuals who record at least 5 purchases each year. In each segment, this conditioning drops less than 3% of drinks transactions.

In table B.1 we compare the demographic composition of the Kantar Worldpanel with the nationally representative Living Costs and Food Survey for a single year (2012). It shows that Kantar Worldpanel households are broadly representative of the UK population.

Table B.1: *Household demographics*

	Kantar	LCFS
Region		
North East	4.6	4.8
	[4.3, 4.9]	[4.3, 5.4]
North West	11.2	11.5
	[10.7, 11.6]	[10.6, 12.3]
Yorkshire and Humber	11.3	9.6
	[10.8, 11.7]	[8.8, 10.4]
East Midlands	8.4	7.8
	[8.0, 8.7]	[7.1, 8.6]
West Midlands	8.9	9.5
	[8.5, 9.3]	[8.7, 10.2]
East of England	10.5	10.4
	[10.1, 10.9]	[9.6, 11.2]
London	8.5	9.0
	[8.1, 8.9]	[8.3, 9.8]
South East	14.6	14.4
	[14.2, 15.1]	[13.5, 15.4]
South West	9.1	9.1
	[8.7, 9.5]	[8.3, 9.9]
Wales	4.6	4.9
	[4.4, 4.9]	[4.3, 5.5]
Scotland	8.2	8.9
	[7.9, 8.6]	[8.1, 9.7]
Socioeconomic status		
Highly skilled	20.9	17.4
	[20.3, 21.4]	[16.1, 18.7]
Semi skilled	55.8	53.0
	[55.1, 56.4]	[51.3, 54.7]
Unskilled	23.4	29.6
	[22.8, 23.9]	[28.1, 31.2]
Number of adults		
1	22.1	32.9
	[21.5, 22.6]	[31.7, 34.2]
2	60.8	55.8
	[60.1, 61.4]	[54.5, 57.2]
3+	17.2	11.3
	[16.7, 17.7]	[10.4, 12.1]
Number of children		
1	14.6	14.1
	[14.1, 15.1]	[13.2, 15.0]
2	15.1	11.0
	[14.6, 15.6]	[10.2, 11.8]
3+	6.1	5.1
	[5.8, 6.5]	[4.6, 5.7]

Notes: Table shows the share of households in the Kantar Worldpanel and Living Costs and Food Survey in 2012 by various demographic groups. Socioeconomic status is based on the occupation of the head of the household and is shown for the set of non-pensioner households. 95% confidence intervals are shown below each share.

B.2 Variation in SSB consumption by income

In Figure 3.1(b) of the paper we show that higher income households consume less sugar-sweetened beverages than poorer households. To what extent is this driven by heterogeneity in preferences or causal income effects? To answer this we estimate the following two regressions:

$$\text{volSSB}_{iyq} = \sum_{k=1}^{5} \beta_k^{NOFE} \text{income quintile}_{iy}^k + \epsilon_{iyq} \tag{B.1}$$

$$\text{volSSB}_{iyq} = \sum_{k=1}^{5} \beta_k^{FE} \text{income quintile}_{iy}^k + \mu_i + \epsilon_{iyq} \tag{B.2}$$

where volSSB_{iyq} denotes the volume of sugar-sweetened beverages purchased by household i in year-quarter (y, q) for at-home consumption. $\text{income quintile}_{iy}^k$ is an indicator variable equal to 1 if household i is in income quintile k in year y, and μ_i is a household fixed effect. We estimate this over the period 2008 to 2012.

Figure B.1 plots the estimated $\hat{\beta}_k^{NOFE}$ and $\hat{\beta}_k^{FE}$. It shows that, although in the cross-section there is a negative relationship between household income and volume of sugar-sweetened beverage consumption, this relationship disappears when we control for household fixed effects. This indicates that preference heterogeneity accounts for the variation in sugar-sweetened beverage consumption across the income distribution, with little evidence of causal income effects.

We repeat the analysis summarized Figure 3.1 in the paper, but using the on-the-go sample. This is shown in Figure B.2.

Figure B.1: *Income correlation with sugar-sweetened beverage consumption*

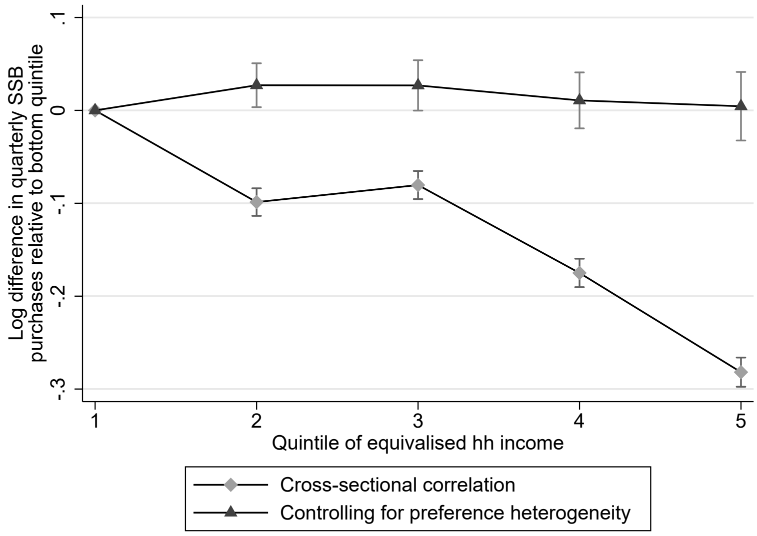

Notes: The light grey markets plot $\hat{\beta}_k^{NOFE}$ from equation (B.1) and the dark grey markets plot $\hat{\beta}_k^{FE}$ from equation (B.2).

Figure B.2: *Variation in volume of sugar-sweetened beverages consumed on-the-go*

(a) Dietary sugar intake

(b) Household income

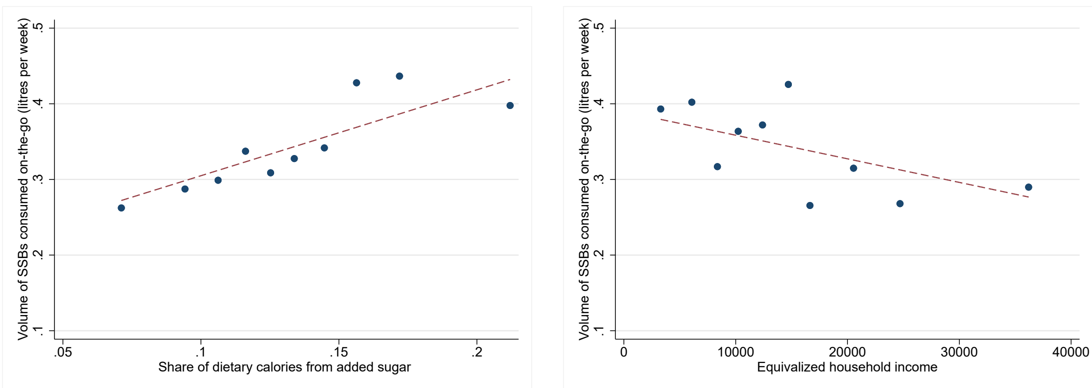

Notes: The left hand panel shows mean volume of sugar-sweetened beverages purchased per person week and consumed on-the-go, by deciles of the share of dietary calories from added sugar (from food consumed at home). The right hand panel shows mean volume of sugar-sweetened beverages purchased per person per week and consumed on-the-go by deciles of equivalized (using the OECD-equivalence scale) household income.

B.3 Brand, products and retailers

In Table B.2 we list the main firms that operate in the drinks market and the brands that they own. The firms Coca Cola Enterprises and Britvic dominate the market, having a combined market share exceeding 65% in the at-home segment and close to 80% in the on-the-go segment. In Table B.3 we list the variants available for each brand. Most brands are available in a regular and diet variant (with some also having an additional zero sugar variant). The table also shows, for each brand-variant, the number of sizes available to consumers in the at-home and on-the-go segments. We include a size option corresponding to multiple units of a single UPC if that UPC-multiple unit combination accounts for at least 10,000 (around 0.2%) transactions. We refer to a brand-variant-size combination as a product.

Table B.2: *Firms and brands*

			Market share (%)		Price (£/l)	
Firm	Brand	Type	At-home	On-the-go	At-home	On-the-go
Coca Cola Enterprises			*33.0*	*59.0*		
	Coke	Soft	20.4	36.3	0.86	2.09
	Capri Sun	Soft	3.1	–	1.08	–
	Innocent fruit juice	Fruit	2.1	1.6	2.04	7.09
	Schweppes Lemonade	Soft	1.7	–	0.44	–
	Fanta	Soft	1.7	5.3	0.79	2.10
	Dr Pepper	Soft	1.2	3.4	0.75	2.08
	Schweppes Tonic	Soft	1.1	–	1.22	–
	Sprite	Soft	1.0	2.8	0.77	2.08
	Cherry Coke	Soft	0.8	4.0	0.96	2.17
	Oasis	Soft	–	5.6	–	2.15
Britvic			*33.6*	*20.0*		
	Robinsons	Soft	10.7	–	1.09	–
	Pepsi	Soft	10.1	11.6	0.64	1.93
	Tropicana fruit juice	Fruit	6.1	3.8	1.62	3.63
	Robinsons Fruit Shoot	Soft	2.7	0.8	1.49	2.83
	Britvic fruit juice	Fruit	1.6	–	2.17	–
	7 Up	Soft	0.9	1.7	0.70	1.88
	Copella fruit juice	Fruit	0.8	–	1.68	–
	Tango	Soft	0.8	2.2	0.66	1.73
GSK			*7.6*	*12.7*		
	Ribena	Soft	3.3	3.4	1.69	2.20
	Lucozade	Soft	3.1	6.4	1.11	2.37
	Lucozade Sport	Soft	1.2	2.9	1.15	2.22
JN Nichols	Vimto	Soft	1.6	–	1.06	–
Barrs	Irn Bru	Soft	0.6	2.6	0.61	1.93
Merrydown	Shloer	Soft	2.0	–	1.79	–
Red Bull	Red Bull	Soft	0.2	3.4	3.66	5.27
Muller	Frijj flavoured milk	Milk	–	1.4	–	1.90
Friesland Campina	Yazoo flavoured milk	Milk	–	0.8	–	1.95
Store brand			*21.3*	*0.0*		
	Store brand soft drinks	Soft	13.1	–	0.62	–
	Store brand fruit juice	Fruit	8.1	–	1.05	–

Notes: Type refers to the type of drinks product: "soft" denotes soft drinks, "fruit" denotes fruit juice, and "milk" denotes flavored milk. The fourth and fifth columns display each firm and brand's share of total spending on all listed drinks brands in the at-home and on-the-go segments of the market; a dash ("–") denotes that the brand is not available in that segment. The final two columns display the mean price (£) per liter for each brand.

Table B.3: *Brands, sugar contents and sizes*

			Sugar	Number of sizes	
Firm	Brand	Variant	(g/100ml)	At-home	On-the-go
Coca Cola Enterprises	Coke	Diet	0.0	10	2
		Regular	10.6	9	2
		Zero	0.0	7	2
	Capri Sun	Regular	10.9	3	–
	Innocent fruit juice	Regular	10.7	4	1
	Schweppes Lemonade	Diet	0.0	2	–
		Regular	4.2	2	–
	Fanta	Diet	0.0	2	1
		Regular	7.9	2	2
	Dr Pepper	Diet	0.0	2	1
		Regular	10.3	2	2
	Schweppes Tonic	Diet	0.0	2	–
		Regular	5.1	2	–
	Sprite	Diet	0.0	2	–
		Regular	10.6	2	2
	Cherry Coke	Diet	0.0	2	1
		Regular	11.2	2	2
	Oasis	Diet	0.0	–	1
		Regular	4.2	–	1
Britvic	Robinsons	Diet	0.0	6	–
		Regular	3.2	6	–
	Pepsi	Diet	0.0	5	2
		Max	0.0	6	2
		Regular	11.0	5	2
	Tropicana fruit juice	Regular	9.6	4	1
	Robinsons Fruit Shoot	Diet	0.0	2	1
		Regular	10.3	2	–
	Britvic fruit juice	Regular	9.9	2	–
	7 Up	Diet	0.0	2	1
		Regular	10.8	2	2
	Copella fruit juice	Regular	10.1	3	–
	Tango	Regular	3.5	3	2
GSK	Ribena	Diet	0.0	2	1
		Regular	10.8	4	2
	Lucozade	Regular	11.3	3	2
	Lucozade Sport	Diet	0.0	1	1
		Regular	3.6	1	1
JN Nichols	Vimto	Diet	0.0	3	–
		Regular	5.9	4	–
Barrs	Irn Bru	Diet	0.0	1	2
		Regular	8.7	1	2
Merrydown	Shloer	Regular	9.1	3	–
Red Bull	Red Bull	Diet	0.0	–	1
		Regular	10.8	1	1
Muller	Frijj flavoured milk	Regular	10.8	–	1
Friesland Campina	Yazoo flavoured milk	Regular	9.5	–	1
Store brand	Store brand soft drinks	Diet	0.0	4	–
		Regular	10.3	2	–
	Store brand fruit juice	Regular	10.4	2	–

Notes: The final two columns displays the number of sizes of each brand-variant in the at-home and on-the-go segments of the market; a dash ("–") denotes that the brand-variant is not available in that segment.

Table B.4 lists retailers and the share of drinks spending that they account for in each segment. In the at-home segment, four large national supermarket chains account for almost 90% of spending, with the remaining spending mostly made in smaller national retailers. Each of these retailers offers all brands, with some variation in the specific sizes available in each retailer. The large four supermarkets are less prominent in the on-the-go segment, collectively accounting for less than 20% of on-the-go spending on drinks. The majority of transactions in the on-the-go segment are in local convenience stores.

Table B.4: *Retailers*

	Expenditure share (%)	
	at-home	on-the-go
Large national chains	87.0	19.9
of which:		
Tesco	34.7	–
Sainsbury's	16.8	–
Asda	19.8	–
Morrisons	15.7	–
Small national chains	10.7	16.4
Vending machines	0.0	9.1
Convenience stores	2.3	54.5
in region:		
South	–	13.6
Central	–	15.5
North	–	25.4

Notes: Numbers show the share of total drinks expenditure, in the at-home and on-the-go segment, made in each retailer.

C Equilibrium model details

C.1 Choice occasions and consumer groups

We observe households for an average of 36 at-home choice occasions and individuals for an average of 44 on-the-go choice occasions each year, with a total of 3.3 million at-home and 286,576 on-the-go choice occasions.

We estimate our demand model allowing all preference parameters to vary by the consumer groups shown in Table C.1. In the at-home segment we split households based on whether there are any children in the household. In the on-the-go segment we separate individuals aged 30 and under from those aged over 30. We also differentiate between those with low, high or very high total dietary sugar. This measure is based on the household's (or, for individuals in the on-the-go sample, the household to which they belong) share of total calories purchased in the form of added sugar across all grocery shops in the preceding year. We classify those that meet the World Health Organization (2015) recommendation of less than 10% of calories from added sugar as "low dietary sugar", those that purchase between 10%

and 15% as “high dietary sugar”, and those that purchase more than 15% of their calories from added sugar as “very high dietary sugar”.

Table C.1: *Consumer groups*

	No. of consumers	% of sample
At-home segment (households)		
No children, low dietary sugar	7500	17
No children, high dietary sugar	11931	27
No children, very high dietary sugar	7292	17
With children, low dietary sugar	3561	8
With children, high dietary sugar	8382	19
With children, very high dietary sugar	5185	12
On-the-go segment (individuals)		
Under 30, low dietary sugar	240	6
Under 30, high dietary sugar	576	15
Under 30, very high dietary sugar	381	10
Over 30, low dietary sugar	601	16
Over 30, high dietary sugar	1319	34
Over 30, very high dietary sugar	757	20

Notes: Columns 2 and 3 show the number and share of consumers (households in the at-home segment, individuals in the on-the-go segment) in each group, respectively. If consumers move group over the sample period (2008-12) they are counted twice, hence the sum of the numbers of consumers in each group is greater than the total number of consumers. Dietary sugar is calculated based on the share of total calories from added sugar purchased in the preceding year; “low” is less than 10%, “high” is 10-15% and “very high” is more than 15%. Households with children are those with at least one household member aged under 18.

C.2 Dependence across at-home and on-the-go segments

Our demand model assumes independence between demand for drinks in the at-home and on-the-go segments of the market. A potential concern is that when people live in a household that has recently purchased drinks for at-home consumption, they will be less likely to purchase drinks on-the-go, thus introducing dependency between the two segments of the market.

We assess evidence for this by looking at the relationship between a measure of a household’s recent at-home drinks purchases and the quantity of drinks an individual from that household purchases on-the-go. We construct a dataset at the individual-day level (we drop days before and after the first and last dates that the individual is observed in the on-the-go sample). The dataset includes the quantity of drinks purchased on-the-go (including zeros), and the total quantity of drinks purchased at home over a variety of preceding time periods.

We estimate:

$$\text{quantity on-the-go}_{it} = \sum_{s=1}^{4} \beta_s \text{week } s \text{ at-home volume}_{it} + \mu_i + \rho_r + \tau_t + \epsilon_{it}$$
$$\text{quantity on-the-go}_{it} = \sum_{d=1}^{7} \beta_d \text{daily } d \text{ at-home volume}_{it} + \mu_i + \rho_r + \tau_t + \epsilon_{it}$$

where week s at-home volume$_{it}$ is the total at-home purchases of drinks made by individual i's household in the s week before day t, and daily d at-home volume$_{it}$ is the total at-home purchases of drinks made by individual i's household on the d day before day t. We estimate both of these regression with and without individual fixed effects to show the importance of individual preference heterogeneity.

Table C.2 shows the estimates. The first two columns show the relationship between the volume of drinks purchased on-the-go and the volume of at-home purchases in the four weeks prior. When we do not include fixed effects, the results are positive and statistically significant. However, in the second column, once we include fixed effects, the results go to almost zero. We see a similar pattern in the final two columns, which show the relationship between volume purchased on-the-go and the daily volume of at-home purchases in the previous 7 days.

These descriptive results provide support for modeling the at-home and on-the-go segments as separate parts of the market. They are also consistent with the formal test of non-separability between the segments conducted in Dubois et al. (2020).

Table C.2: *Dependence across at-home and on-the-go*

	(1) Volume	(2) Volume	(3) Volume	(4) Volume
At-home purchases 1 week before	0.0008***	0.0001**		
	(0.0000)	(0.0000)		
At-home purchases 2 weeks before	0.0008***	0.0001***		
	(0.0000)	(0.0000)		
At-home purchases 3 weeks before	0.0007***	0.0001*		
	(0.0000)	(0.0000)		
At-home purchases 4 weeks before	0.0007***	0.0001*		
	(0.0000)	(0.0000)		
At-home purchases 1 day before			0.0011***	-0.0002
			(0.0001)	(0.0001)
At-home purchases 2 days before			0.0014***	0.0000
			(0.0001)	(0.0002)
At-home purchases 3 days before			0.0012***	-0.0002
			(0.0001)	(0.0001)
At-home purchases 4 days before			0.0015***	0.0002
			(0.0001)	(0.0001)
At-home purchases 5 days before			0.0016***	0.0002
			(0.0001)	(0.0001)
At-home purchases 6 days before			0.0017***	0.0004**
			(0.0001)	(0.0001)
At-home purchases 7 days before			0.0018***	0.0005***
			(0.0001)	(0.0001)
N	2668585	2668585	2776989	2776989
Mean of dependent variable	0.0452	0.0452	0.0452	0.0452
Time effects?	Yes	Yes	Yes	Yes
Decision maker fixed effects?	No	Yes	No	Yes

Notes: Dependent variable in all regressions is the volume of drinks purchased on-the-go (in liters). An observation is an individual-day; data include zero purchases of drinks. Robust standard errors shown in parentheses.

C.3 Price variation

In Figure C.1 we depict graphically price variation, by showing the path of price over one year for two example products in two different retailers. In the example shown in panel (a) a 2l bottle of Coke costs £2 in either retailer for the whole period. However, for most of the time two units of 2l Coke (which we treat as a separate product) is available on a multi-buy offer – where the price per liter is less when the consumer purchases two units. This kind of multi-buy offer is common, accounting for 30% of transactions. Both the depth and timing of the discount varies over time differentially by retailers. Panel (b) shows an example of a product, 12× 330ml cans of Coke, that does not have a multi-buy offer, but rather where the promotion takes the form of a ticket price reduction, or temporarily low price – this type of promotion accounts for 20% of transactions. Again the timing and depth of promotions vary across the retailers.

Figure C.1: *Examples of price variation for Coke options*

(a) 2l bottle (b) 12×330ml cans

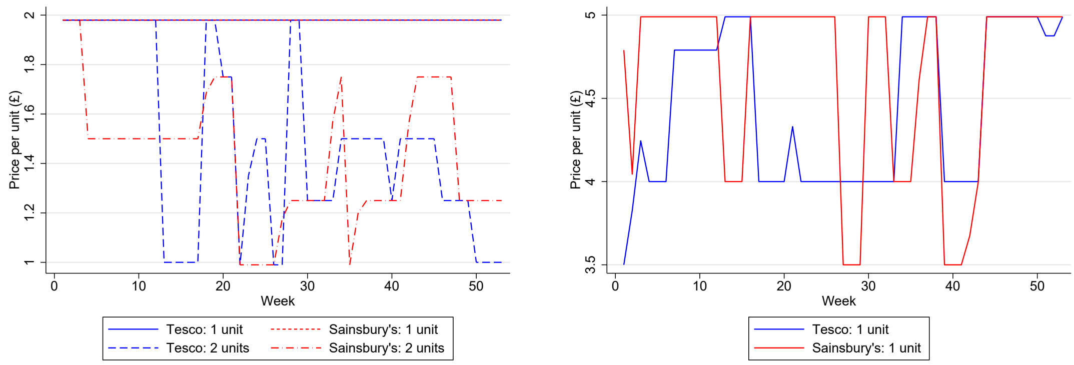

Notes: Panel (a) shows the weekly price series for a 2l bottle of Coke in Tesco and Sainsbury's when either one unit or two units are purchased. Prices are expressed per unit. Panel (b) shows the weekly price series for a pack of 12×330ml cans of Coke in Tesco and Sainsbury's when one unit is purchased.

C.4 Stockpiling

We present evidence regarding whether households in the at-home segment stockpile drinks by conducting a number of checks based on implications of stockpiling behavior highlighted by Hendel and Nevo (2006b). Hendel and Nevo (2006b) highlight the importance of controlling for preference heterogeneity across consumers; throughout our analysis, we focus on within-consumer predictions and patterns of stockpiling behavior.

We construct a dataset that, for each household, has an observation for every day that they visit a retailer. The data set contains information on: (i) whether the household purchased a drink on that day, (ii) how much they purchased, and (iii) the share of volume of drinks purchased on sale. To account for households who do not record purchasing any groceries for a sustained period of time (for instance, because they are on holiday), we construct "purchase strings" for each household. These are sequences that do not contain a period of non-reporting of any grocery purchases longer than 3 or more weeks.

Inventory. One implication of stockpiling behavior highlighted in Hendel and Nevo (2006b) is that the probability a consumer purchases and, conditional on purchasing, the quantity purchased decline in the current inventory of the good.

Inventory is unobserved; following Hendel and Nevo (2006b) we construct a measure of each household's inventory as the cumulative difference in purchases from the household's mean purchases (within a purchase string). Inventory increases if today's purchases are higher than the household's average, and inventory declines if today's purchases are lower than the household's average.

Let i index household, $\tau = (1, \ldots, \tau_i)$ index days on which we observe the household shopping – we refer to this as a shopping trip – r index retailer and t index year-weeks. We estimate:

$$\begin{aligned}
\text{buysoftdrink}_{i\tau} &= \beta^{\text{inv,pp}}\text{inventory}_{i\tau} + \mu_i + \rho_r + t_\tau + \epsilon_{i\tau} \\
q_{i\tau} &= \beta^{\text{inv,q}}\text{inventory}_{i\tau} + \mu_i + \rho_r + t_\tau + \epsilon_{i\tau} \quad \text{if buysoftdrink}_{i\tau} = 1
\end{aligned}$$

where $\text{buysoftdrink}_{i\tau}$ is a dummy variable equal to 1 if household i buys any drinks on shopping trip τ; $q_{i\tau}$ is the quantity of drink purchased, and $\text{inventory}_{i\tau}$ is household i's inventory on shopping trip τ, constructed as described above. μ_i are household-purchase string fixed effects, ρ_r are retailer effects and t_τ are year-week effects.

If stockpiling behavior is present we would expect that $\beta^{\text{inv,pp}} < 0$ and $\beta^{\text{inv,q}} < 0$; when a household's inventory is high it is less likely to purchase, and conditional on purchasing it will buy relatively little. The first two columns of Table C.3 summarize the estimates from these regressions. There is a small positive relationship between inventory and purchase probability and quantity purchased, conditional on buying. An increase in inventory of 1 liter is associated with an increase in the probability of buying of 0.001, relative to a mean of 0.23, and an increase in the quantity purchased, conditional on buying a positive amount, of 0.013, relative to a mean of 3.925. These effects are both very small and go in the opposite direction to that predicted by Hendel and Nevo (2006b) if stockpiling behavior was present.

Time between purchases. The second and third implications of stockpiling behavior highlighted in Hendel and Nevo (2006b) are that, on average, the time to the next purchase is longer after a household makes a purchase on sale, and that the time since the previous purchase is shorter.

We check for this by estimating:

$$\begin{aligned}
\text{timeto}_{i\tau} &= \beta^{\text{lead}}\text{sale}_{i\tau} + \mu_i + \rho_r + t_\tau + \epsilon_{i\tau} \\
\text{timesince}_{i\tau} &= \beta^{\text{lag}}\text{sale}_{i\tau} + \mu_i + \rho_r + t_\tau + \epsilon_{i\tau}
\end{aligned}$$

where $\text{timeto}_{i\tau}$ is the number of days to the next drinks purchase, $\text{timesince}_{i\tau}$ is the number of days since the previous purchase, $\text{sale}_{i\tau}$ is the quantity share of drinks purchased on sale on shopping trip τ by household i, and μ_i, ρ_r, and t_τ are household-purchase string, retailer and time effects.

Stockpiling behavior should lead to $\beta^{\text{lead}} > 0$ and $\beta^{\text{lag}} < 0$. Columns (3) and (4) of Table C.3 summarize the estimates from these regressions. We estimate that purchasing on sale is associated with an increase of 0.14 days to the next purchase and 0.23 days less since the previous purchase. The sign of these effects are consistent with stockpiling, however their magnitudes are very small; the average gap between purchases of drinks is 12 days.

Probability of previous purchase being on sale. A fourth implication highlighted by Hendel and Nevo (2006b) is that stockpiling behavior implies that if a household makes a non-sale purchase today, the probability of the previous purchase being non-sale is higher than if the current purchase was on sale.

We estimate:

$$\text{nonsale}_{i\tau-1} = \beta^{\text{ns}}\text{sale}_{i\tau} + \mu_i + \rho_r + t_\tau + \epsilon_{i\tau}$$

where $\text{nonsale}_{i\tau} = \mathbb{1}[\text{sale}_{i\tau} < 0.1]$ indicates a non-sale purchase, and the other effects are as defined above.

The Hendel and Nevo (2006b) prediction is that $\beta^{\text{ns}} < 0$. Column (5) shows the estimated β^{ns} from this regression. We find that there is a negative relationship between buying on sale today and the previous purchase not being on sale, however, the magnitude of this effect is relatively small.

Table C.3: *Stockpiling evidence*

	(1) Buys drink	(2) Vol. cond. on buying	(3) Days to next	(4) Days since previous	(5) Prev purch on sale
Inventory	0.0009***	0.0127***			
	(0.0001)	(0.0006)			
Purchase on sale?			0.1451***	-0.2263***	-0.0892***
			(0.0198)	(0.0198)	(0.0016)
Mean of dependent variable	0.2271	3.9250	12.1625	12.1625	0.4638
N	8027010	1823157	1692245	1692245	1712051
Time effects?	Yes	Yes	Yes	Yes	Yes
Retailer effects?	Yes	Yes	Yes	Yes	Yes
Decision maker fixed effects?	Yes	Yes	Yes	Yes	Yes

Notes: The dependent variable in column (1) is a dummy variable equal to 1 if the household purchases a non-alcoholic drink on shopping trip τ; in column (2) it is the quantity of drink purchased by household i on shopping trip τ, conditional on buying a positive quantity; in column (3) it is the number of days to the next drink purchase; in column (4) it is the number of days since the previous purchase; and in column (5) it is a dummy variable equal to 1 if the previous purchase was not on sale. Robust standard errors are shown in parentheses.

Sales and product switching. While the evidence suggests that people do not change the timing of their purchases when they buy on sale, this does not imply consumer choice does not respond to price variation resulting from sales. We quantify the propensity of people to switch brands, sizes and pack types (e.g. from bottles to cans) by estimating the following:

$$\begin{aligned}
\text{brandswitch}_{i\tau} &= \beta^{brandswitch}\text{sale}_{i\tau} + \mu_i + \rho_r + t_\tau + \epsilon_{i\tau} \\
\text{sizeswitch}_{i\tau} &= \beta^{sizeswitch}\text{sale}_{i\tau} + \mu_i + \rho_r + t_\tau + \epsilon_{i\tau} \\
\text{packtypeswitch}_{i\tau} &= \beta^{packtypeswitch}\text{sale}_{i\tau} + \mu_i + \rho_r + t_\tau + \epsilon_{i\tau} \\
\text{retailerswitch}_{i\tau} &= \beta^{retailerswitch}\text{sale}_{i\tau} + \mu_i + \rho_r + t_\tau + \epsilon_{i\tau}
\end{aligned}$$

where $\text{brandswitch}_{i\tau}$ is a dummy variable equal to 1 if the household purchased a brand that is different from the brand they bought last, $\text{sizeswitch}_{i\tau}$ is a dummy variable equal to 1 if the household purchased a size that is different from the size they bought last, $\text{packtypeswitch}_{i\tau}$ is a dummy variable equal to 1 if the household purchased a pack type that is different from the pack type they bought last, and $\text{retailerswitch}_{i\tau}$ a dummy variable equal to 1 if the household shopped in a different retailer to their last shopping trip.

Table C.4 shows the estimated β coefficients. We find that buying on sale leads to an increase in the probability of switching brands, sizes and pack types. The percentage effect is largest for pack type switching: buying on sale is associated with an 12.5% (0.016/0.127) increase in the probability that the household switches to buying a new pack type (i.e., cans instead of bottles or vice versa). Buying on sale is associated with a 3.3% and 4.5% increase in probability of switching between brands and sizes, respectively. In contrast, although statistically significant, there is less than a 1% change in the probability of switching retailer. Switching across pack types, brand and sizes in response to sales contributes to the identification of the price preference parameters in our demand model.

Table C.4: *Sales and product switching*

	(1) Brand switch	(2) Size switch	(3) Pack type switch	(4) Retailer switch
Purchase on sale?	0.0181***	0.0234***	0.0160***	0.0035***
	(0.0012)	(0.0012)	(0.0007)	(0.0010)
Mean of dependent variable	0.5432	0.5183	0.1272	0.3566
N	1823157	1823157	1823157	1823157
Time effects?	Yes	Yes	Yes	Yes
Retailer effects?	Yes	Yes	Yes	Yes
Decision maker fixed effects?	Yes	Yes	Yes	Yes

Notes: The dependent variable in column (1) is a dummy variable equal to 1 if the household buys a brand on shopping trip τ that they did not buy on the last trip on which they made a soft drinks purchase; in column (2) it is a dummy variable equal to 1 if the household buys a size on shopping trip τ that they did not buy on the last trip on which they made a soft drinks purchase; in column (3) it is a dummy variable equal to 1 if the household buys a pack type on shopping trip τ that they did not buy on the last trip on which they made a soft drinks purchase; in column (4) it is a dummy variable equal to 1 if the household visits a different retailer on shopping trip τ to their previous trip. Robust standard errors are shown in parentheses.

To summarize, we find very limited evidence of stockpiling behavior in our data; although we cannot conclusively rule it out, any effects are likely to be extremely small.

D Additional demand and supply estimates

We allow all parameters to vary by consumer group and estimate the choice model separately by groups. For estimation: in the at-home segment, for each group, we use a random sample of 1,500 households and 10 choice occasions per household; in the on-the-go sample we use data on all individuals in each group and randomly sample 50 choice occasions per individual, weighting the likelihood function to account for differences in the frequency of choice occasion across consumers.

To calculate the confidence intervals, we obtain the variance-covariance matrix for the parameter vector estimates using standard asymptotic results. We then take 50 draws of the parameter vector from the joint normal asymptotic distribution of the parameters and, for each draw, compute the statistic of interest, using the resulting distribution across draws to compute Monte Carlo confidence intervals (which need not be symmetric).

Demand parameter estimates

Table D.1 summarizes our demand estimates. The top half of the table shows estimates for the at-home segment of the market and the bottom half shows estimates for the on-the-go segment. These include a set of random coefficients over price, a dummy variable for drinks products, a dummy for variable for whether the product contains sugar, a dummy variable for whether the product is 'large' (more than 2l in size for the at-home segment, and 500ml in size in the on-the-go segment), and dummy variables for whether the product is a cola, lemonade, fruit juice, store brand soft drink (at-home only), or a flavored milk (on-the-go only).

Conditional on consumer group, the price random coefficient is log-normally distributed and the other random coefficients are normally distributed; the unconditional distribution of consumer preferences is a mixture of normals. We normalize the means of the random coefficients for the drinks, large, cola, lemonade, store soft drinks and fruit juice effects to zero as they are collinear with the brand-size effects. We allow for correlation within consumer group between preferences for sugar and drinks. For the coefficients on price, branded soft drinks, store brand soft drinks, fruit juice and sugar we allow the mean preferences (within consumer group) to vary by household equivalized income. Note that across 10 of the 12 consumer groups (across both segments) the interaction with the price coefficient is negative and statistically significant – this indicates higher income households are less price sensitive than lower income households. Higher income households also tend to have weaker preferences for store brand soft drinks and products with high sugar content, but stronger preferences for pure fruit juice.

D.1 Elasticities, costs and margins

Table D.2 reports mean market elasticities for a set of popular products in the at-home and on-the-go segments of the market. For each segment, we show elasticities for the most popular size belonging to each of the 10 most popular brand-variants (where variants refer to regular/diet/zero versions). Table D.3 reports the average price, marginal cost and price-cost margin (all per liter) for each brand, as well as the average price-cost mark-up. Numbers in brackets are 95% confidence intervals.

In Figure D.1 we show how prices, marginal costs, and price-cost margins vary with product size. There is strong non-linear pricing; in per liter terms, smaller products are, on average, more expensive. Average marginal costs are broadly constant across the size distribution, with the exception of small single portion sizes, which, on average, have higher costs. Price-cost margins are declining in size – the average margin (per liter) is more than twice as large for the smallest options compared with the largest. This pattern has important implications for tax policy. A tax levied on the sugar in sweetened beverages will result in a higher tax burden (per liter) on large products. To the extent that this causes consumers to switch more strongly away from large products, relative to smaller products, consumers' baskets of taxed products will become more dominated by small, high margin products, which will exacerbate distortions associated with the market power of sugar-sweetened beverages.

Table D.1: *Estimated preference parameters*

At-home		No children			Children		
		low dietary sugar	med. dietary sugar	high dietary sugar	low dietary sugar	med. dietary sugar	high dietary sugar
Mean	Price	0.257	0.356	0.316	0.378	0.411	0.399
		(0.052)	(0.045)	(0.045)	(0.039)	(0.034)	(0.031)
	Sugary:<10g/100ml	1.076	1.048	1.119	0.507	0.851	1.001
		(0.135)	(0.123)	(0.125)	(0.112)	(0.106)	(0.099)
	Sugary:≥10g/100ml	0.541	0.441	0.645	0.102	0.507	0.843
		(0.110)	(0.099)	(0.102)	(0.089)	(0.087)	(0.080)
	Advertising	0.252	0.289	0.230	0.268	0.246	0.311
		(0.055)	(0.053)	(0.051)	(0.045)	(0.040)	(0.039)
Interaction with income	× Price	-0.008	-0.009	-0.011	-0.010	-0.012	-0.010
		(0.002)	(0.002)	(0.002)	(0.002)	(0.002)	(0.002)
	× Branded soft drinks	0.005	-0.005	0.009	-0.015	-0.015	-0.025
		(0.006)	(0.006)	(0.006)	(0.006)	(0.007)	(0.007)
	× Store brand soft drinks	-0.013	-0.006	0.018	-0.029	-0.023	-0.036
		(0.006)	(0.006)	(0.007)	(0.007)	(0.008)	(0.008)
	× Pure fruit juice	0.051	0.021	0.022	0.033	0.037	0.018
		(0.008)	(0.007)	(0.008)	(0.009)	(0.009)	(0.009)
	× Sugary:<10g/100ml	-0.027	-0.009	-0.014	-0.009	-0.018	-0.006
		(0.006)	(0.005)	(0.006)	(0.006)	(0.006)	(0.006)
	× Sugary:≥10g/100ml	-0.029	-0.008	-0.006	-0.022	-0.034	-0.016
		(0.005)	(0.005)	(0.005)	(0.005)	(0.005)	(0.005)
Variance	Price	0.127	0.175	0.165	0.069	0.061	0.075
		(0.019)	(0.020)	(0.019)	(0.010)	(0.009)	(0.009)
	Sugary	2.205	2.308	1.851	1.309	1.644	1.766
		(0.210)	(0.188)	(0.176)	(0.115)	(0.132)	(0.133)
	Drinks	2.217	1.790	1.422	1.296	1.750	1.481
		(0.220)	(0.165)	(0.177)	(0.134)	(0.149)	(0.142)
	Large	0.388	0.458	0.454	0.770	0.360	0.407
		(0.237)	(0.142)	(0.153)	(0.183)	(0.117)	(0.106)
	Cola	2.376	2.026	2.454	1.929	1.960	2.131
		(0.303)	(0.233)	(0.317)	(0.208)	(0.184)	(0.199)
	Lemonade	2.071	2.951	1.574	1.838	1.346	2.280
		(0.466)	(0.495)	(0.288)	(0.457)	(0.267)	(0.304)
	Store brand soft drinks	2.562	2.638	2.191	2.357	2.040	1.805
		(0.241)	(0.250)	(0.224)	(0.202)	(0.167)	(0.147)
	Pure fruit juice	3.176	3.283	3.993	2.823	2.449	2.613
		(0.286)	(0.327)	(0.459)	(0.271)	(0.240)	(0.218)
Covariance	Sugary-Drinks	-1.751	-1.607	-0.704	-0.851	-1.131	-1.041
		(0.192)	(0.157)	(0.146)	(0.100)	(0.117)	(0.121)
On-the-go		Aged under 30			Aged over 30		
		low dietary sugar	med. dietary sugar	high dietary sugar	low dietary sugar	med. dietary sugar	high dietary sugar
Mean	Price	1.069	1.207	0.966	0.868	1.499	1.263
		(0.129)	(0.088)	(0.146)	(0.123)	(0.054)	(0.083)
	Sugary:<10g/100ml	2.641	3.134	2.701	2.806	2.271	0.994
		(0.299)	(0.167)	(0.224)	(0.159)	(0.118)	(0.144)
	Sugary:≥10g/100ml	0.629	1.230	1.215	1.566	0.821	0.064
		(0.205)	(0.104)	(0.130)	(0.119)	(0.095)	(0.090)
	Advertising	0.786	0.666	0.545	0.553	0.457	0.603
		(0.077)	(0.045)	(0.060)	(0.046)	(0.031)	(0.046)
Interaction with income	× Price	0.022	-0.013	0.013	-0.038	-0.076	-0.081
		(0.014)	(0.010)	(0.012)	(0.009)	(0.006)	(0.007)
	× Branded soft drinks	-0.016	0.042	0.047	0.025	-0.021	-0.108
		(0.016)	(0.010)	(0.014)	(0.010)	(0.007)	(0.008)
	× Pure fruit juice	0.148	0.154	0.065	0.028	0.009	-0.145
		(0.027)	(0.018)	(0.023)	(0.021)	(0.011)	(0.013)
	× Flavored milk	-0.090	0.089	-0.019	-0.003	-0.070	-0.106
		(0.023)	(0.015)	(0.020)	(0.018)	(0.011)	(0.014)
	× Sugary:<10g/100ml	-0.001	-0.050	-0.015	-0.081	-0.054	0.103
		(0.010)	(0.006)	(0.008)	(0.007)	(0.005)	(0.006)
	× Sugary:≥10g/100ml	0.037	-0.037	-0.011	-0.080	-0.036	0.082
		(0.009)	(0.005)	(0.007)	(0.006)	(0.004)	(0.005)
Variance	Price	0.531	0.083	0.030	0.273	0.120	0.115
		(0.117)	(0.013)	(0.009)	(0.049)	(0.011)	(0.018)
	Sugary	8.770	4.380	7.230	8.576	7.970	6.333
		(0.724)	(0.235)	(0.534)	(0.423)	(0.320)	(0.354)
	Drinks	3.411	5.532	3.495	5.551	2.968	3.300
		(0.396)	(0.282)	(0.320)	(0.301)	(0.168)	(0.203)
	Large	4.839	4.985	3.630	7.787	4.157	5.667
		(0.397)	(0.299)	(0.248)	(0.412)	(0.171)	(0.304)
	Cola	4.927	5.358	3.660	7.536	7.191	7.472
		(0.398)	(0.294)	(0.300)	(0.418)	(0.284)	(0.346)
	Lemonade	3.383	4.984	6.205	0.793	1.285	5.507
		(0.408)	(0.680)	(0.666)	(0.183)	(0.160)	(0.492)
	Pure fruit juice	17.307	3.295	4.448	8.997	3.006	2.728
		(2.501)	(0.513)	(0.613)	(0.776)	(0.304)	(0.381)
	Flavored milk	5.667	2.251	9.466	4.636	4.140	2.727
		(1.015)	(0.485)	(1.097)	(0.919)	(0.556)	(0.503)
Covariance	Sugary-Drinks	-4.422	-4.503	-3.877	-5.903	-3.879	-3.494
		(0.514)	(0.227)	(0.416)	(0.319)	(0.222)	(0.239)
Brand-size effects		Yes	Yes	Yes	Yes	Yes	Yes
Brand-retailer effects		Yes	Yes	Yes	Yes	Yes	Yes
Size-retailer effects		Yes	Yes	Yes	Yes	Yes	Yes
Brand-time effects		Yes	Yes	Yes	Yes	Yes	Yes
Size-time effects		Yes	Yes	Yes	Yes	Yes	Yes

Notes: Standard errors are reported below the coefficients. For the variance estimates: "Sugary" is an indicator variable equal to 1 if the drink contains any sugar; "Drinks" is an indicator variable equal to 1 for all the inside drinks options; "Large" is an indicator variable equal to 1 if the option is larger than 2l in the at-home segment and 500ml in the on-the-go-segment.

Table D.2: *Price elasticities for popular products*

At-home		Coca Cola Enterprises				Pepsico/Britvic					GSK
		Coke		Capri Sun	Schweppes	Robinsons		Pepsi		Tropicana	Lucozade
		Reg. 2l	Diet 2x2l	10x200ml	Reg. 2x2l	Squash 1l	Fruit diet 1l	Reg. 2l	Max 2x2l	1l	Reg. 6x380ml
Coke	Regular 2l	-2.204	0.020	0.016	0.007	0.019	0.004	0.028	0.027	0.024	0.021
	Diet 2x2l	0.009	-2.721	0.009	0.005	0.008	0.008	0.010	0.058	0.016	0.016
Capri Sun	10x200ml	0.007	0.010	-2.582	0.009	0.025	0.007	0.011	0.013	0.025	0.033
Schweppes Lemonade	Regular 2x2l	0.006	0.010	0.017	-2.261	0.019	0.005	0.007	0.013	0.028	0.037
Robinsons	Squash 1l	0.007	0.008	0.023	0.009	-1.297	0.007	0.011	0.012	0.027	0.028
	Fruit diet 1l	0.004	0.015	0.013	0.005	0.015	-1.319	0.006	0.021	0.019	0.016
Pepsi	Regular 2l	0.023	0.020	0.019	0.007	0.023	0.005	-1.391	0.032	0.022	0.021
	Max 2x2l	0.009	0.047	0.010	0.005	0.010	0.009	0.013	-2.379	0.016	0.016
Topicana	1l	0.004	0.006	0.010	0.006	0.012	0.004	0.005	0.008	-1.944	0.016
Lucozade	Regular 6x380ml	0.006	0.010	0.021	0.012	0.020	0.005	0.007	0.013	0.026	-2.555
Outside option		0.004	0.007	0.009	0.005	0.014	0.006	0.006	0.010	0.019	0.012

On-the-go		Coca Cola Enterprises						Pepsico/Britvic		GSK	
		Coke		Fanta	Dr Pepper	Cherry Coke	Oasis	Pepsi		Ribena	Lucozade
		Reg 500ml	Diet 500ml	Reg 500ml	Reg 500m	Reg 500m	Reg 500ml	Reg 500ml	Max 500ml	Reg 500ml	Reg 330ml
Coke	Regular 500ml	-1.785	0.135	0.051	0.041	0.035	0.079	0.211	0.067	0.022	0.019
	Diet 500ml	0.229	-2.186	0.023	0.019	0.016	0.038	0.067	0.218	0.010	0.008
Fanta	Regular 500ml	0.202	0.054	-2.419	0.113	0.095	0.208	0.066	0.031	0.062	0.042
Dr Pepper	Max 500ml	0.215	0.059	0.150	-2.580	0.104	0.245	0.069	0.034	0.063	0.041
Cherry Coke	Regular 500ml	0.204	0.054	0.138	0.115	-2.401	0.197	0.061	0.028	0.056	0.046
Oasis	Regular 500ml	0.196	0.055	0.129	0.116	0.084	-2.165	0.059	0.029	0.049	0.038
Pepsi	Regular 500ml	0.761	0.143	0.060	0.048	0.038	0.086	-2.256	0.090	0.027	0.024
	Regular 500ml	0.248	0.479	0.029	0.024	0.018	0.043	0.092	-2.355	0.013	0.010
Ribena	Regular 500ml	0.220	0.060	0.159	0.121	0.097	0.202	0.075	0.035	-2.606	0.046
Lucozade	Regular 330ml	0.099	0.026	0.056	0.042	0.043	0.081	0.035	0.014	0.024	-1.813
Outside option		0.065	0.045	0.029	0.023	0.019	0.042	0.023	0.025	0.011	0.042

Notes: Numbers show the mean price elasticities of market demand in the most recent year covered by our data (2012). Number shows price elasticity of demand for option in column 1 with respect to the price of option in row 1.

Table D.3: *Average price-cost margins by brands*

Firm	Brand	Price (£/l)	Marginal cost (£/l)	Price-cost margin (£/l)	(Price-cost) /Price
Coca Cola Enterprises	Coke	1.13	0.34	0.79	0.60
			[0.31, 0.38]	[0.76, 0.83]	[0.58, 0.62]
	Capri Sun	1.17	0.57	0.60	0.50
			[0.55, 0.59]	[0.58, 0.62]	[0.49, 0.52]
	Innocent fruit juice	3.34	1.54	1.80	0.56
			[1.40, 1.65]	[1.69, 1.94]	[0.54, 0.59]
	Schweppes Lemonade	0.52	0.14	0.38	0.71
			[0.13, 0.16]	[0.36, 0.39]	[0.69, 0.74]
	Fanta	1.44	0.41	1.03	0.70
			[0.34, 0.47]	[0.96, 1.10]	[0.67, 0.74]
	Dr Pepper	1.33	0.42	0.91	0.66
			[0.38, 0.47]	[0.87, 0.95]	[0.64, 0.69]
	Schweppes Tonic	1.65	0.72	0.93	0.64
			[0.66, 0.75]	[0.89, 0.99]	[0.62, 0.67]
	Sprite	1.26	0.34	0.92	0.72
			[0.28, 0.38]	[0.88, 0.98]	[0.70, 0.76]
	Cherry Coke	1.53	0.50	1.04	0.62
			[0.43, 0.56]	[0.98, 1.11]	[0.60, 0.66]
	Oasis	2.31	0.53	1.79	0.77
			[0.33, 0.68]	[1.64, 1.98]	[0.71, 0.86]
Pepsico/Britvic	Robinsons	1.20	0.32	0.88	0.74
			[0.29, 0.34]	[0.86, 0.91]	[0.72, 0.76]
	Pepsi	1.02	0.40	0.62	0.61
			[0.37, 0.43]	[0.59, 0.65]	[0.59, 0.63]
	Tropicana fruit juice	2.20	0.94	1.25	0.56
			[0.87, 1.01]	[1.19, 1.33]	[0.54, 0.59]
	Robinsons Fruit Shoot	1.81	0.63	1.18	0.64
			[0.58, 0.68]	[1.13, 1.23]	[0.62, 0.66]
	Britvic fruit juice	2.06	0.90	1.16	0.56
			[0.86, 0.93]	[1.12, 1.20]	[0.55, 0.58]
	7 Up	1.22	0.45	0.77	0.68
			[0.41, 0.48]	[0.74, 0.82]	[0.65, 0.71]
	Copella fruit juice	1.40	0.23	1.17	0.83
			[0.19, 0.27]	[1.13, 1.21]	[0.80, 0.86]
	Tango	1.13	0.34	0.79	0.74
			[0.31, 0.38]	[0.75, 0.82]	[0.71, 0.76]
GSK	Ribena	1.77	0.89	0.88	0.50
			[0.85, 0.92]	[0.86, 0.92]	[0.48, 0.52]
	Lucozade	1.62	0.77	0.85	0.53
			[0.73, 0.82]	[0.80, 0.89]	[0.51, 0.55]
	Lucozade Sport	1.49	0.83	0.65	0.44
			[0.80, 0.86]	[0.63, 0.69]	[0.42, 0.46]
	Vimto	1.09	0.50	0.59	0.57
			[0.48, 0.51]	[0.58, 0.61]	[0.56, 0.59]
	Irn Bru	1.56	0.63	0.93	0.61
			[0.55, 0.72]	[0.84, 1.00]	[0.57, 0.65]
Merrydown	Shloer	1.59	0.71	0.88	0.55
			[0.68, 0.73]	[0.86, 0.91]	[0.54, 0.57]
Red Bull	Red Bull	4.74	2.60	2.15	0.44
			[2.32, 2.81]	[1.93, 2.42]	[0.40, 0.49]
Total		1.44	0.55	0.89	0.62
			[0.51, 0.58]	[0.85, 0.93]	[0.60, 0.64]

Notes: We recover marginal costs for each product in each market. We report averages by brand for the most recent year covered by our data (2012). Margins are defined as price minus cost and expressed in £ per liter. 95% confidence intervals are given in square brackets.

Figure D.1: *Price-cost margins, by product size*

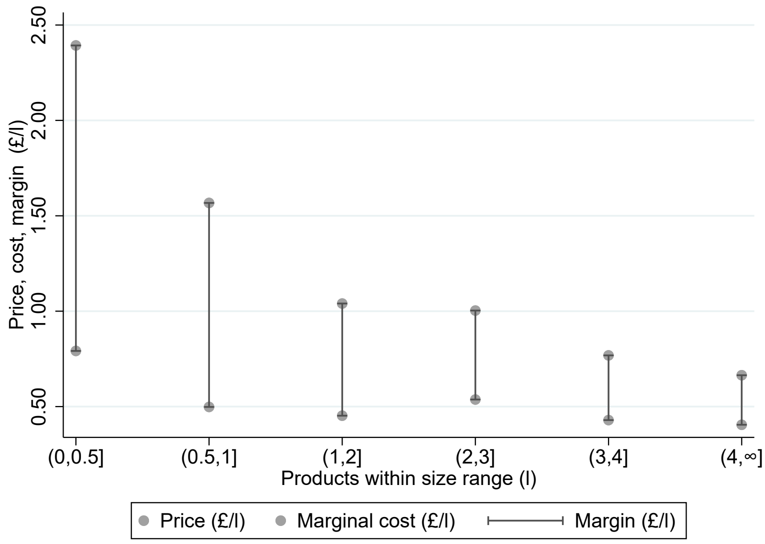

Notes: We group products by size. The figure shows the mean price, cost, and margin (all expressed in £/l) across products within each size range. Numbers are for the more recent year covered by our data (2012).

D.2 Model validation

We use data on the price changes of drinks following the introduction of the UK's Soft Drinks Industry Levy (SDIL) in 2018 to validate our empirical model's tax pass-through predictions. We use a weekly database of UPC level prices and sugar contents for drinks products, collected from the websites of 6 major UK supermarkets (Tesco, Asda, Sainsbury's, Morrisons, Waitrose and Ocado), that cover the period 12 weeks before and 18 weeks after the introduction of the tax (on April 1, 2018).[29] We use data on all the brands included in our demand model, excluding data on minor brands (some of which benefit from a small producers' exemption from the levy).

The SDIL tax is levied per liter of product, with a lower rate of 18p/liter for products with sugar contents of 5-8g/100ml and a higher rate of 24p/liter for products with sugar content $>$ 8g/100m. The tax applies to sugar-sweetened beverages; milk-based drinks and pure fruit juices are exempt from the tax.

We define three sets of products. First, the "higher rate treatment group" are those products with at least 8g of sugar per 100ml, at the time the tax was introduced and therefore are subject to the higher tax rate. Second, the "lower rate treatment group" are those products that have 5-8g of sugar per 100ml, and

[29]We are grateful to the University of Oxford for providing us with access to these data, which were collected as part of the foodDB project.

therefore are subject to the lower tax rate. The remaining set of products are exempt, either because their sugar content is less than 5g per 100ml, or because they are milk-based or fruit juice. There was some reformulation in anticipation of the introduction of the SDIL. We categorize products based on the post reformulation sugar contents.[30]

We estimate price changes for the two treatment and the exempt groups. Let j index product, r retailer, and t week. We define the dummy variables $\text{TreatHi}_j = 1$ if product j is in the high treatment group, $\text{TreatLo}_j = 1$ if product j is in the low treatment group, and $\text{TreatExempt}_j = 1$ if product j is exempt from the tax. Let Post_t denote a dummy variable equal to 1 if $t >= 13$ i.e. weeks following the introduction of the tax. We estimate the following regression, pooling across products in each of the three groups:

$$p_{jrt} = \beta^{hi}\text{TreatHi}_j \times \text{Post}_t + \beta^{lo}\text{TreatLo}_j \times \text{Post}_t + \sum_{t \neq 12} \tau_t + \xi_j + \rho_r + \epsilon_{jrt} \quad \text{(D.1)}$$

where p_{jrt} denotes the price per liter of product j in retailer r in week t,[31] τ_t are week effects, ξ_j are product fixed effects, and ρ_r are retailer fixed effects.

Figure D.2(a) plots the estimated price changes, relative to the week preceding the introduction of the tax, for the higher rate treatment group ($= \hat{\beta}^{hi} \times \text{Post}_t + \sum_{t \neq 12} \hat{\tau}_t$). Figure D.2(b) plots the analogous estimates for the lower rate treatment group ($= \hat{\beta}^{lo} \times \text{Post}_t + \sum_{t \neq 12} \hat{\tau}_t$). Figure D.2(c) plots the estimates for the group of products exempt from the tax ($\sum_{t \neq 12} \hat{\tau}_t$). The solid blue line plots the tax per liter. The data suggest that there was slight overshifting of the tax, with an average price increase among the high treatment group of 26p per liter (a pass-through rate of 108%), and the average price increase among the low treatment group of 19p per liter (a pass-through rate of 105%). The prices of products not subject to the tax do not change following its introduction.

We simulate the introduction of the SDIL using our estimated model of demand and supply in the non-alcoholic drinks market (based on product sugar contents when the SDIL was implemented). The red lines plot the average price increase for each of the three group predicted by our model. These match very closely the actual price increases following the policy's introduction.

[30] We exclude a small number of products belonging to the Irn Bru and Shloer brands that were reformulated approximately 10 weeks after the introduction of the tax.

[31] This is the VAT-exclusive price per liter.

Figure D.2: *Out-of-sample model validation: UK Soft Drinks Industry Levy*

(a) High treatment group

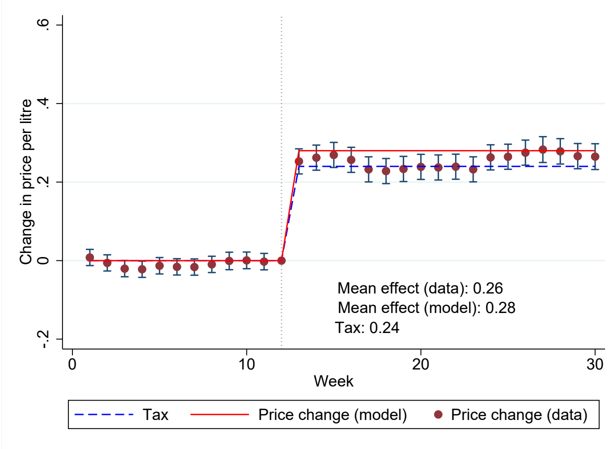

(b) Low treatment group

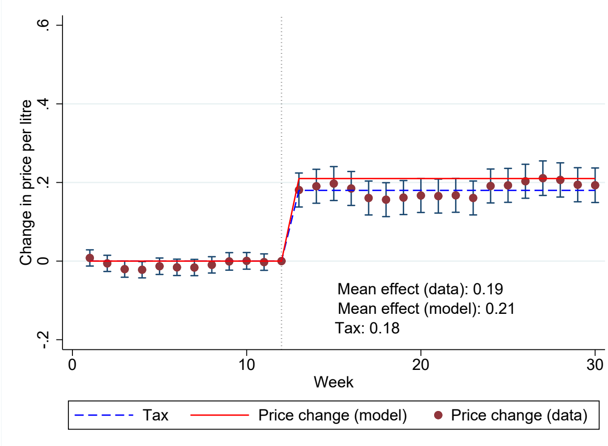

(c) Exempt group

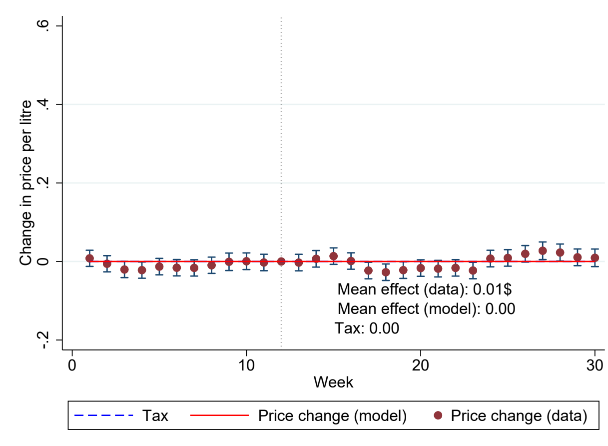

Notes: Grey markers show the estimated price changes (relative to the week preceding the introduction of the tax). For the higher rate treatment group (top panel), the estimated prices changes are $= \hat{\beta}^{hi} Post_t + \sum_{t \neq 12} \hat{\tau}_t$*, for the lower rate treatment group (middle panel), the estimated price changes are* $= \hat{\beta}^{lo} Post_t + \sum_{t \neq 12} \hat{\tau}_t$*, and for the exempt group (bottom panel) they are* $= \hat{\tau}_t$ *All coefficients are estimated jointly (equation (D.1)). 95% confidence intervals shown. The blue line shows the value of the tax, and the red line shows the predicted price changes from our estimated demand and supply model.*

E Implementation of optimal tax problem

E.1 Externality calibration

We use two different approaches to estimate the average externality costs associated with reductions in sugar-sweetened beverage consumption. Both yield similar estimates.

Approach using evidence from Wang et al. (2012)

Wang et al. (2012) consider the impact of a fall of approximately 15% in sugar-sweetened beverage consumption among adults aged 25-64 on health care costs in the US. They conclude it would result in savings of $17.1 billion realized over 10 years, discounted at a rate of 3% per year.

As a baseline, they use an average daily serving of 0.56 and serving size of 170kcal. They simulate a reduction in sugar-sweetened beverage consumption to 0.47 daily servings, which translates into a fall in calories from these products of (0.56-0.47)*170=15kcal per adult per day.[32] This corresponds to a 3.75g fall in sugar per adult per day. Their estimate of health care cost savings of $17.1 billion over 10 years corresponds to an average daily fall of $4.7 million, or 2.7¢ per adult (based on 171 million Americans aged 15-64). Hence, the implied health cost saving is $2.7/0.375 \approx 7$¢ per 10g of sugar.

We convert the average health care saving to UK numbers by applying a $-£ exchange rate of 0.75 and deflating by an estimate of the cost of providing health care in the UK relative to US (equal to 0.83 and based on OECD (2019)). This yields an average health care cost saving of approximately 4 pence per 10g of sugar. Health care in the UK is almost entirely provided by the taxpayer funded National Health Service, so we treat this as an externality.

Approach using evidence from Briggs et al. (2013)

Briggs et al. (2013) use a comparative risk assessment model, which maps dietary changes into health outcomes, to estimate the impact of a 1g reduction in sugar per adult per day into changes in the prevelance of obesity and overweight.[33] They find that the fall in consumption would lead to a 1.3% fall in the number of obese adults and a 0.9% fall in the number of overweight adults. This implies a 1.1% reduction

[32]Note, they assume 40% of the calories are displaced so refer to an 9kcal reduction.

[33]The 1g reduction corresponds to their estimate of the impact of 20% tax on sugar-sweetened beverages.

in the number of adults who are overweight or obese (27% of UK adults are obese and 36% are overweight).

Public Health England (2017) estimate that the NHS spent £6.1 billion on overweight and obesity-related ill health in 2014/15. A 1.1% reduction in these costs implies a fall of 0.37 pence per adult per day (based on 49.3 million adults in the UK in 2014). This yields a healthcare cost saving of approximately 3.7 pence per 10g sugar.

Distribution of the external costs

Based on the World Health Organization's official recommendation that individual added sugar consumption should be below 10% of dietary calories we assume that only consumers with dietary sugar above this threshold create externalities. This group comprises around 80% of consumers, so this implies an externality per 10g of sugar of 5 pence per 10g of sugar for this group. Since, on average, sugar-sweetened beverages have 26g of sugar per 10 oz, this implies an *average* externality of 14 pence per 10 oz of sugar-sweetened beverage for people in this group; however individual marginal externalities will vary with the sugar content of sweetened beverages chosen.

E.2 Distribution of profits

In Section 6.2 of the paper we investigate how distributional concerns change the optimal tax rate on sugar-sweetened beverages. This requires information on how profits are distributed between the government, domestic residents and the portion that flow overseas. Measuring the distribution of profits across individuals is challenging and a topic of much recent work. It forms a key part of the "Distributional National Accounts" (DINA) approach, pioneered by Piketty et al. (2018), whose goal is to allocate all national income to individuals. Nonetheless, the methods developed in the literature are subject to considerable debate. We use a method to allocate profits that is inspired by this research, and implement it to the best extent possible using publicly available data.

The government's share of profits comprises corporate tax revenue, as well as revenue from the personal taxation of individual profit holdings. We set the share of profits that the government collects through corporate taxation to 25%, based on the effective average corporate tax rate in 2012 (estimated by the University of Oxford's Centre for Business Taxation (Bilicka and Devereux (2012)). We assume that 30% of profits flow overseas. We calculate this using information from the

UK National Accounts: we take the ratio of income distributed by corporations that flows overseas to net operating surplus excluding imputed rents from owner-occupied housing.

We assume that the remaining 45% of profits are distributed to UK households in proportion to the share of dividend income that they receive. Saez and Zucman (2016) use a combination of dividend income and realized capital gains to estimate stock ownership. Smith et al. (2020) use a weighted average of dividend income and capital gains, but with most of the weight assigned to dividend income, which they find a better predictor of stock ownership. There is no publicly available data that contains information on the joint distribution of capital gains and taxable income for UK individuals. Instead, we use the Survey of Personal Incomes (SPI), which records dividend income received by individuals to estimate the relationship between individuals' total income and the amount they receive from dividends. Table E.1 shows the mean dividend income for individuals with different levels of total income. Individuals earning more than £40,000 (roughly the top 10%), receive approximately 70% of dividend income recorded on tax records.

We map this into the share of dividend income received by *households* (as opposed to individuals). To do this, we use the Living Costs and Food Survey, which contains information on the total (but not dividend) income received by individual household members. We use the mean dividend income by banded personal income shown in Table E.1 to impute dividend income for individuals in participating households in the Living Costs and Food Survey 2012, which we then sum for all members in the household. Table E.2 shows the distribution of dividend income across households. Note that the distribution across households is less skewed than the distribution across individuals, reflecting the fact that many households consist of one high and one lower earner.

Table E.1: *Mean dividend income by banded personal total income*

Total income	Mean dividend income
0-2.5k	39
2.5-5k	51
5-7.5k	46
7.500-10k	56
10-12.5k	73
12.5-15k	98
15-20k	143
20-30k	302
30-40k	758
40k+	3436

Notes: We use data from the Survey of Personal Incomes in 2012. The table shows the mean dividend income (excluding dividends received from owner-managed companies) for individuals with total personal incomes in the bands shown in the first column.

Table E.2: *Distribution of dividends across household equivalized income distribution*

(1) Equivalized hh income	(2) % hh	(3) Mean div income	(4) ATR divs	(5) % dividends Pre-tax	(6) % dividends Post-tax
0-5k	12.8	60	0.00	0.8	0.9
5-10k	11.9	141	0.00	1.8	2.0
10-15k	17.6	173	0.00	3.3	3.7
15-25k	29.6	508	0.00	16.3	17.7
25-35k	13.5	1721	0.04	25.2	25.8
34-45k	9.3	2933	0.10	29.4	28.5
45k+	5.3	4021	0.17	23.2	21.3

Notes: We use the mean dividend income by banded personal income shown in Table E.1 to impute dividend income for individuals in participating households in the Living Costs and Food Survey 2012. We sum dividend income for all members in the household. We construct equivalized total household income (using the OECD-modified equivalence scale) and put households into bands, listed in column (1). Column (2) shows the share of households in each band, column (3) shows the mean amount of dividend income per household for each band, and column (4) shows the average personal tax paid on dividends for households in each band. Columns (5) and (6) show the share of total dividend income (pre and post- dividend tax, respectively) that households in each band receive.

Dividend income is subject to personal taxation. Table E.1 reports the average tax rate on dividends for each household income band. After taking account of this (and corporate tax), the government share in profits is 29%. Post-tax profits are distributed to households according to column (6) in the table.

E.3 Solution algorithm

Obtaining the optimal tax rate (or vector of rates) entails solving an algorithm that consists of an outer loop and several inner loops. The solution of the outer loop is the optimal tax vector, the solution to the inner loops are, given a candidate tax vector, the equilibrium price vector and the matrix of derivatives of the optimal price vector with respect to the tax vector.

Inner loops Given the tax vector $(\tau_1, \ldots, \tau_K)$, equilibrium prices, $\boldsymbol{p'} = (p'_1, \ldots, p'_J)$, are obtained as the solution to the system of equations: for $j = 1, \ldots, J$

$$q_j(\boldsymbol{p'}) + \sum_{j' \in \mathcal{J}_f} (p'_{j'} - \mathbb{1}\{j' \in \mathcal{J}_k\}\tau_k - c_{j'}) \frac{\partial q_{j'}(\boldsymbol{p'})}{\partial p_j} = 0.$$

The $J \times K$ matrix of derivatives $\frac{d\mathbf{p}'}{d\boldsymbol{\tau}}$ is obtained by solving $k = 1, \ldots, K$ systems of equations of the form: for $j = 1, \ldots, J$

$$\sum_{j' \in \mathcal{M}} \frac{\partial q_j}{\partial p'_{j'}} \frac{dp_{j'}}{d\tau_k} + \sum_{j' \in \mathcal{J}_f} \left(\frac{dp'_{j'}}{d\tau_k} - \mathbb{1}\{j' \in \mathcal{J}_k\} \right) \frac{\partial q_{j'}}{\partial p_j} +$$
$$\sum_{j' \in \mathcal{J}_f} (p'_{j'} - \mathbb{1}\{j' \in \mathcal{J}_k\}\tau_k - c_{j'}) \sum_{j'' \in \mathcal{M}} \frac{\partial^2 q_{j'}}{\partial p_j \partial p_{j''}} \frac{dp'_{j''}}{d\tau_k} = 0.$$

Outer loop We use three alternative methods for solving the outer loop:

1. The optimal tax vector can be expressed in the form: $\boldsymbol{\tau}^* = G(\boldsymbol{\tau}^*)$ (see equation (2.2), for the case of a single sugar-sweetened beverage tax rate). One solution method involves iterating on this equation: (1) guess a tax vector $\boldsymbol{\tau}^r$, (2) solve the inner loops, (3) compute $G(\boldsymbol{\tau}^r)$, (4) set $\boldsymbol{\tau}^{r+1} = G(\boldsymbol{\tau}^r)$ and repeat until convergence. This method is relatively quick but has the disadvantage that it is not suitable for imposing constraints on the government's objective function.

2. When solving for the multi-tax rate system subject to constraints (see Section 6.3 of the paper) we instead numerically maximize the social welfare function subject to the constraint. For each iteration of the algorithm, we must solve the inner loops.

3. A third solution method is a grid search over the tax rate (feasible in the single rate, but not multi rate, case). We use this method to draw Figure E.1.

E.4 Multi-rate tax system

We consider two variants of the multi-rate system. The first variant allows the government to vary taxes rates across 12 different drink types,[34] subject to there being no deterioration in the government's budget, inclusive of the budgetary externality. The second variant requires all tax rates be non-negative, ruling out subsidies on alternatives to sugar-sweetened beverages.

Table E.3 summarises the results. The optimal system under no budget deterioration involves subsidizing non-sugar sweetened drinks, and entails, on average, lower tax rates on sugar-sweetened beverages than under the optimal single rate system. This leads to increases in consumer surplus and profits, with the fall in tax revenue offset by a reduction in budgetary external costs. Overall, welfare is 80% higher than under the optimal single rate system. The third column shows the optimal policy when subsidies are prohibited – this leads to a smaller rise in welfare (of 17%) relative to the optimal single rate system.

[34]These are: 5 sugar-sweetened beverage drinks types (cola, lemonade, other sodas, juices and energy/sports drinks) and 7 drinks types comprising alternatives to sugar-sweetened drinks (pure fruit juices, milk drinks, plus diet counterparts of cola, lemonade etc.).

Table E.3: *Multi-rate taxation*

	(1) Single rate	(2) Multi-rate: No budget deterioration	(3) Multi-rate: No subsidy
Tax rate (p/10oz) for:			
Sin products	5.97	4.28*	5.53*
Alternatives	0.00	-4.01*	0.00*
Price change for:			
Sin products	27.3%	22.5%	27.1%
Alternatives	-0.7%	-15.8%	-0.6%
Consumption change for:			
Sin products	-37.7%	-37.7%	-35.8%
Alternatives	9.4%	35.1%	8.5%
Welfare components (£m):			
Private welfare	-790	299	-670
Consumption	*-747*	*276*	*-634*
Profit holdings	*-43*	*23*	*-37*
Gov. budget	957	0	866
Excise tax rev.	*522*	*-572*	*424*
Profits tax rev.	*-74*	*40*	*-63*
Ext. cost savings	*509*	*532*	*505*
Total welfare (£m)	167	299	196

* *average tax rate*

Notes: Each column corresponds to a tax system as described in the text. Numbers summarize the effect of policy when the social marginal welfare weight on foreign individuals is 0 and on domestic individuals is $1/\tilde{y}_i$*. Welfare numbers are per annum and report the change relative to no drinks taxation. Total welfare = Private welfare+Government budget.*

E.5 Sugar tax

Table E.4 compares the optimal sugar tax, assuming no changes in products' sugar contents, to the optimal single rate volumetric tax. The sugar tax, which entails a rate of 2.28 pence per 10g of sugar, results in a larger increase in the average price of sugar-sweetened beverages (30.3% relative to 27.3% under the volumetric tax). This leads to slightly larger falls in consumer welfare and lower tax revenue, but a considerably larger fall in externality costs (£572m compared with £509m under the volumetric tax). Overall the sugar tax raises welfare by £208m, which is higher than the £167m rise under the optimal single rate volumetric tax. Similarly to volumetric taxation, the costs, in terms of forgone welfare gains, of ignoring market power when setting the sugar tax are substantial (at 31%).

Table E.4: *Sugar taxation*

	% Δ in SSB price	Change (relative to zero tax) in: Welfare components (£m) Private, welfare, from: Cons.	Profits	Tax revenue: Sin tax	Profit tax	Ext. cost savings	Total welfare	% loss under naivety
Volumetric tax	27.3%	-747	-43	522	-74	509	167	40%
Sugar tax	30.3%	-762	-42	511	-72	572	208	31%

Notes: Row (1) shows the effects of the optimal volumetric tax for reference. It repeats some of the information in row (1) of Table 6.2. Row (2) shows the effects of the optimal tax on the sugar in sweetened beverages under the assumption of no product reformulation. Numbers in the final column show the % of welfare gains from optimal policy forgone when the government sets policy ignoring market power. Numbers summarize the effect of policy when the social marginal welfare weight on foreign individuals is 0 and on domestic individuals is $1/\tilde{y}_i$. Welfare numbers are per annum and report the change relative to no drinks taxation. Total welfare = Private welfare+Tax revenue+External cost savings. Numbers in the final column show the % of welfare gains from optimal policy forgone when the government sets policy ignoring market power.

We also consider the effect of a sugar tax when firms reoptimize both price and the sugar content of their products. We model firms' decision over product sugar content following Barahona et al. (2021). In their model a firm can substitute sugar for an alternative input, whilst keeping the taste (and hence consumer valuation) of the product unchanged, but at the cost of increasing production costs. With no tax in place the firm will choose the cost minimizing sugar level. With a tax in place the firm may change product sugar content, trading-off increased production costs with a reduced tax liability.

Consider firm $f = 1, \dots, F$, which owns products $\mathcal{J}_f$ – it chooses the vector of tax-inclusive prices for these products $\{p_j\}_{j\in\mathcal{J}_f}$. Denote the subset of products in $\mathcal{J}_f$ that are sugar-sweetened beverages by $\mathcal{J}_f^{\mathcal{S}}$ and the remaining products by $\mathcal{J}_f^{\mathcal{N}}$. The firm chooses the sugar content of each product in set $\mathcal{J}_f^{\mathcal{S}}$. We denote by z_j^* the production cost-minimizing sugar content of product $j \in \mathcal{J}_f^{\mathcal{S}}$ (conditional on the taste of the product).

In the absence of a sugar tax, the firm's problem is

$$\max_{\{p_j\}_{j\in\mathcal{J}_f},\{z_j\}_{j\in\mathcal{J}_f^{\mathcal{S}}}} \sum_{j\in\mathcal{J}_f^{\mathcal{S}}} (p_j - c_j(z_j))q_j(\mathbf{p}) + \sum_{j\in\mathcal{J}_f^{\mathcal{N}}} (p_j - c_j)q_j(\mathbf{p})$$

The first order conditions are: for $f = 1, \ldots, F$

$$q_j + \sum_{j' \in \mathcal{J}_f^{\mathcal{S}}} (p_{j'} - c_{j'}(z_{j'})) \frac{\partial q_{j'}}{\partial p_j} + \sum_{j' \in \mathcal{J}_f^{\mathcal{N}}} (p_{j'} - c_{j'}) \frac{\partial q_{j'}}{\partial p_j} = 0 \quad \text{for all } j \in \mathcal{J}_f$$

$$c_j'(z_j) = 0 \quad \text{for all } j \in \mathcal{J}_f^{\mathcal{S}}$$

By definition, the sugar contents that satisfy these conditions are $z_j = z_j^*$ for all $j \in \mathcal{J}_f^{\mathcal{S}}$ and all f.

With a sugar tax in place, we can define the tax-inclusive marginal cost as $C_j(z_j) = \tau z_j + c_j(z_j)$ for all $j \in \mathcal{J}_f^{\mathcal{S}}$ and f. The first order conditions that characterize the firms' optimal choices are then: for $f = 1, \ldots, F$

$$q_j + \sum_{j' \in \mathcal{J}_f^{\mathcal{S}}} (p_{j'} - C_{j'}(z_{j'})) \frac{\partial q_{j'}}{\partial p_j} + \sum_{j' \in \mathcal{J}_f^{\mathcal{N}}} (p_{j'} - c_{j'}) \frac{\partial q_{j'}}{\partial p_j} = 0 \quad \text{for all } j \in \mathcal{J}_f$$

$$C_j'(z_j) = 0 \quad \text{for all } j \in \mathcal{J}_f^{\mathcal{S}}$$

Hence the optimal sugar choice of product k satisfies: $\tau + c_k'(z_k) = 0$.

We assume that the marginal costs function takes the following quadratic form:

$$c_j = \bar{c}_j + \frac{\nu}{z_j^*}(z_j^* - z_j)^2,$$

where $\bar{c}_j$ denotes the cost-minimizing marginal cost (which corresponds to production decisions in the absence of a sugar tax) and ν controls the marginal cost of reformulation. Along with the firm's first order condition for sugar choice, this implies that with a sugar tax in place:

$$\frac{(z_j^* - z_j)}{z_j^*} = \frac{\tau}{2\nu}.$$

Hence the percentage reduction in a product's sugar content is proportional to the sugar tax rate τ and inversely proportional to the reformulation cost ν. Under a sugar tax the increase in the tax-inclusive marginal cost of product j is:

$$\begin{aligned}\Delta C_j(\nu, \tau) =& \tau \left(z_j^* - \frac{z_j^* \tau}{2\nu} \right) + \frac{\nu}{z_j^*} \left(\frac{\tau z_j^*}{2\nu} \right)^2 \\ =& \tau z_j^* - \frac{z_j^* \tau^2}{4\nu}.\end{aligned}$$

Note that, the sugar tax changes the relative marginal cost of two sugary products according to:

$$\frac{\Delta C_k(\nu,\tau)}{\Delta C_j(\nu,\tau)} = \frac{z_k^*}{z_j^*}.$$

Hence, for every reformulation cost, ν, there is a sugar tax rate that results in the same vector of tax-inclusive costs, $\{C_j(\nu,\tau)\}_{j\in\mathcal{M}}$, and hence equilibrium prices and quantities and consumer surplus and profits.

In Figure E.1 we compare the implications of a sugar tax when reformulation costs are prohibitive and when they are relatively low (ν=5 pence). Instead of plotting how welfare varies with the tax rate, τ, we plot how it varies with $\frac{\Delta C}{z^*}$ (which is a monotonically increasing function of τ). Conditional on $\frac{\Delta C}{z^*}$ the market equilibrium is the same (regardless of reformulation costs), and welfare differences as reformulation costs fall are driven purely by whether larger reductions in external costs offset higher production costs. In Figure E.2 we plot how the optimal sugar tax rate and associated welfare gain vary with the reformulation costs. As the cost of reformulation falls, firms choose to remove more of the sugar from their sugar-sweetened beverages. The figure shows that this results in larger welfare gains from optimal sugar taxation. This reason for this is that larger falls in external costs from sugar outweigh raised production costs. Even though firms make privately optimal decisions over product sugar content, the externalities form sugar are sufficiently large that these private decision improve social welfare.

Figure E.1: *Impact of sugar tax with input substitution*

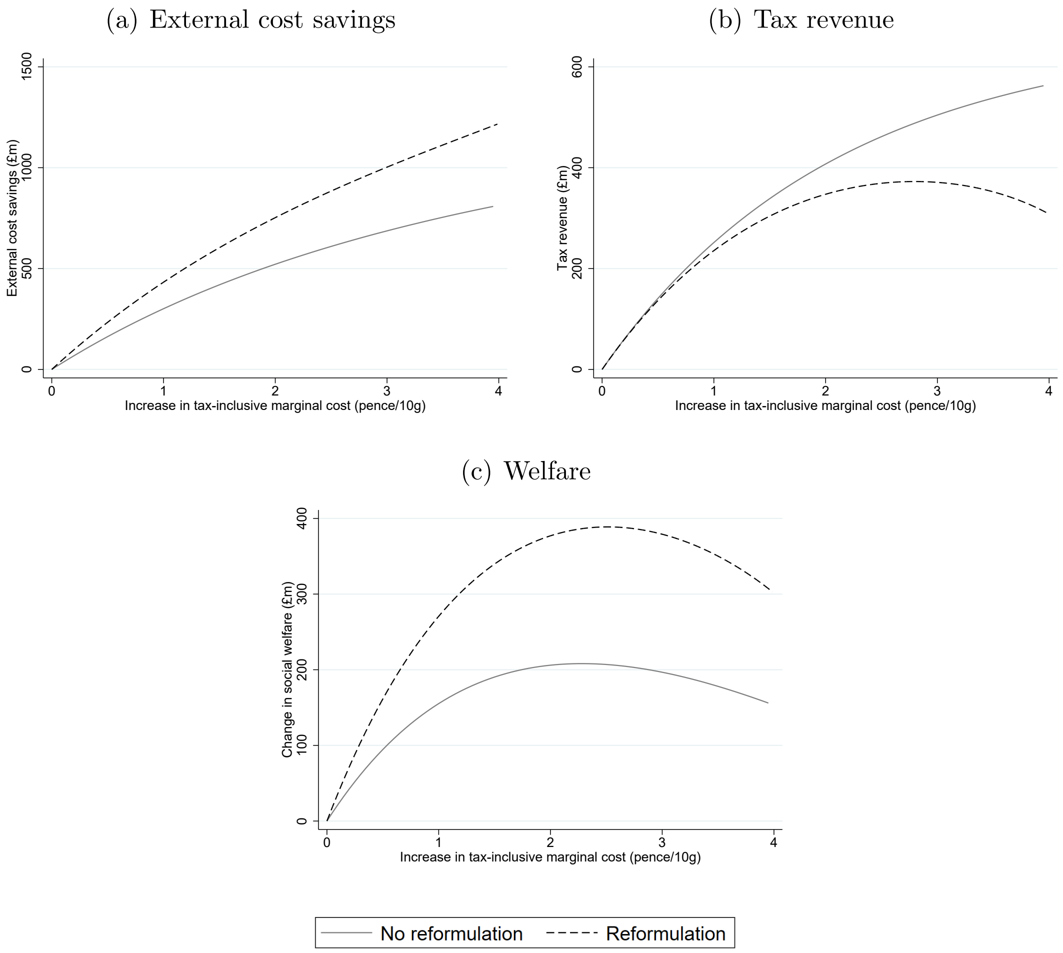

Notes: Graphs show the impact on external costs (panel (a)), tax revenue (panel (b)) and social welfare (panel (c)) from sugar taxation when firms do not reformulate products, and optimally reformulate. On the horizontal axis we plot the increase in the tax-inclusive marginal cost due to a sugar tax. This is a monotonic function of the sugar tax rate, and means that conditional on a given marginal cost increase, the implications of the sugar tax for private welfare are the same in the two scenarios.

Figure E.2: *Variation in optimal sugar tax and welfare gain with reformulation costs*

(a) Optimal tax rate

(b) Welfare gain

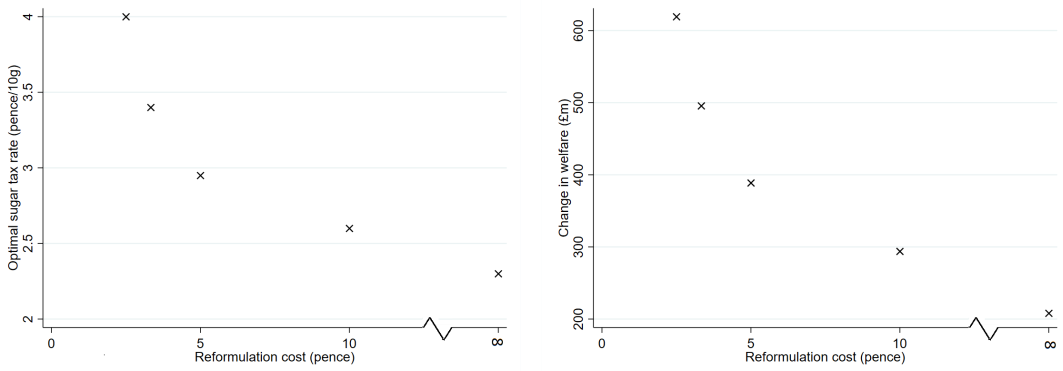

Notes: Graphs show how the optimal sugar tax rate (panel (a)) and associated welfare gain (panel (b)) vary with the reformulation cost parameter ν.